

# CS 480

## *Introduction to Artificial Intelligence*

March 28, 2024

# Announcements / Reminders

- Please follow the Week 11 To Do List instructions (if you haven't already):
- Programming Assignment #02: due on Sunday ~~(03/31 04/07)~~ at 11:59 PM CST
- Written Assignment #04 due on Sunday (03/31) at 11:59 PM CST
- Quiz #06 due on Sunday (03/31) at 11:59 PM CST

# Plan for Today

- Decision Networks
- A Casual Introduction to Machine Learning (if time permits)

# Value of Perfect Information

The value/utility of best action  $\alpha$  without additional evidence (information) is :

$$MEU(\alpha) = \max_{\alpha} \sum_{s'} P(Result(\alpha) = s') * U(s')$$

If we include new evidence/information ( $E_j = e_j$ ) given by some variable  $E_j$ , value/utility of best action  $\alpha$  becomes:

$$MEU(a_{e_j} | e_j) = \max_{\alpha} \sum_{s'} P(Result(\alpha) = s' | e_j) * U(s')$$

The value of additional evidence/information from  $E_j$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(\alpha)$$

using our current beliefs about the world.

# Value of Perfect Information

The value of additional evidence/information from  $E_j$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(a)$$

Maximum expected utility based  
on everything we know so far:

- a) NO evidence whatsoever
- b) OTHER (not value of  $E_j$ )  
evidence (could be more than  
one “piece” of evidence)

# Value of Perfect Information

The value of additional evidence/information from  $E_j$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(a)$$

Value of perfect information for some NOT YET known (but we are considering obtaining / observing it) evidence (specific value of variable  $E_j$ )  $E_j$

# Value of Perfect Information

The value of additional evidence/information from  $E_j$  is:

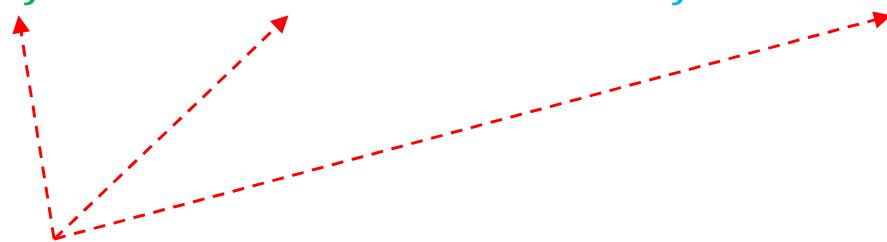
$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(a)$$

Measure of how “much  
EXPECTED utility” we can gain by  
introduction evidence (specific  
value of variable  $E_j$ )  $E_j$

# Value of Perfect Information

The value of additional evidence/information from  $E_j$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(a)$$

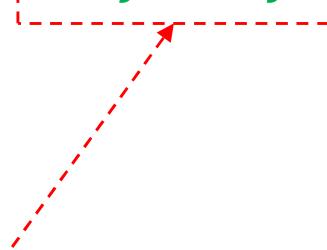


Possible value of variable  $E_j$

# Value of Perfect Information

The value of additional evidence/information from  $E_j$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(a)$$



Probability of variable  $E_j$  having  
value  $e_j$

# Value of Perfect Information

The value of additional evidence/information from  $E_j$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(a)$$

Calculate for every  $e_j$  and sum it

# Value of Perfect Information

The value of additional evidence/information from  $E_j$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * \boxed{MEU(a_{e_j} | E_j = e_j)} \right) - MEU(a)$$

Maximum Expected Utility in the system under the assumption

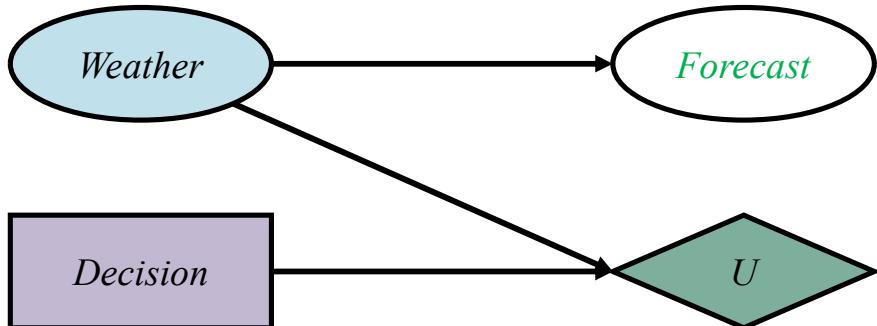
that is  $E_j = e_j$  is given

and everything else (could be nothing) that contributed to

$MEU(a)$  is given

# Decision Network: Example

Decision network



The value of best action  $\alpha$  without additional evidence

$$MEU(\alpha) = MEU(\text{leave}) = 70$$

With evidence information ( $E_j = e_j$ ) given by Forecast:

$$MEU(a_{e_1} | e_1) = MEU(\text{take} | F = \text{rain}) = 53$$

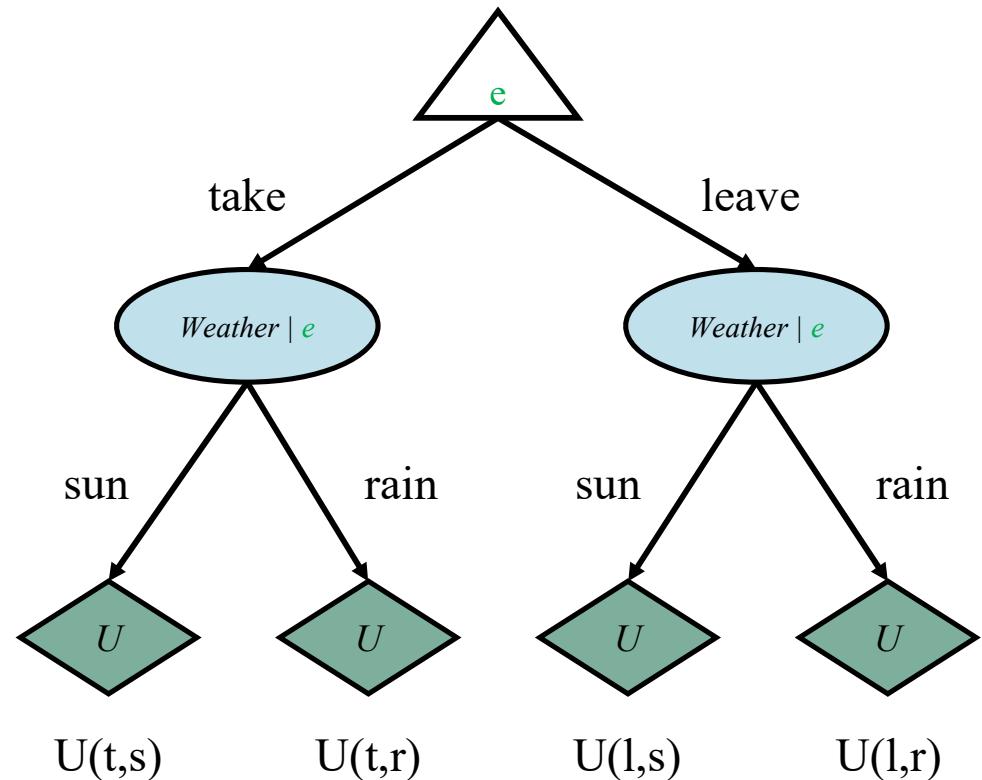
$$MEU(a_{e_2} | e_2) = MEU(\text{leave} | F = \text{sun}) = 95$$

The value of additional evidence / information from  $F$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(\alpha)$$

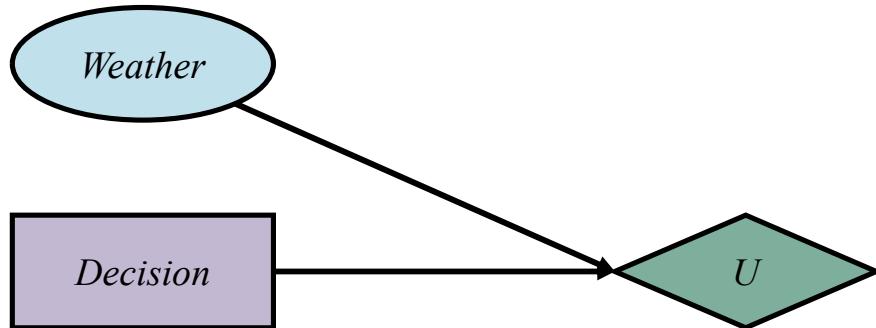
$$\begin{aligned} VPI(F) &= (P(F = \text{rain}) * MEU(\text{take} | F = \text{rain}) + P(F = \text{sun}) * \\ &\quad MEU(\text{leave} | F = \text{sun})) - MEU(\text{leave}) = \\ &\quad (0.41 * 53 + 0.59 * 95) - 70 = 7.78 \end{aligned}$$

Outcome tree



# Decision Networks: Example

## Decision: leave umbrella



$$EU(\text{leave}) = 70$$

The value of best action  $\alpha$  without additional evidence

$$MEU(\alpha) = MEU(\text{leave}) = 70$$

With evidence information ( $E_j = e_j$ ) given by Forecast:

$$MEU(a_{e_1} | e_1) = MEU(\text{take} | F = \text{rain}) = 53$$

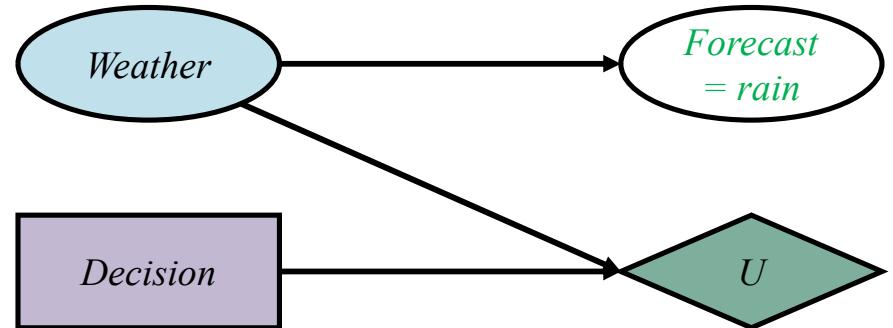
$$MEU(a_{e_2} | e_2) = MEU(\text{leave} | F = \text{sun}) = 95$$

The value of additional evidence / information from  $F$  is:

$$VPI(E_j) = \left( \sum_{e_j} P(E_j = e_j) * MEU(a_{e_j} | E_j = e_j) \right) - MEU(\alpha)$$

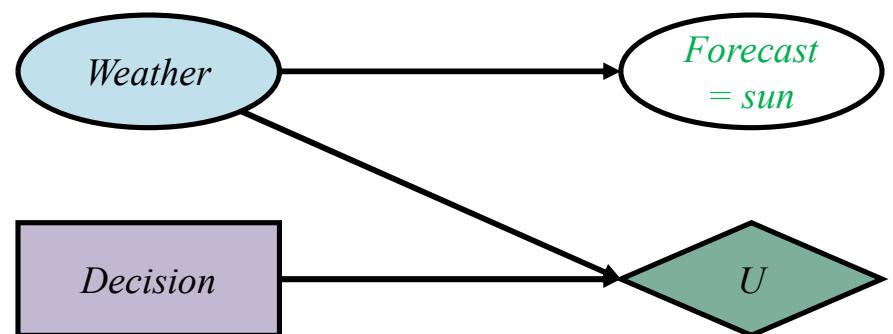
$$\begin{aligned} VPI(F) &= (P(F = \text{rain}) * MEU(\text{take} | F = \text{rain}) + P(F = \text{sun}) * \\ &\quad MEU(\text{leave} | F = \text{sun})) - MEU(\text{leave}) = \\ &\quad (0.41 * 53 + 0.59 * 95) - 70 = 7.78 \end{aligned}$$

## Decision: take umbrella given rain



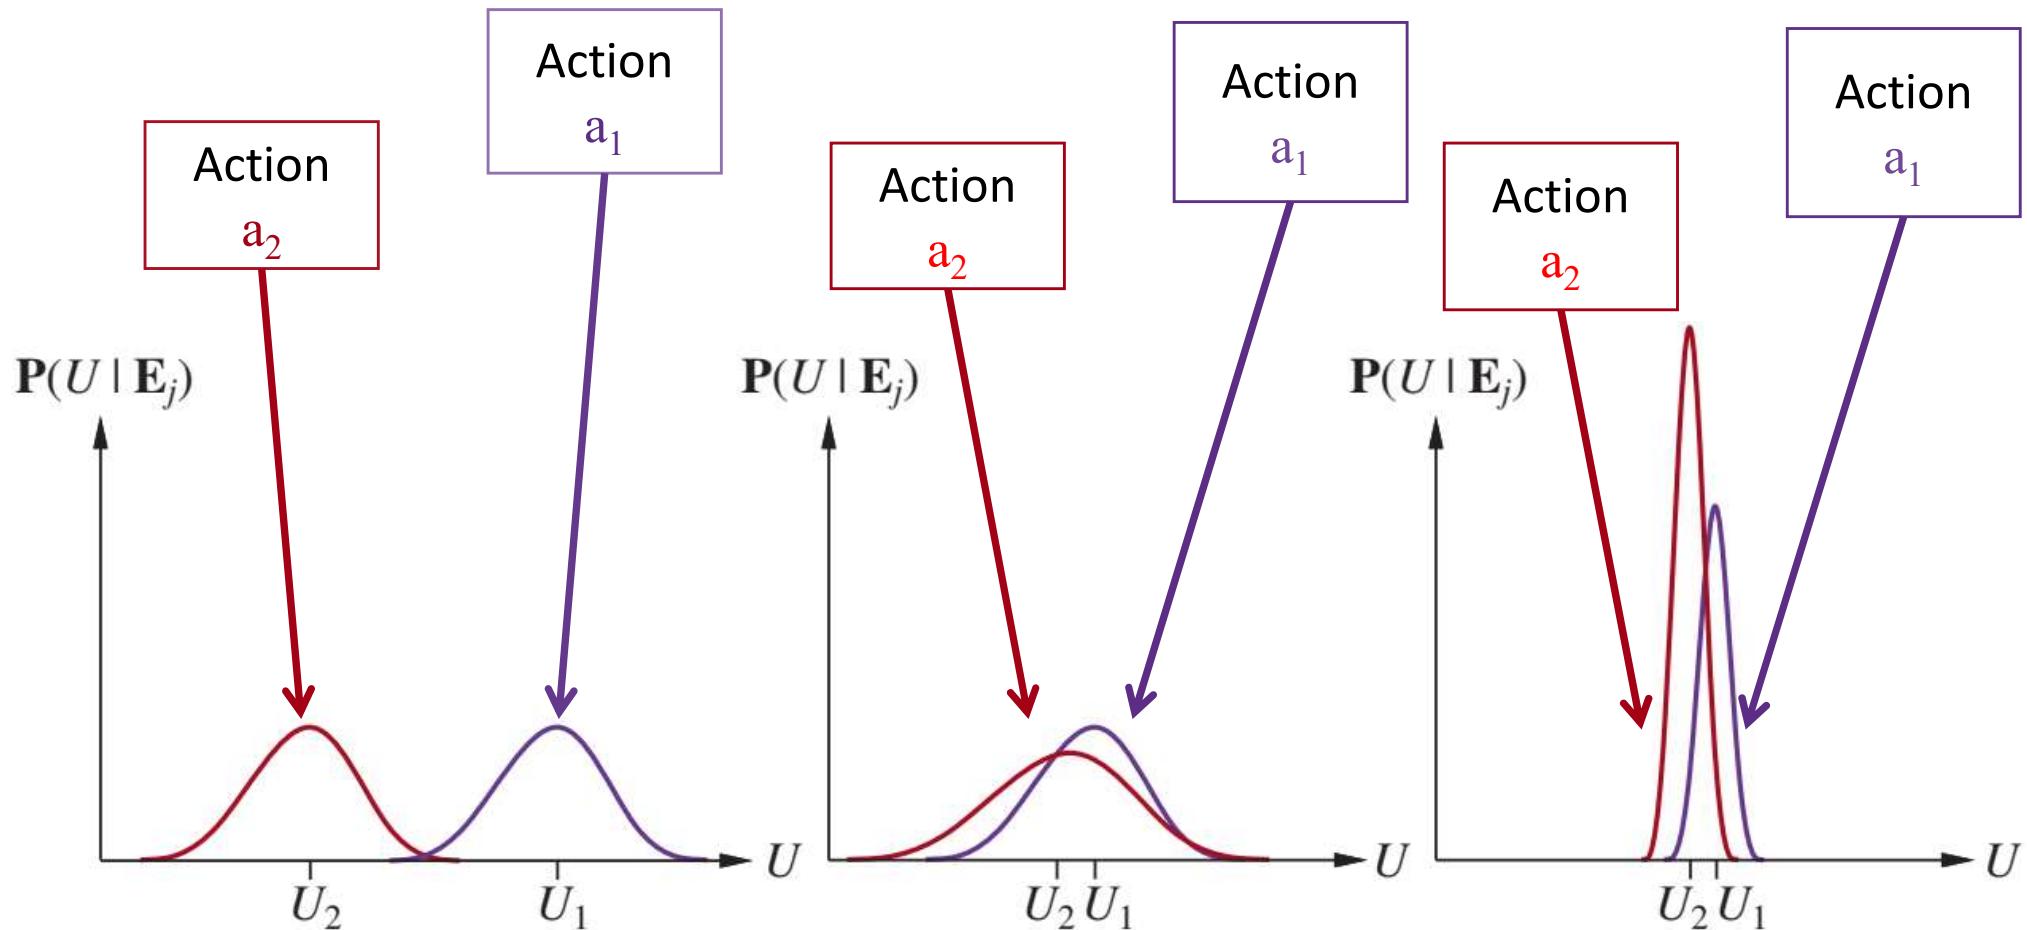
$$EU(\text{take given rain forecast}) = 53$$

## Decision: leave umbrella given sun



$$EU(\text{leave given sun forecast}) = 95$$

# Utility & Value of Perfect Information



New information will not help here.

New information may help a lot here.

New information may help a bit here.

# VPI Properties

Given a decision network with possible observations  $E_j$  (sources of new information / evidence):

- The expected value of information is nonnegative:

$$\forall_j VPI(E_j) \geq 0$$

- VPI is not additive:

$$VPI(E_j, E_k) \neq VPI(E_j) + VPI(E_k)$$

- VPI is order-independent:

$$VPI(E_j, E_k) = VPI(E_j) + VPI(E_k | E_j) = VPI(E_k) + VPI(E_j | E_k) = VPI(E_k, E_j)$$

# Information Gathering Agent

**function** INFORMATION-GATHERING-AGENT(*percept*) **returns** an *action*  
**persistent:**  $D$ , a decision network

integrate *percept* into  $D$

$j \leftarrow$  the value that maximizes  $VPI(E_j) / C(E_j)$

**if**  $VPI(E_j) > C(E_j)$

**then return**  $Request(E_j)$

**else return** the best action from  $D$

# Expected Action Utility

The **expected utility of an action  $a$  given the evidence is the average utility value of all possible outcomes  $s'$  of action  $a$ , weighted by their probability (belief) of occurrence:**

$$EU(a) = \sum_{s'} \sum_s P(s) * P(s' | s, a) * U(s') = \sum_{s'} P(Result(a) = s') * U(s')$$

Rational agent should choose an action that **maximizes the expected utility:**

$$\text{chosen action} = \underset{a}{\operatorname{argmax}} \ EU(a)$$

# How Did We Get Here?

Let's start with relationships (and related notation) between agent's preferences:

- agent **prefers** A over B:

$$A > B$$

- agent is **indifferent** between A and B:

$$A \sim B$$

- agent prefers A over B or is indifferent between A and B (**weak preference**):

$$A \geqslant B$$

# The Concept of Lottery

Let's assume the following:

- an **action a** is a lottery ticket
- the **set of outcomes (resulting states)** is a lottery

A lottery L with possible outcomes  $S_1, \dots, S_n$  that occur with probabilities  $p_1, \dots, p_n$  is written as:

$$L = [p_1, S_1; p_2, S_2; \dots; p_n, S_n]$$

**Lottery outcome  $S_i$ :** atomic state or another lottery.

# Lottery Constraints: Orderability

Given two lotteries A and B, a rational agent must either prefer one or else rate them as equally preferable:

Exactly one of  $(A \succ B)$ ,  $(B \succ A)$ , or  $(A \sim B)$  holds

# Lottery Constraints: Transitivity

**Given three lotteries A, B, and C, if an agent prefers A to B AND prefers B to C, then the agent must prefer A to C:**

$$(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$$

# Lottery Constraints: Continuity

If some lottery B is between A and C in preference, then there is some probability p for which the rational agent will be indifferent between getting B for sure or some other lottery that yields A with probability p and C with probability 1 - p:

$$(A \succ B \succ C) \Rightarrow \exists p [p, A; 1-p, C] \sim B$$

# Lottery Constraints: Substitutability

If an agent is indifferent between two lotteries A and B, then the agent is indifferent between two more complex lotteries that are the same, except that B is substituted for A in one of them:

$$(A \sim B) \Rightarrow [p, A; 1-p, C] \sim [p, B; 1-p, C]$$

# Lottery Constraints: Monotonicity

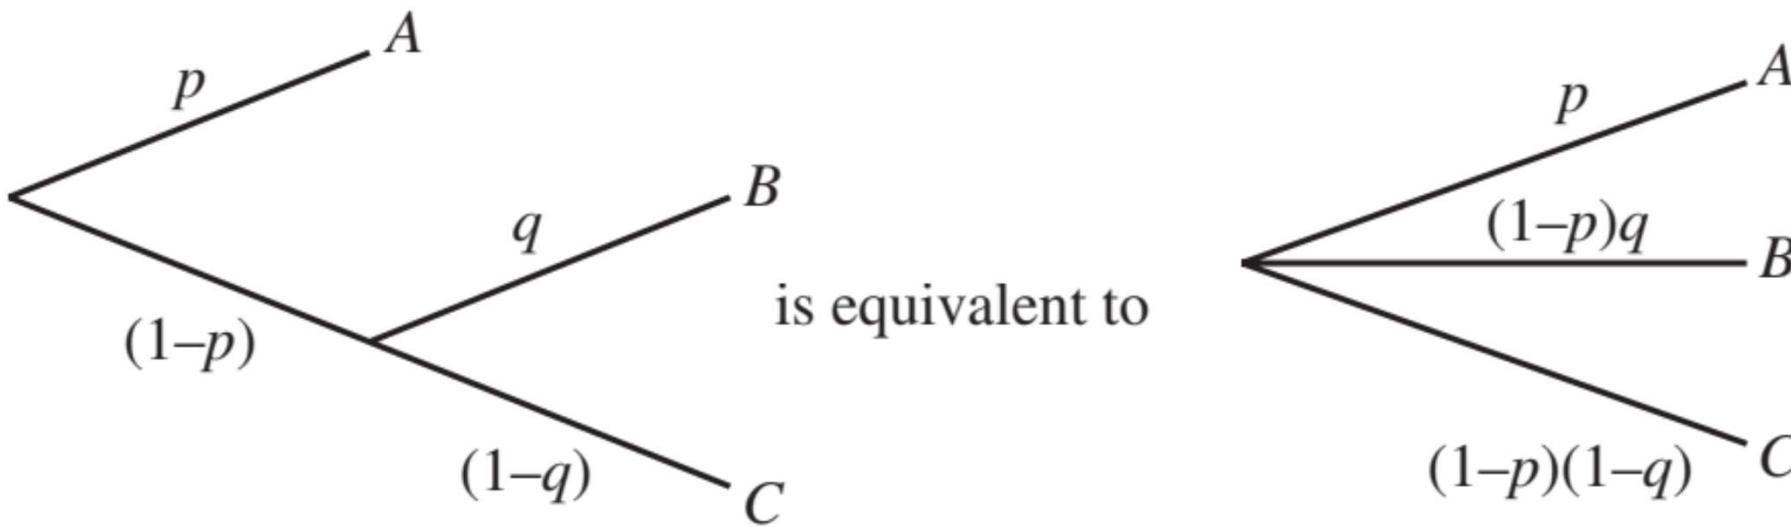
Suppose two lotteries have the same two possible outcomes, A and B. If an agent prefers A to B, then the agent must prefer the lottery that has a higher probability for A:

$$(A \succ B) \Rightarrow (p > q \Leftrightarrow [p, A; 1-p, B] \succ [q, A; 1-q, B])$$

# Lottery Constraints: Decomposability

Compound lotteries can be reduced to smaller ones using the laws of probability:

$$[p, A; 1-p, [q, B; 1-q, C]] \sim [p, A; (1-p)*q, B; (1-p)*(1-q), C]$$



# Preferences and Utility Function

An agent whose preferences between lotteries follow the set of axioms (**of utility theory**) below:

- Orderability
- Transitivity
- Continuity
- Substitutability
- Monotonicity
- Decomposability

can be described as possesing a utility function and maximize it.

# Preferences and Utility Function

If an agent's preferences obey the axioms of utility theory, then there exist a function  $U$  such that:

$U(A) = U(B)$  if and only if  $(A \sim B)$

and

$U(A) > U(B)$  if and only if  $(A > B)$

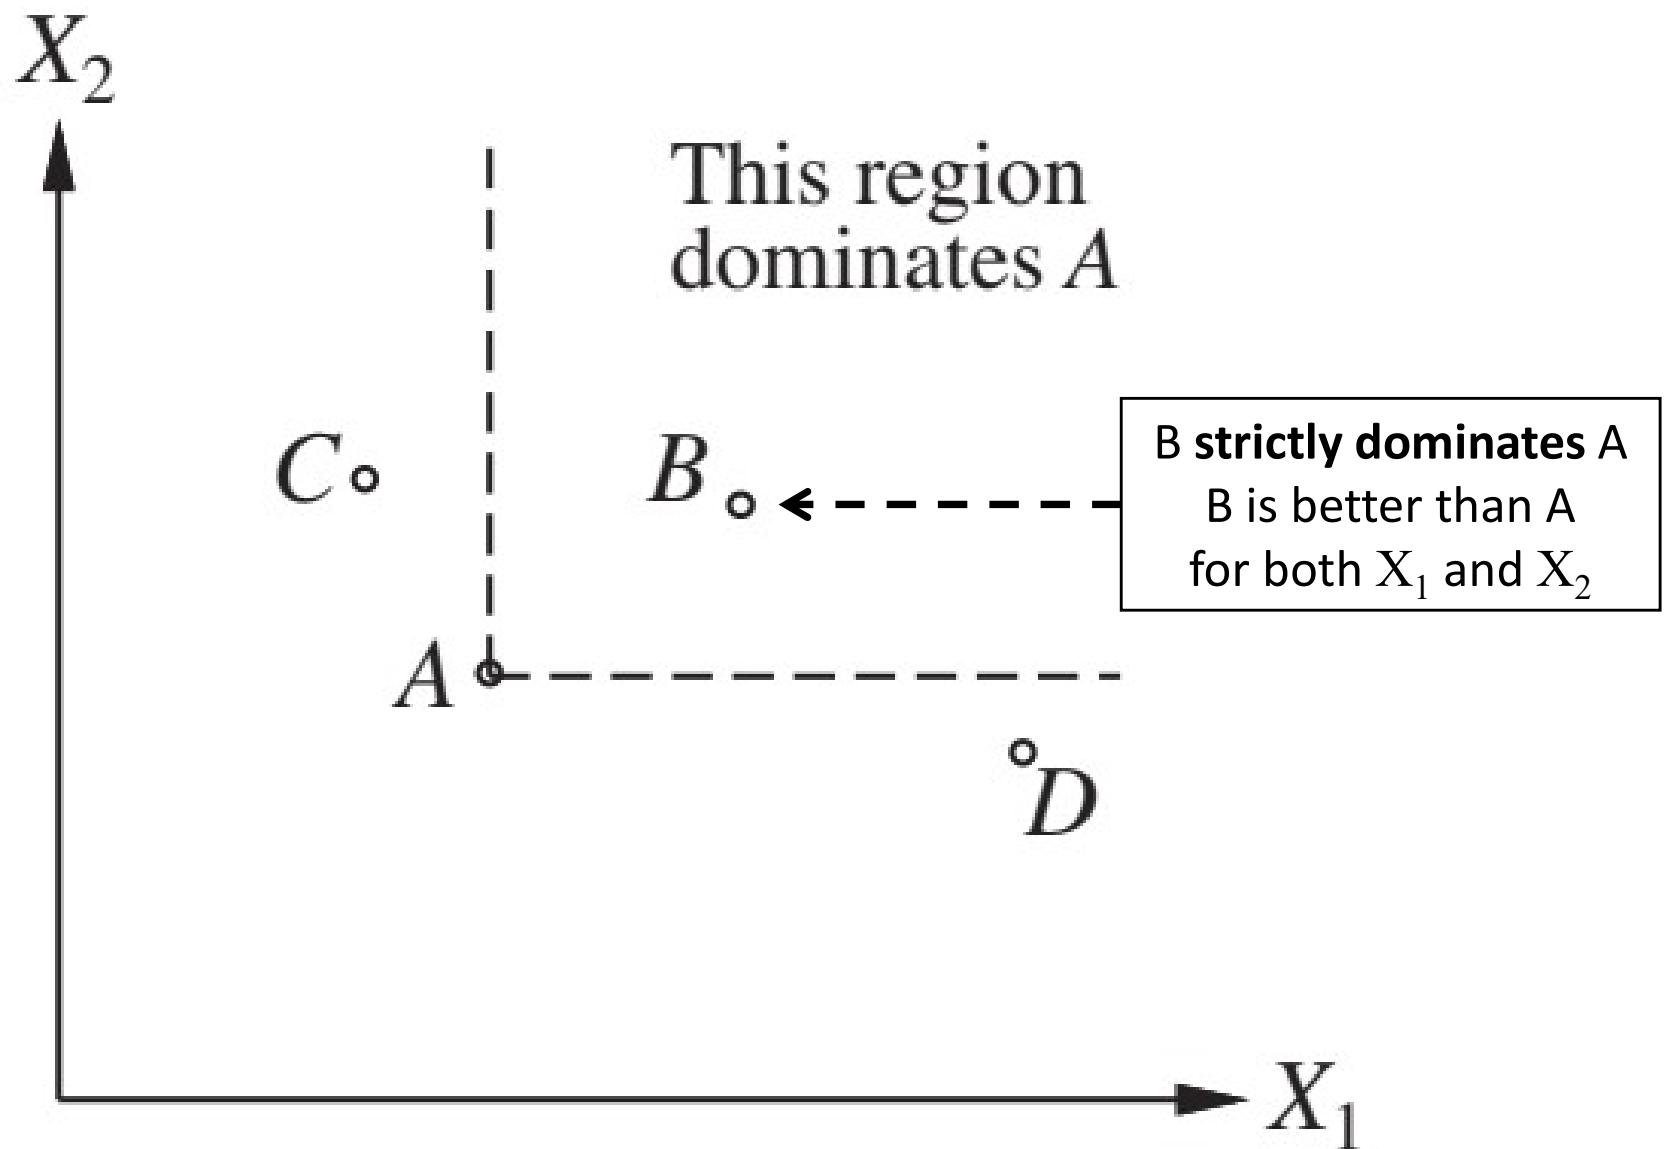
# Multiatribute Outcomes

**Outcomes can be characterized by more than one attribute. Decisions in such cases are handled by Multiatribute Utility Theory.**

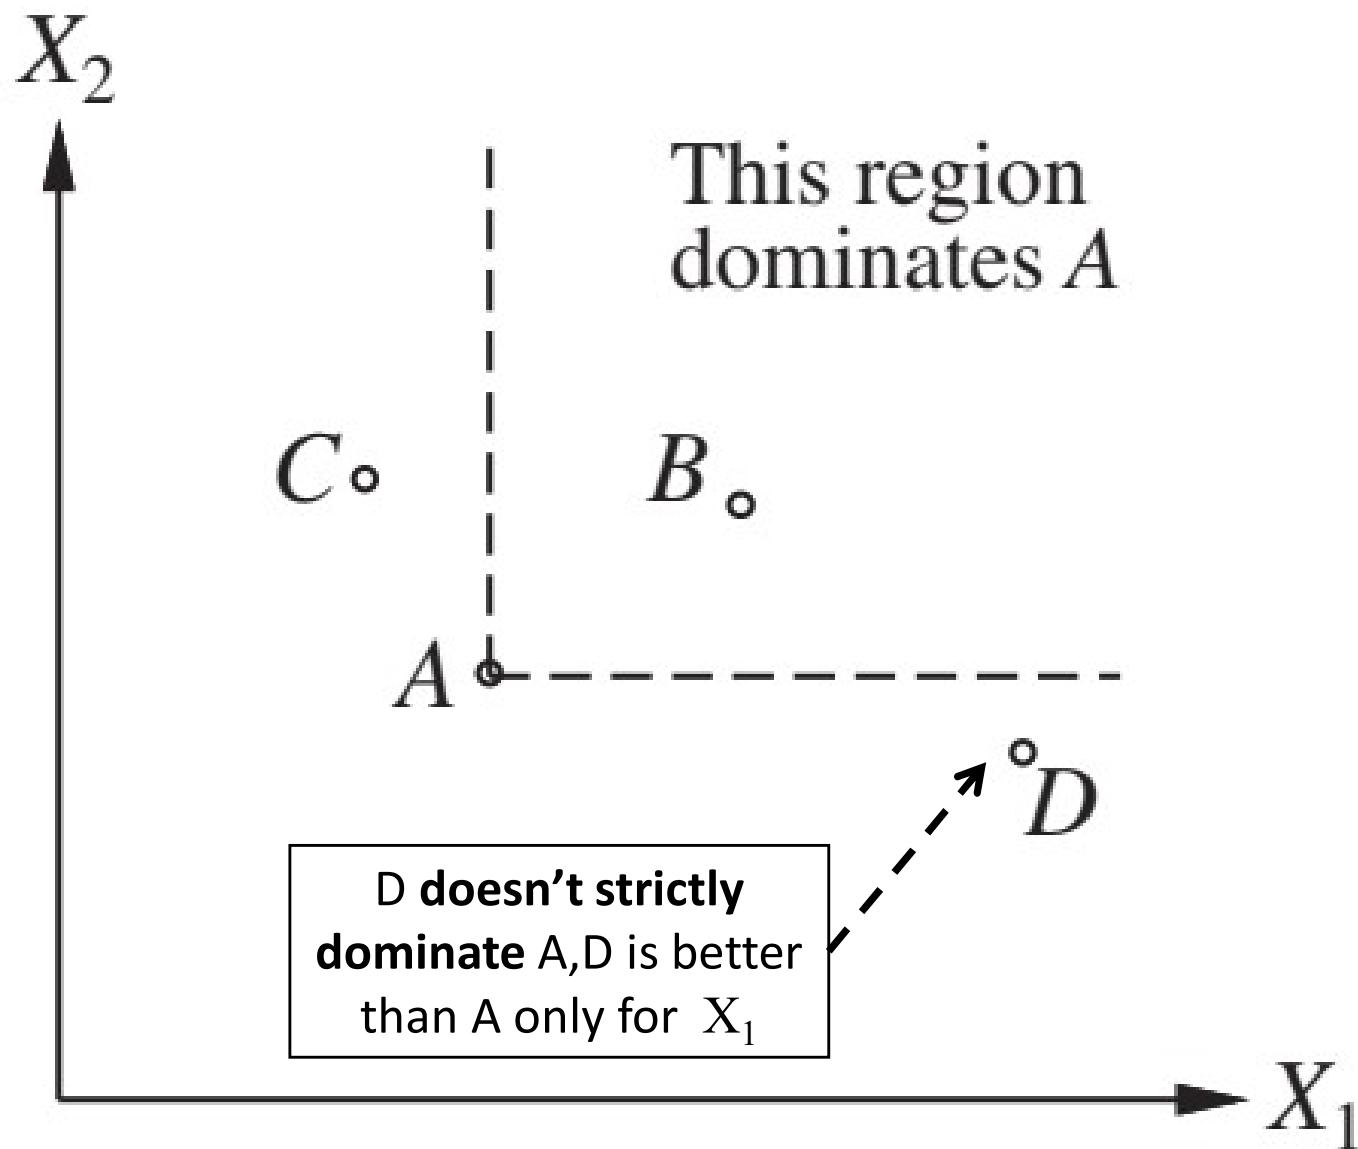
**Attributes:**  $X = X_1, \dots, X_n$

**Assigned values:**  $x = \langle x_1, \dots, x_n \rangle$

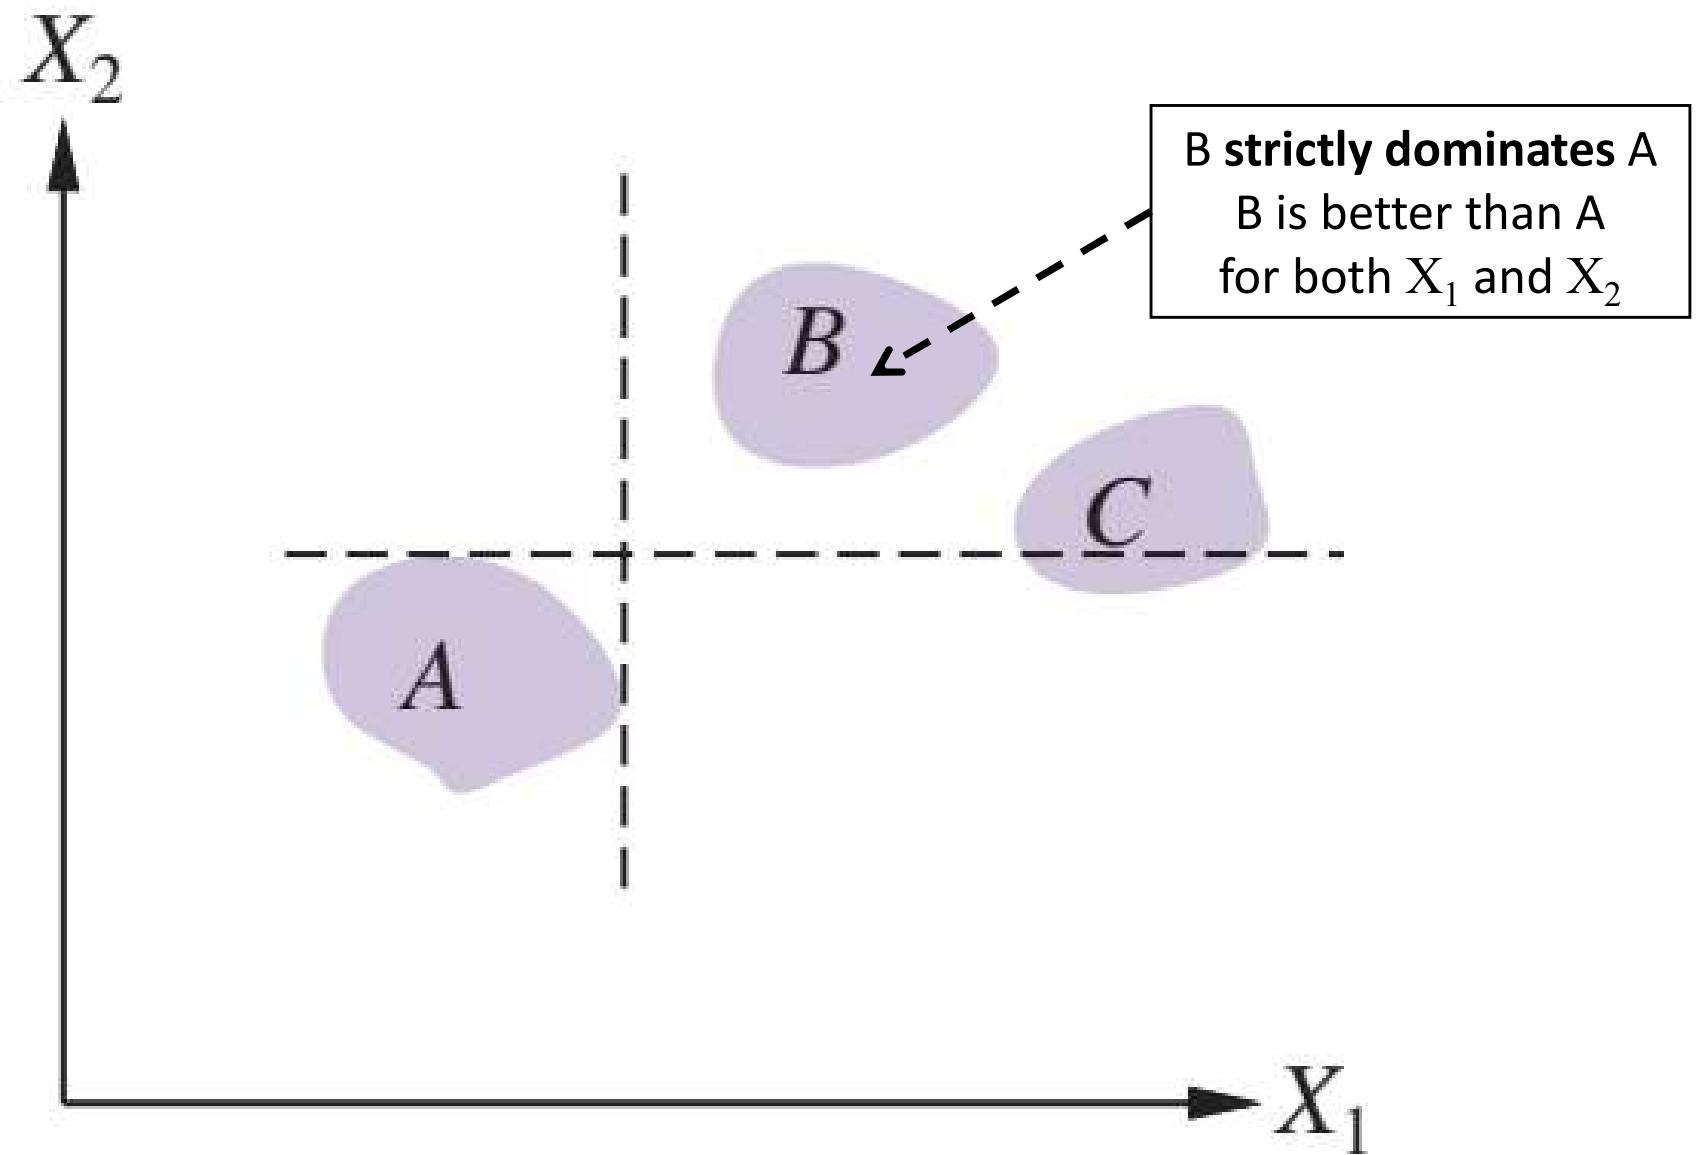
# Strict Dominance: Deterministic



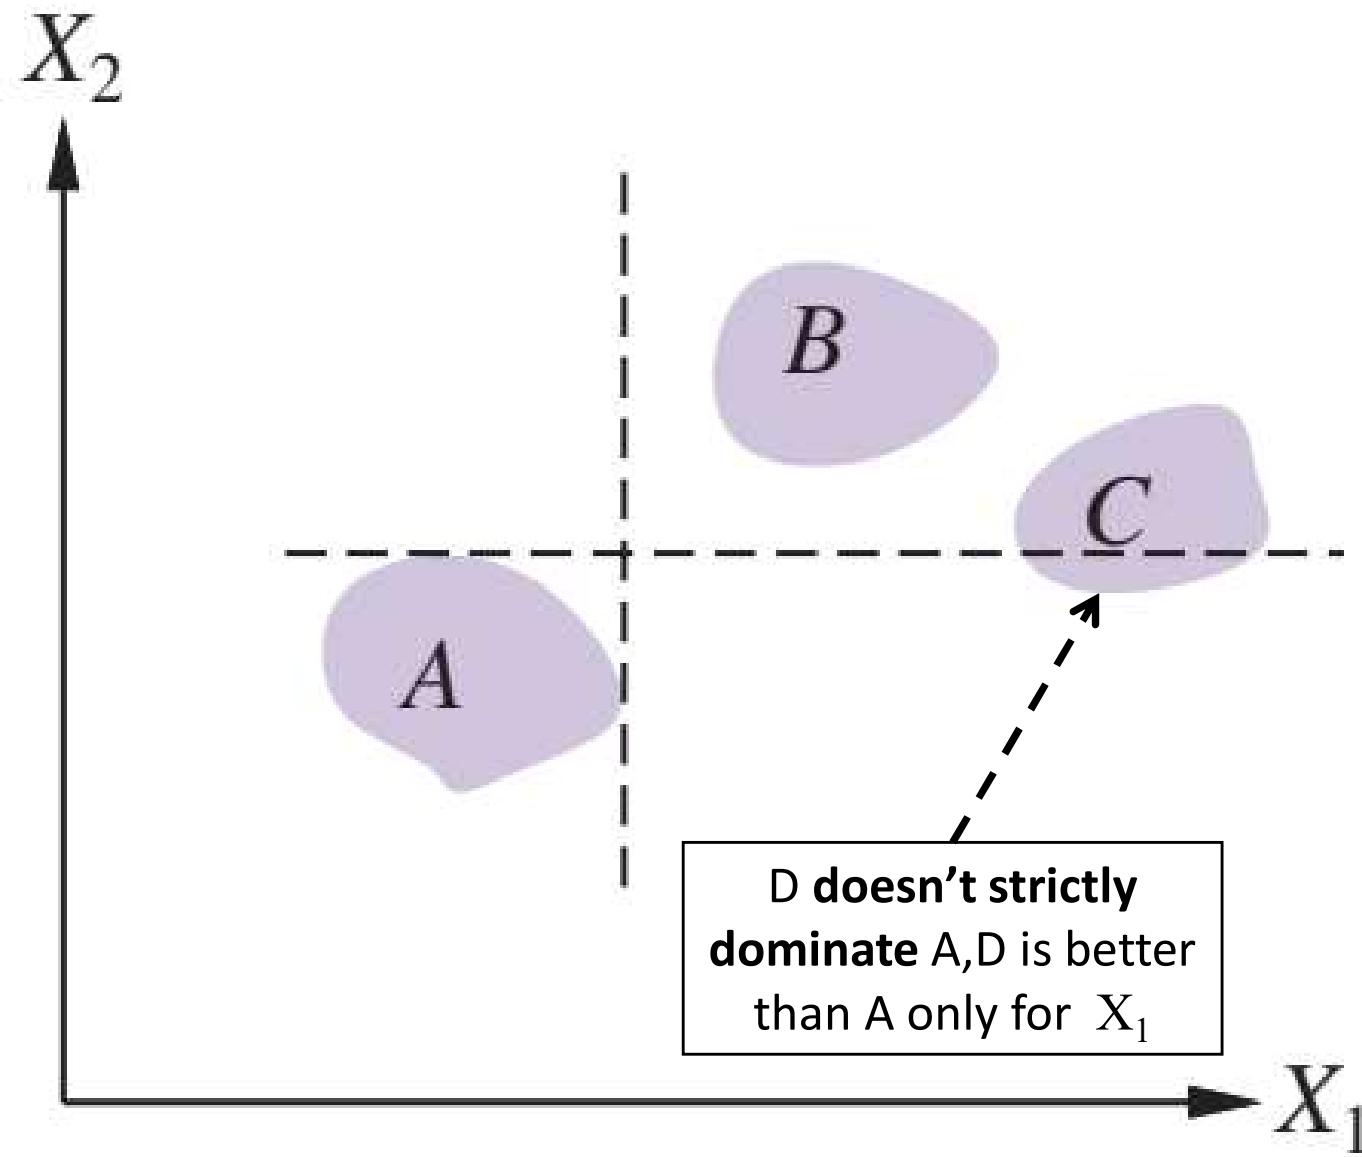
# Strict Dominance: Deterministic



# Strict Dominance: Uncertain



# Strict Dominance: Uncertain



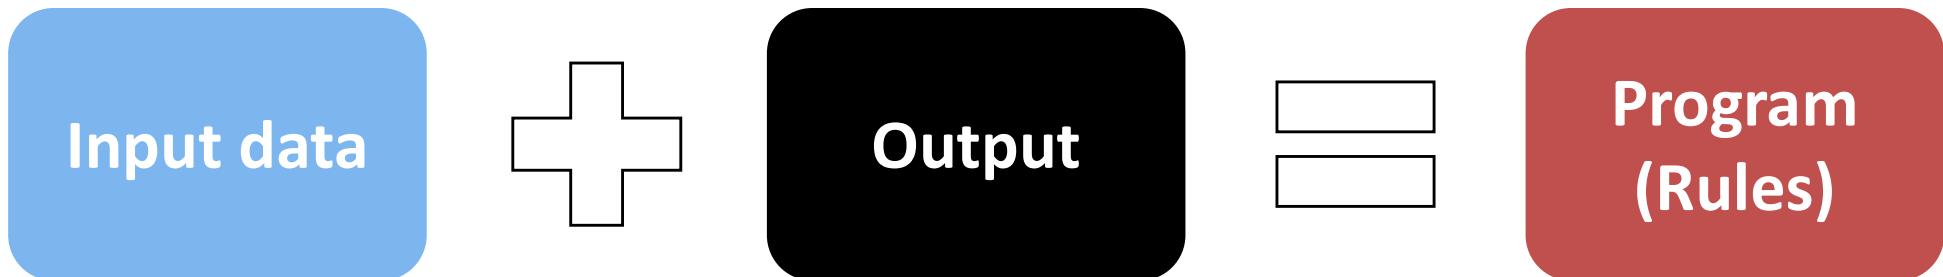
# What is Machine Learning?

# Traditional Programming vs ML

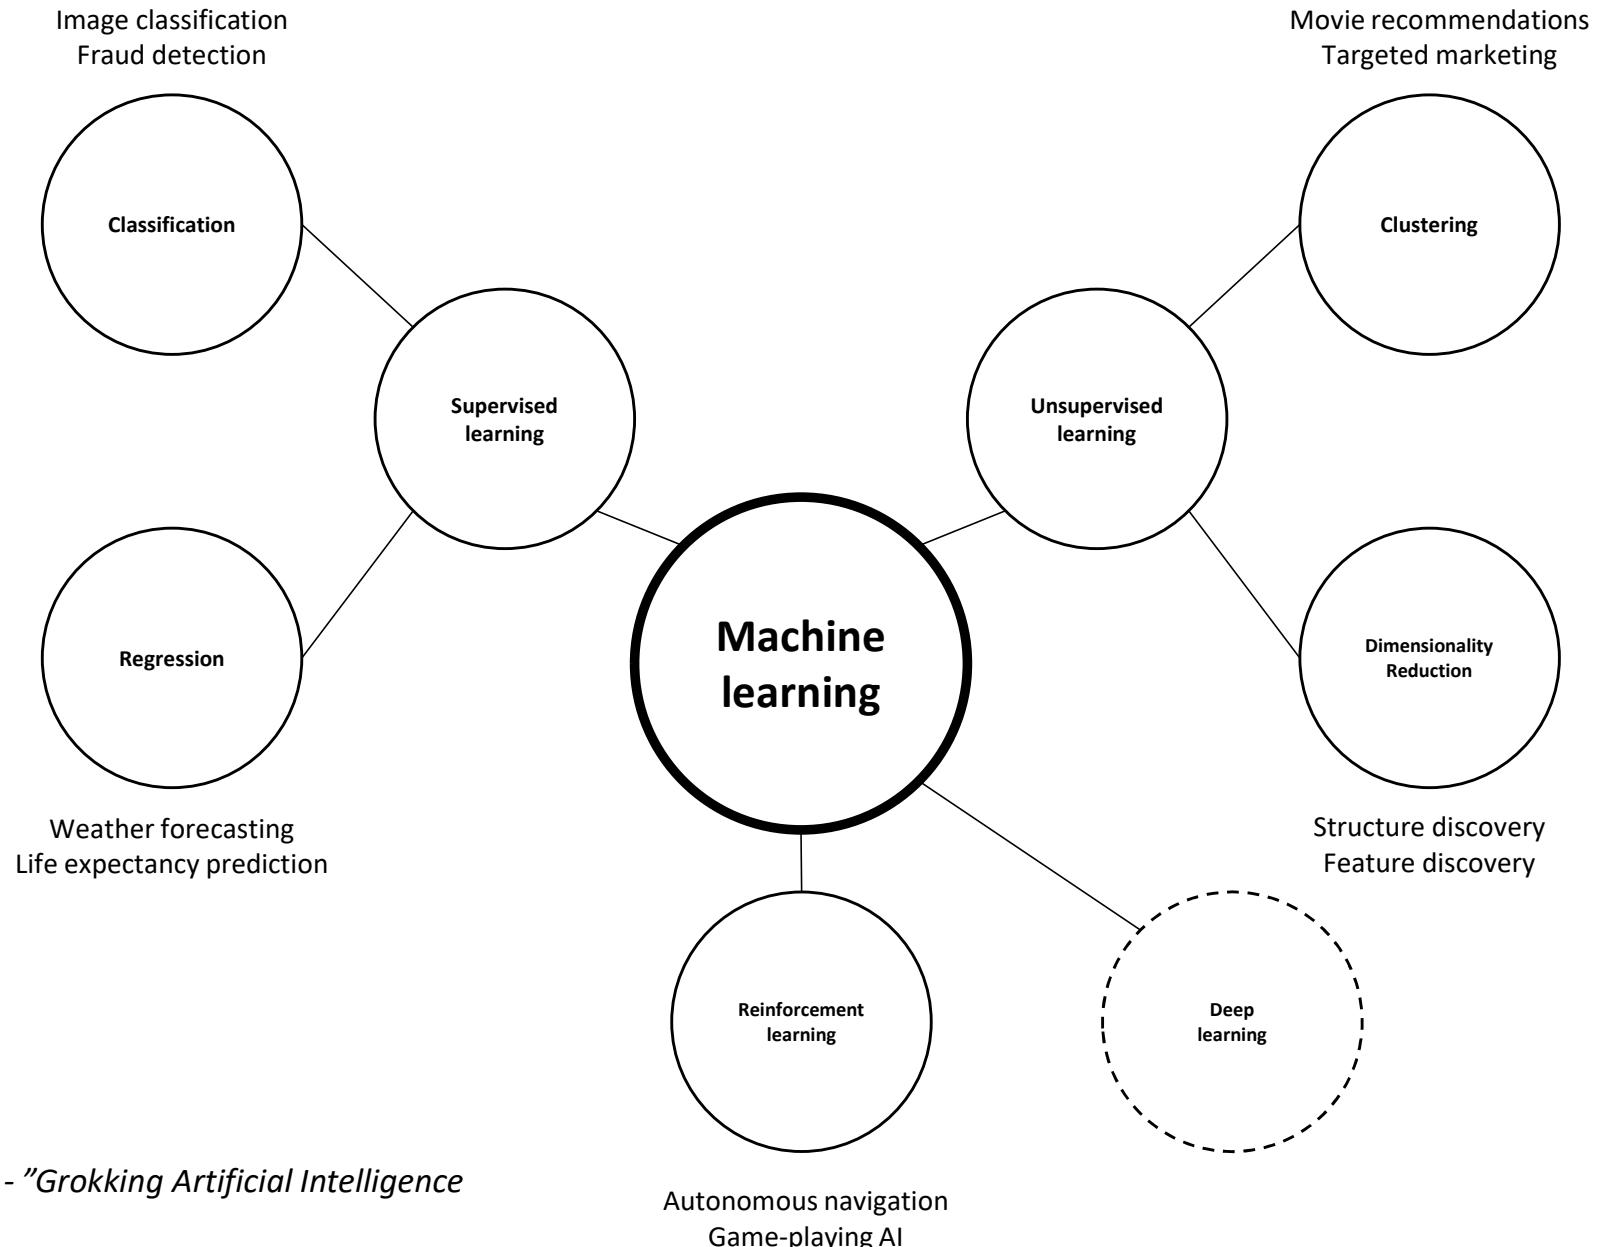
Traditional programming:



Machine learning:



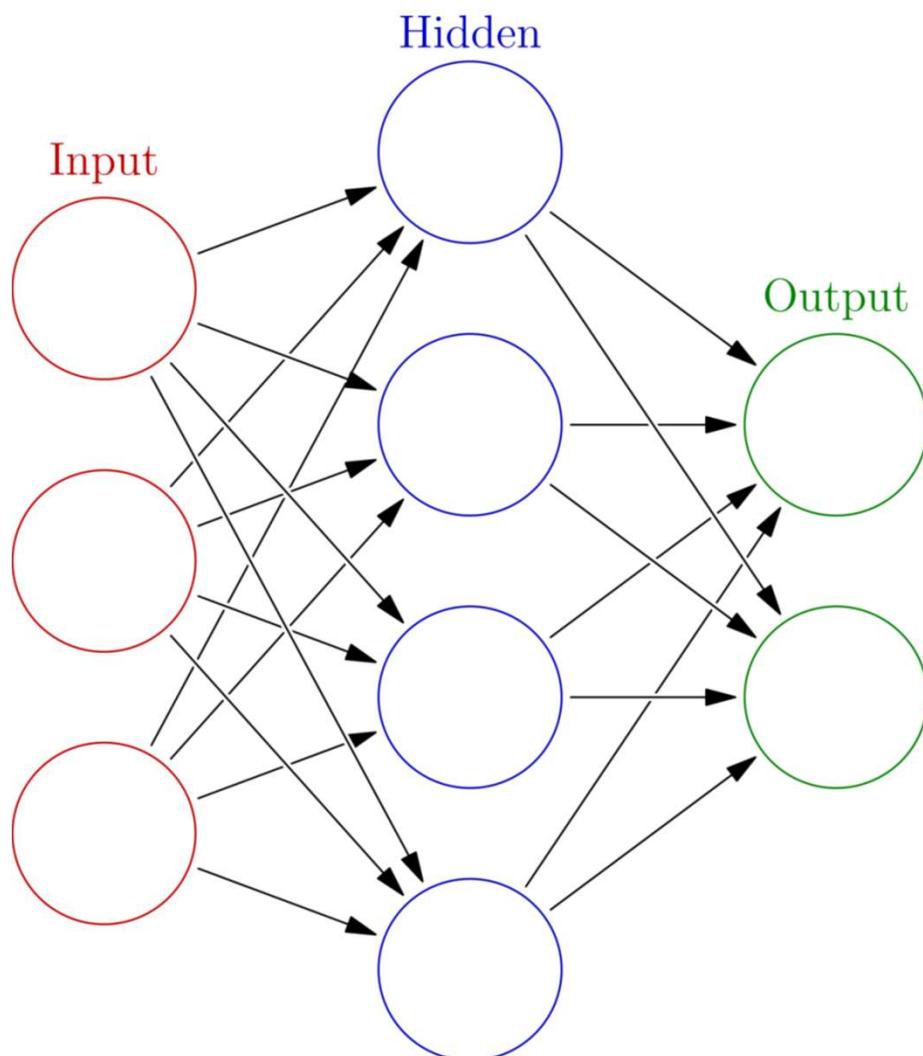
# Machine Learning Categories



# What Kind of Questions ML Answers?

Question	ML Category	Example
Is this A or B?	Classification	Will this car fail in the next two months? Yes or no?
Is this weird?	Anomaly detection	Is this credit card charge normal?
How much / many?	Regression	What will the temperature be tomorrow?
How is this organized?	Clustering	Which car models have the most brake problems?
What should I do next?	Reinforcement learning	Adjust room humidity or leave as is?

# Artificial Neural Network (1943)



First computational models of an **Artificial Neural Network** (loosely inspired by biological neural networks) were proposed by Warren McCulloch and Walter Pitts in 1943. Their ideas are a **key component of modern day machine and deep learning**.

Source:

[https://en.wikipedia.org/wiki/Artificial\\_neural\\_network](https://en.wikipedia.org/wiki/Artificial_neural_network)

# Main Machine Learning Categories

## Supervised learning

**Supervised learning** is one of the most common techniques in machine learning. It is based on **known relationship(s) and patterns within data** (for example: relationship between inputs and outputs).

Frequently used types:  
**regression**, and  
**classification**.

## Unsupervised learning

**Unsupervised learning** involves finding underlying patterns within data. Typically used in **clustering** data points (similar customers, etc.)

## Reinforcement learning

Reinforcement learning is inspired by behavioral psychology. It is **based on a rewarding / punishing an algorithm**.

Rewards and punishments are based on algorithm's action within its environment.

# Choosing Hypothesis / Model

Given a **training set** of  $N$  example input-output (feature-label) pairs

$$(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)$$

where each pair was generated by

$$y = f(x)$$

Ideally, we would like our **model**  $h(x)$  (**hypothesis**) that approximates the true function  $f(x)$  to be:

$$h(x) = y = f(x) \text{ (consistent hypothesis)}$$

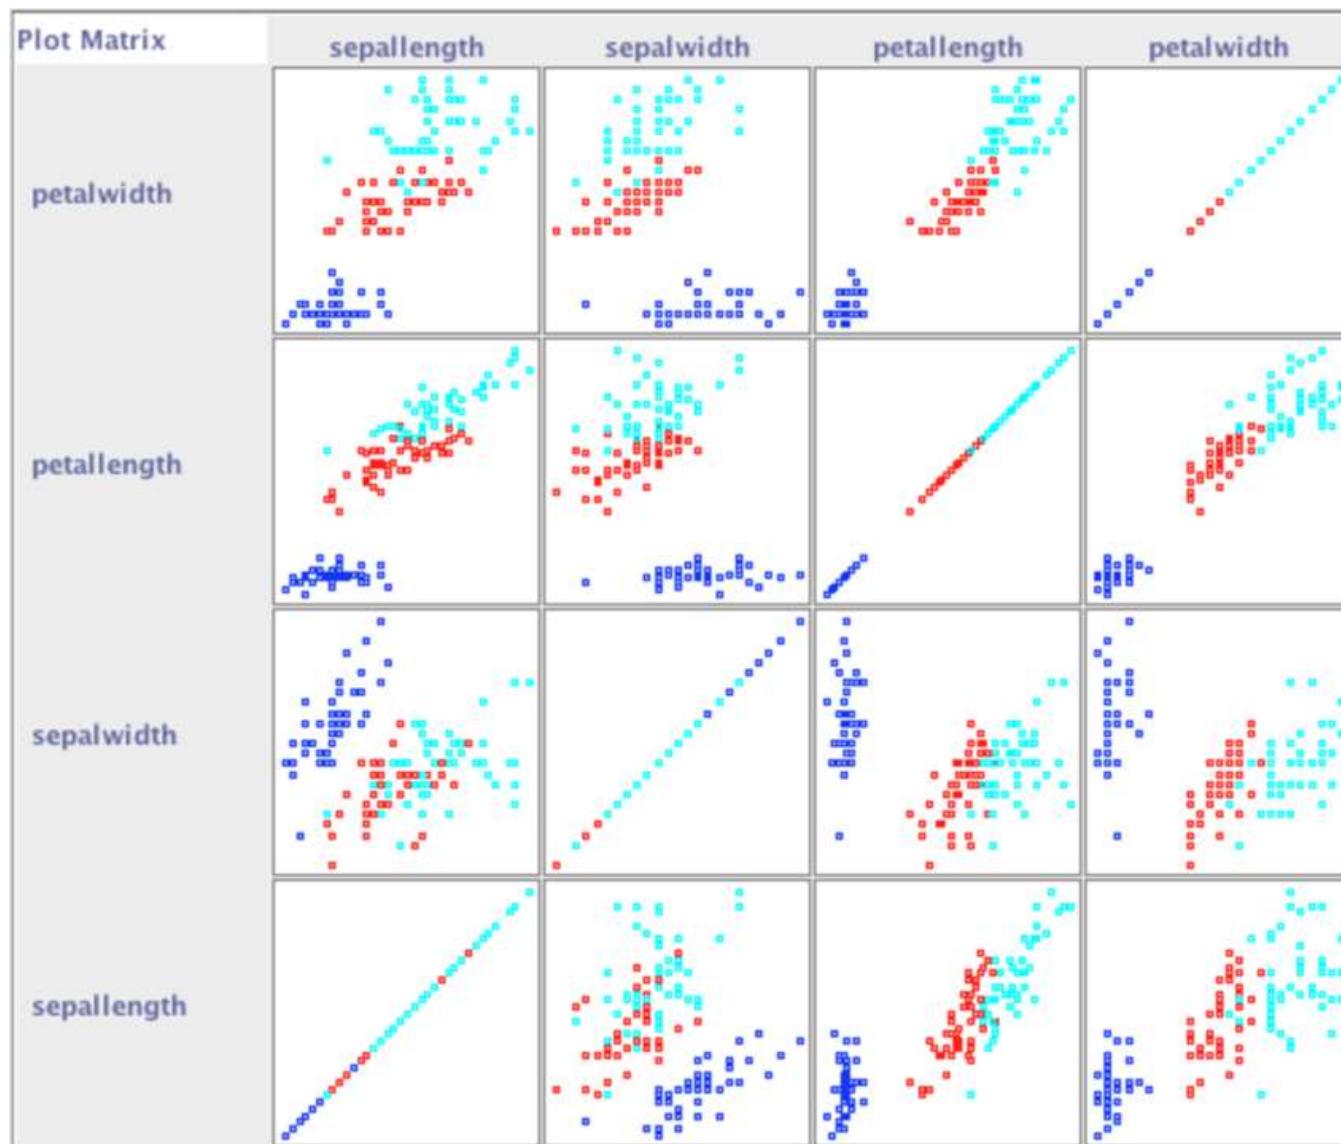
# Choosing Hypothesis / Model

Typically consistent hypothesis is impossible or difficult to achieve:

- use best-fit model / hypothesis

Our model needs to be tested on the test set inputs (data the model has not “seen” yet) to see how well it generalizes (**how accurately it predicts the outputs of the test set**).

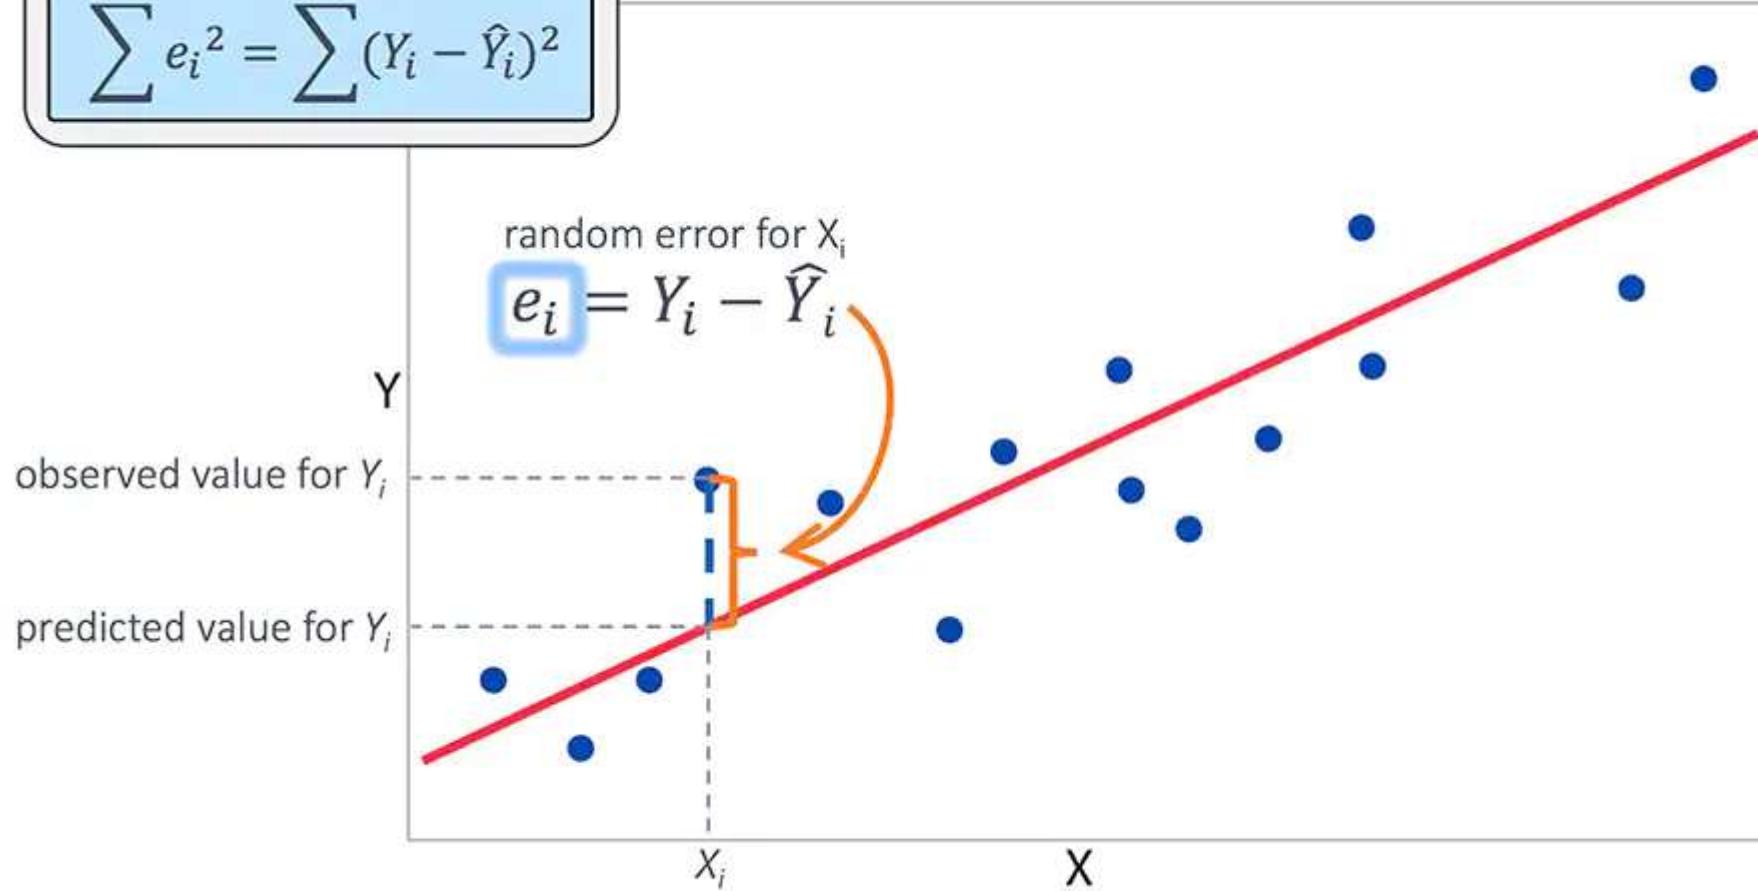
# Exploratory Data Analysis



# Linear Regression Using Least-Squares

Method of Least Squares

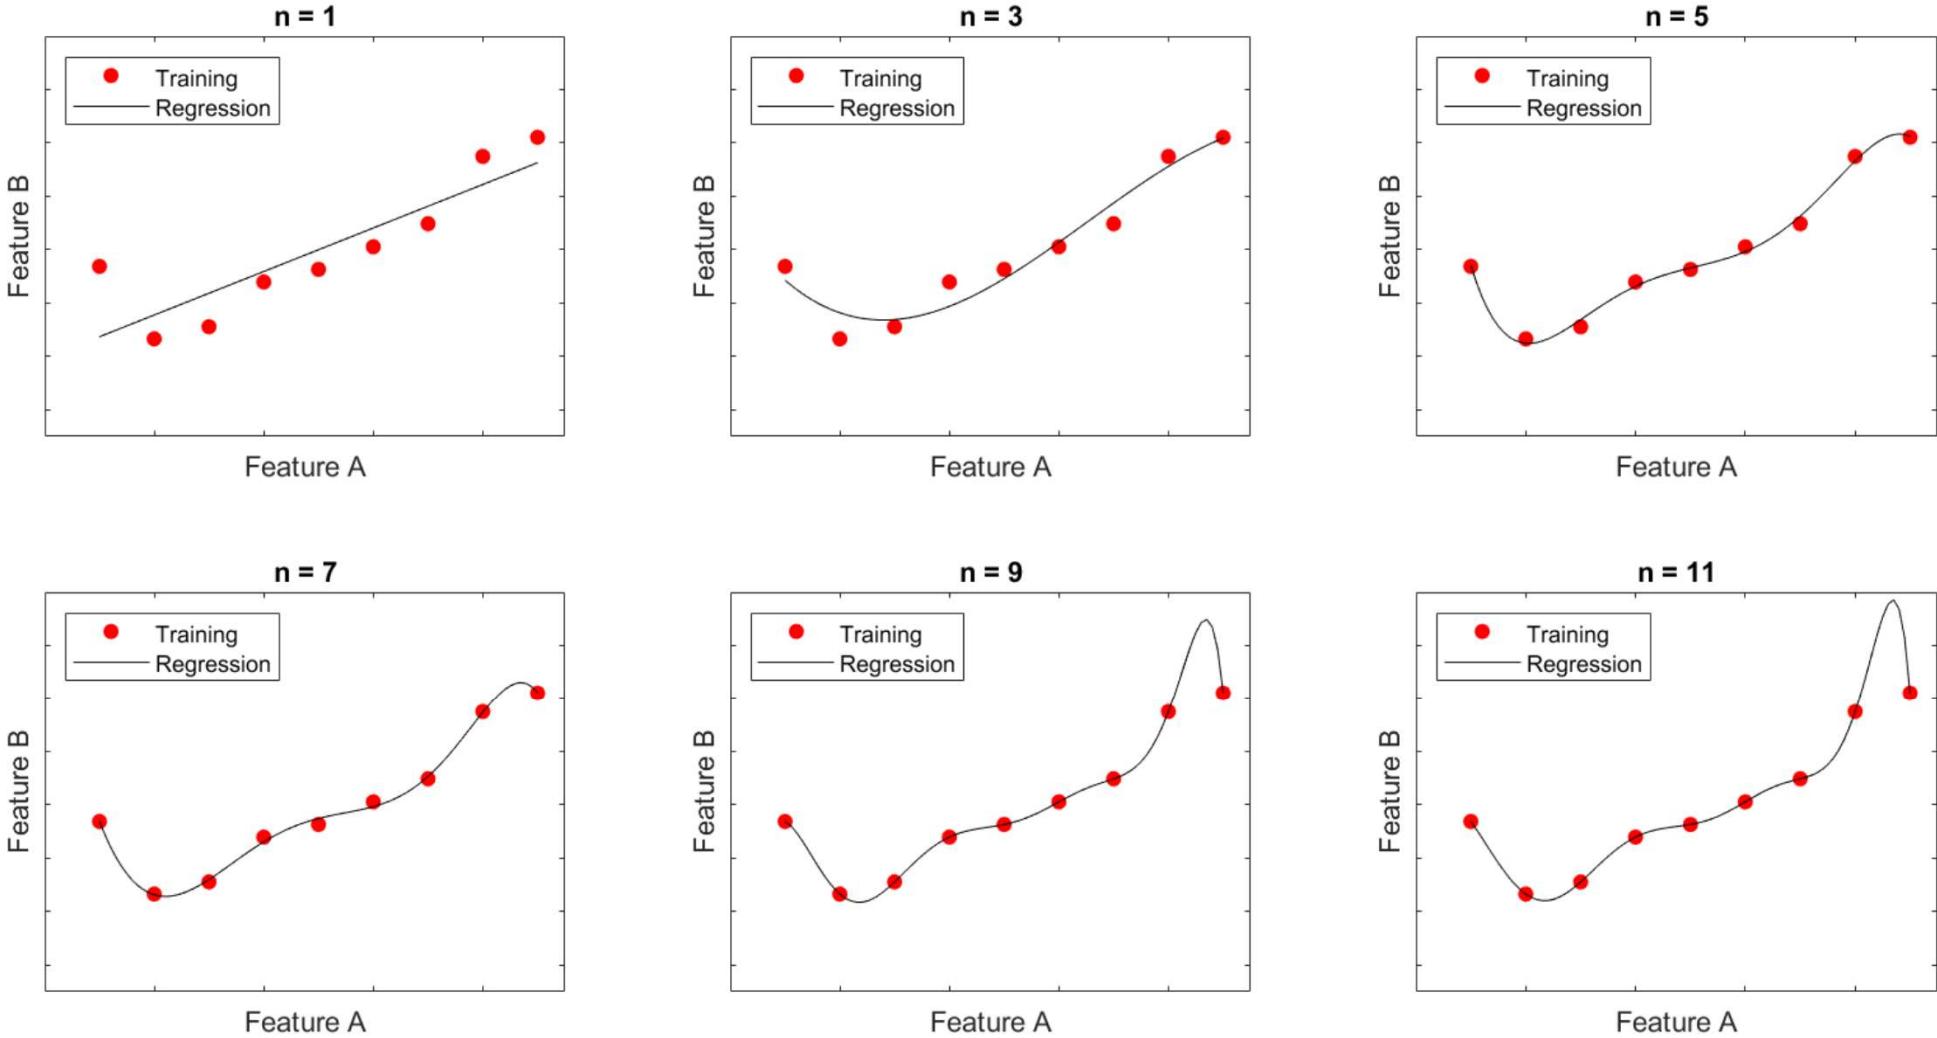
$$\sum e_i^2 = \sum (Y_i - \hat{Y}_i)^2$$



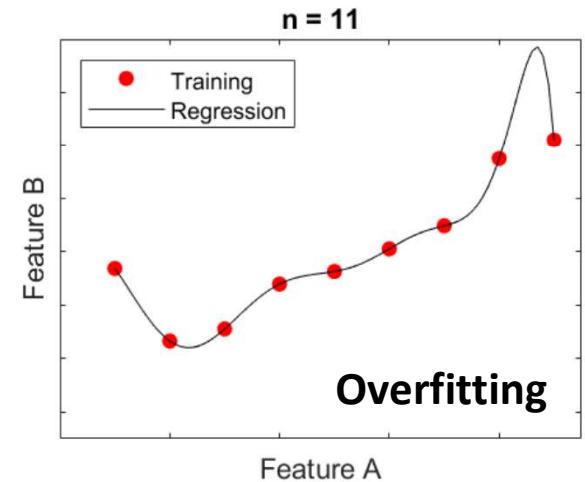
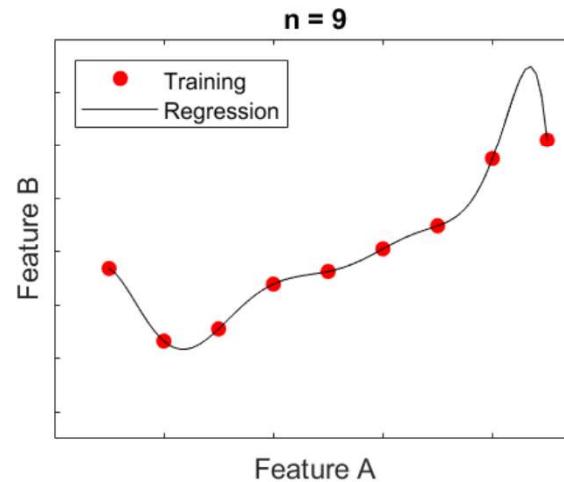
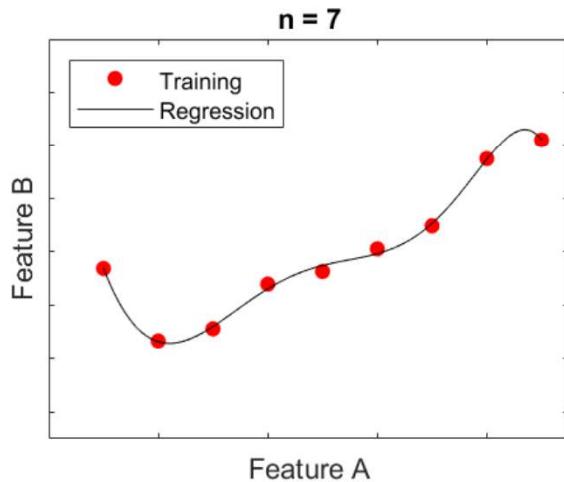
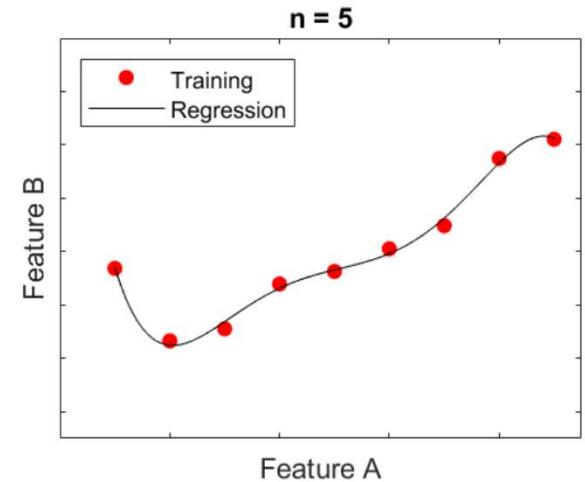
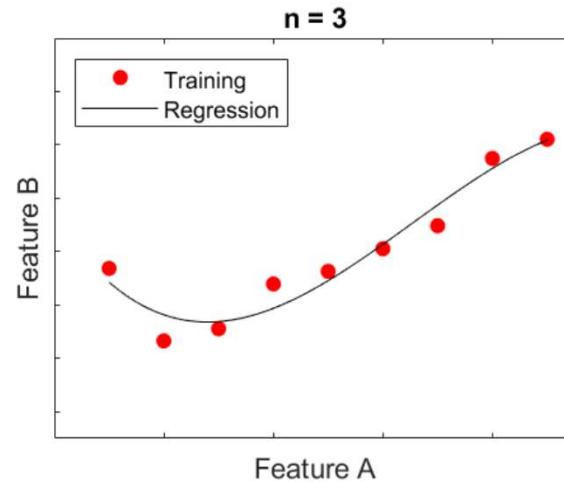
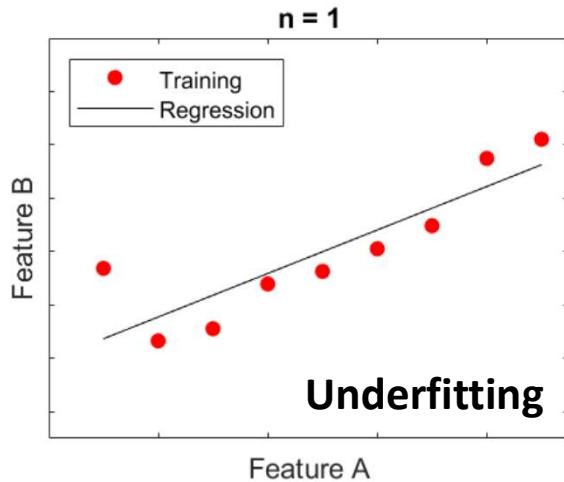
The goal is to find the line  $y = ax + b$  that minimizes the amount of error.

Source: [https://www.jmp.com/en\\_us/statistics-knowledge-portal/what-is-multiple-regression/fitting-multiple-regression-model.html](https://www.jmp.com/en_us/statistics-knowledge-portal/what-is-multiple-regression/fitting-multiple-regression-model.html)

# Regression

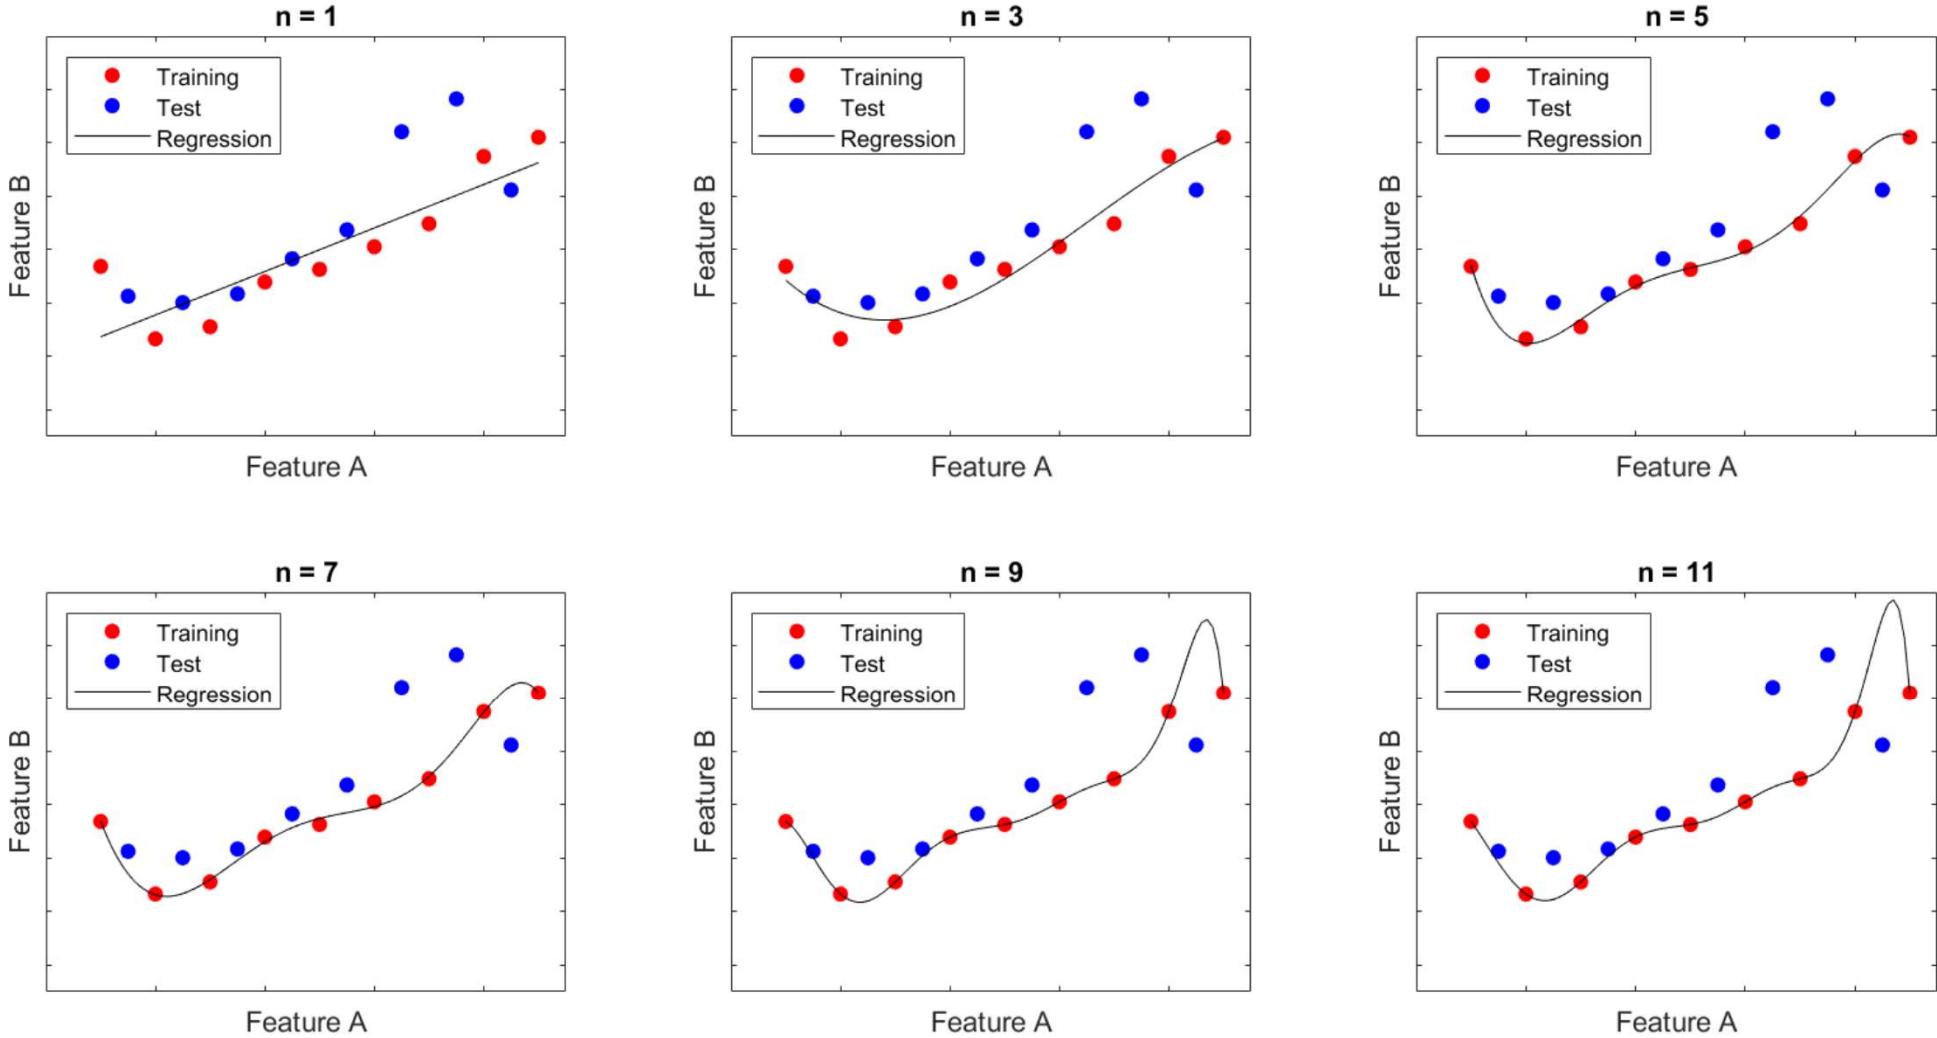


# Regression: Underfitting / Overfitting

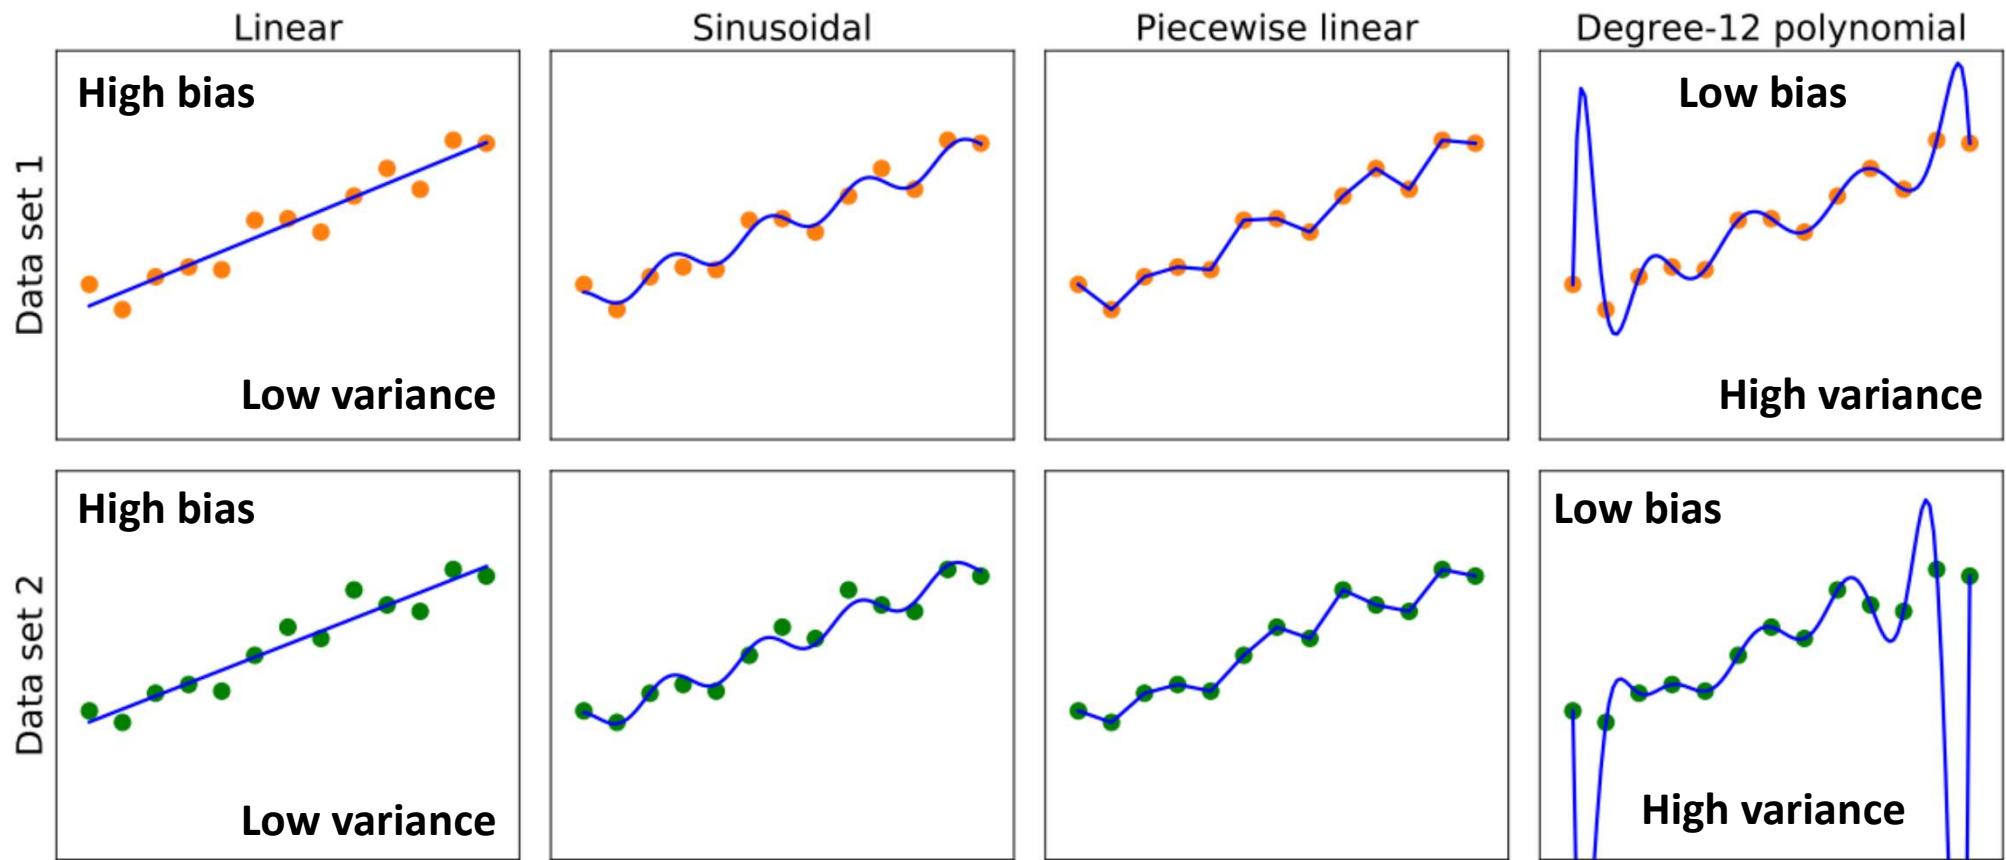


**Underfitting:** “failing” to find pattern in the data.

# Regression



# Bias vs. Variance



**Bias:** the tendency of a predictive hypothesis to deviate from the expected value when averaged over different training sets

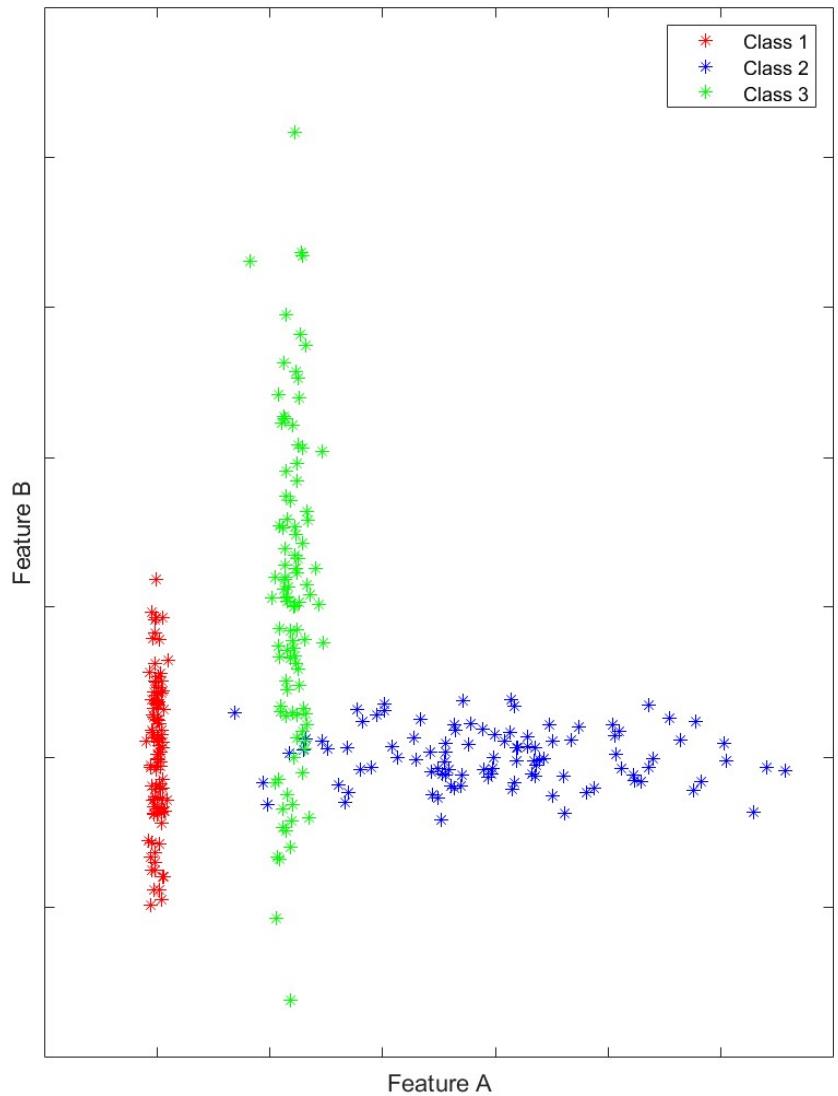
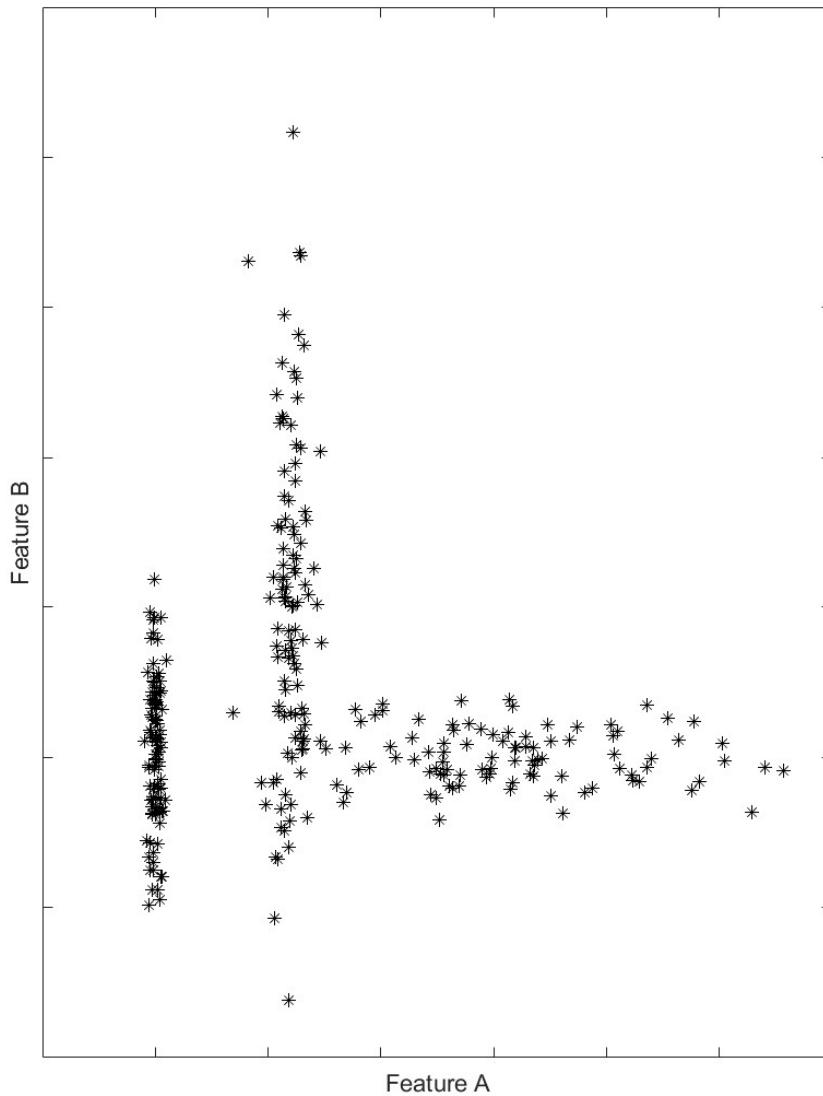
**Variance:** the amount of change in the hypothesis (model) due to fluctuations in training data

# Training / Validation / Test Sets

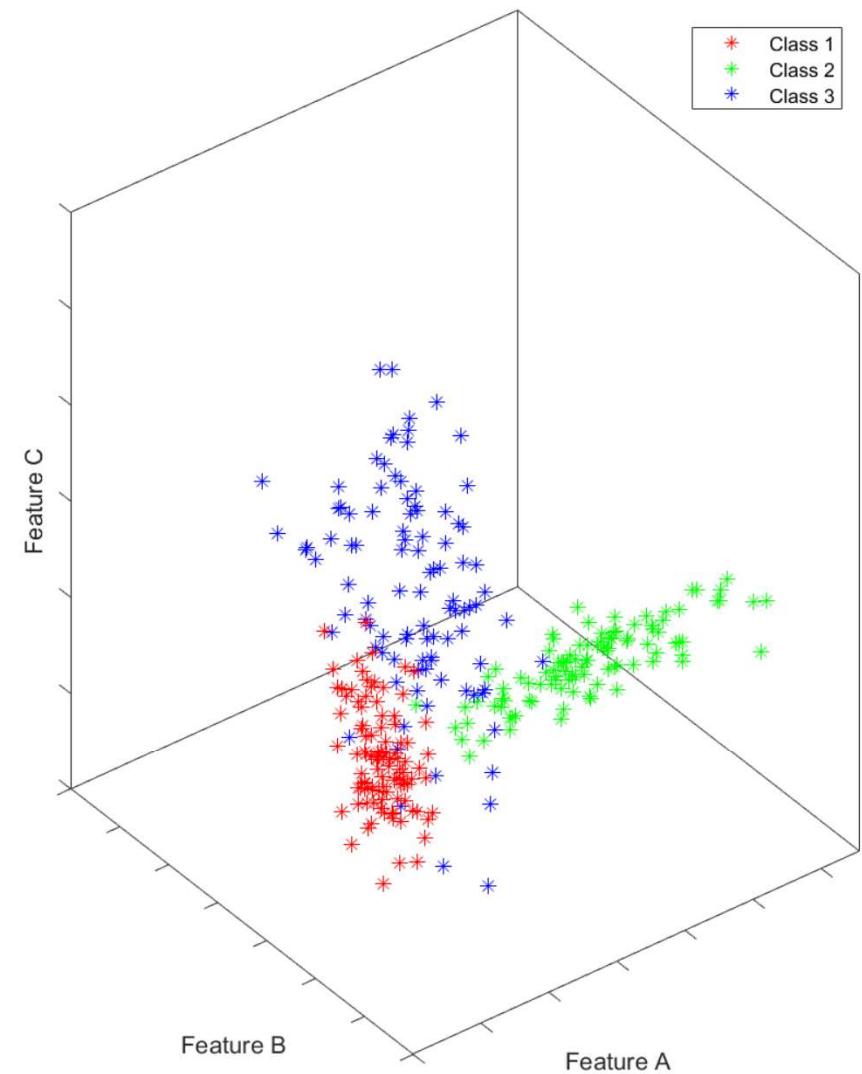
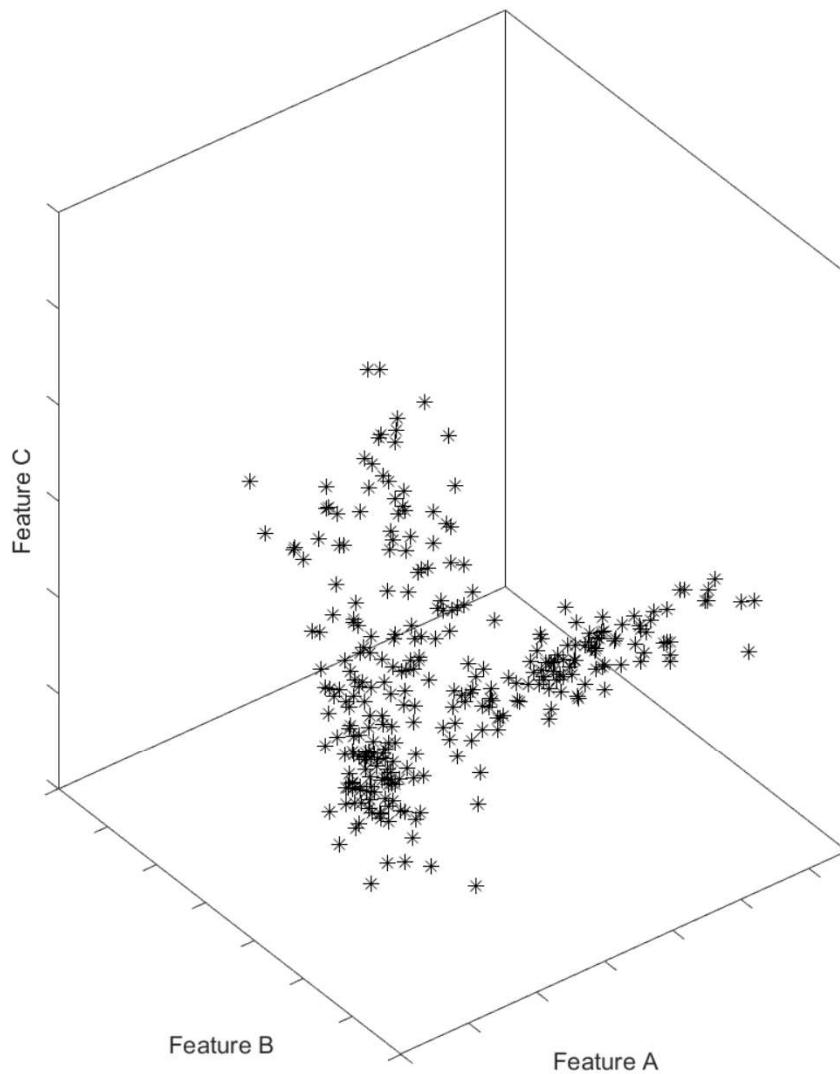
In order to create the best model possible, given some (relatively large) data set, we should divide it into:

- **training** set: to train candidate models
- **validation** set: to evaluate candidate models and pick the best one
- **test** set: to do the final evaluation of the model

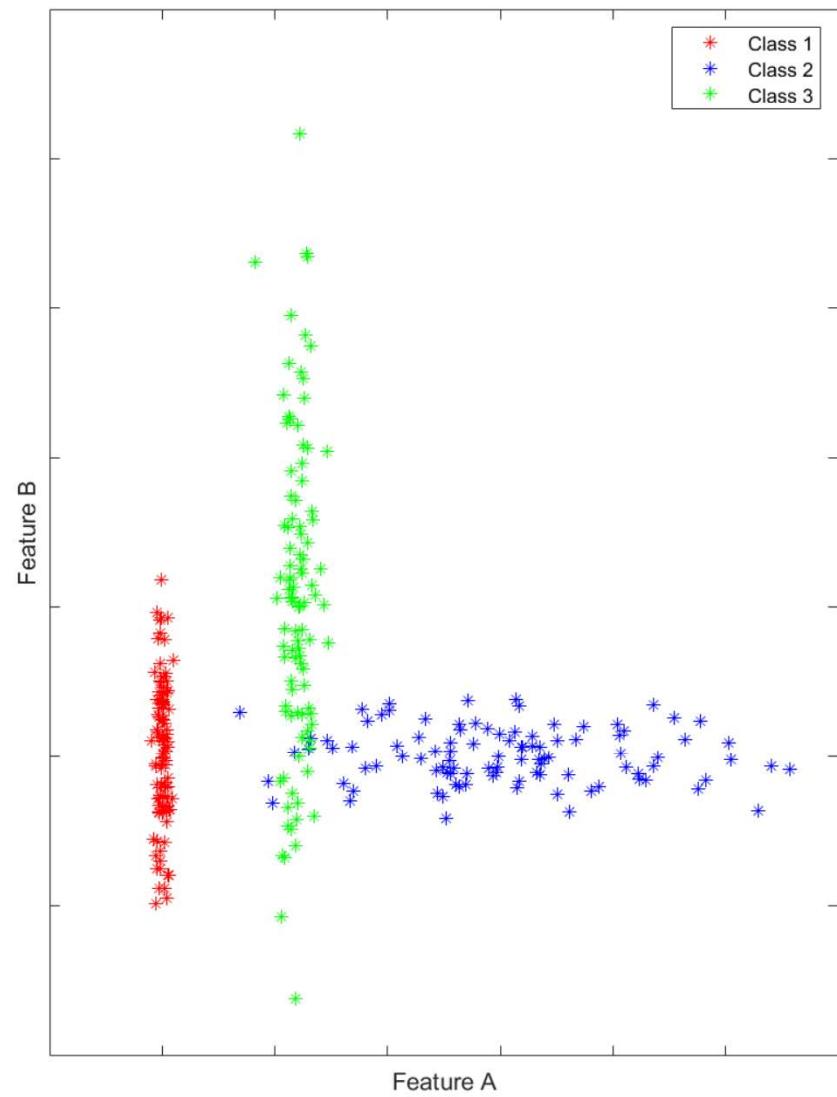
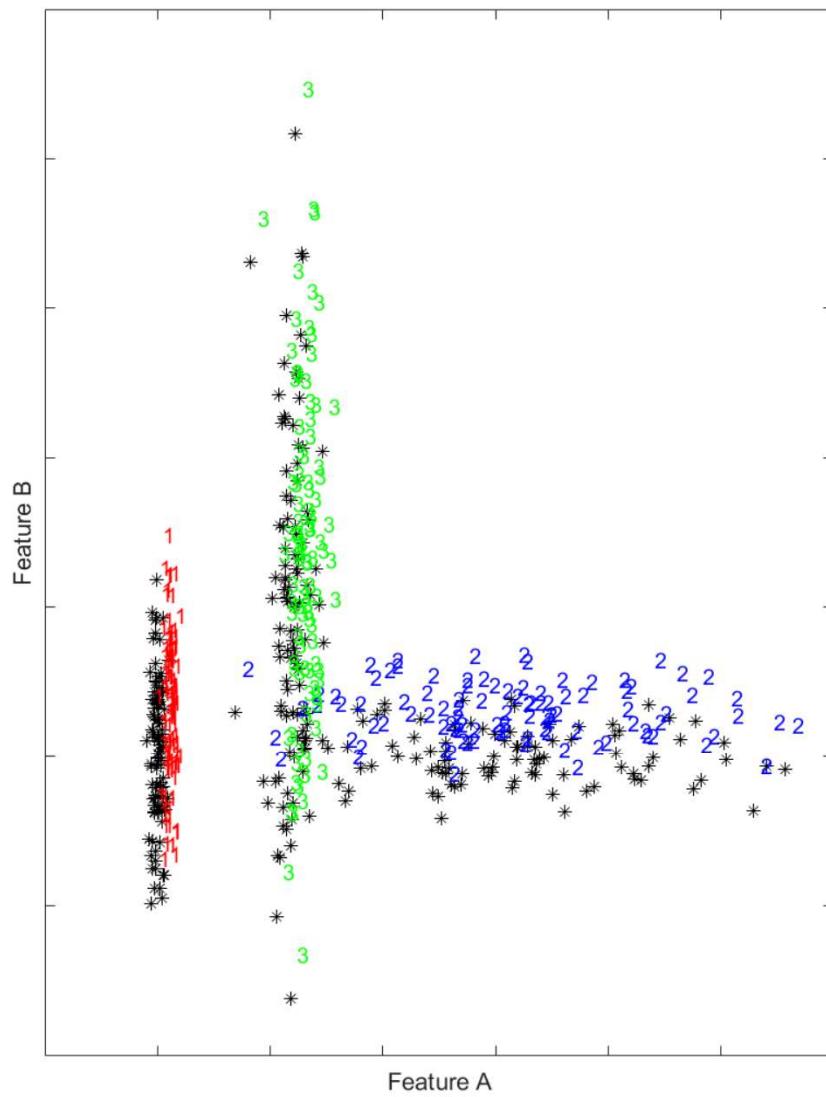
# Supervised Learning: Classification



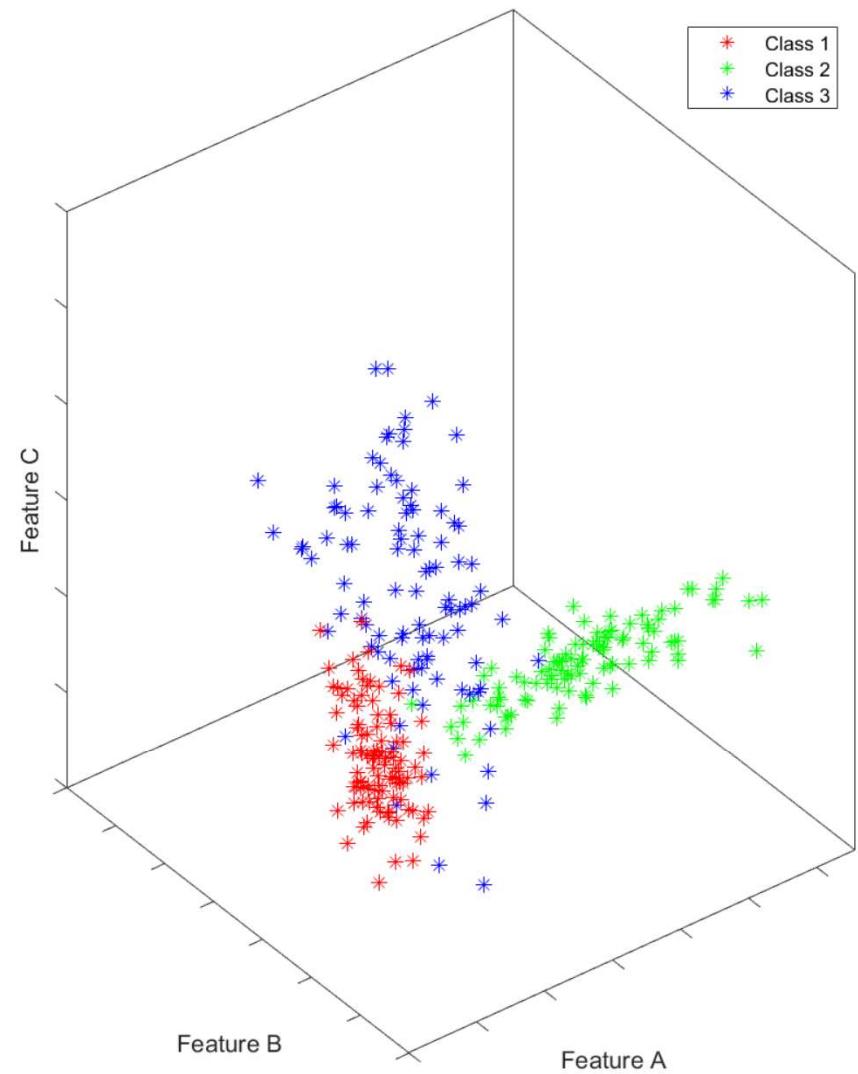
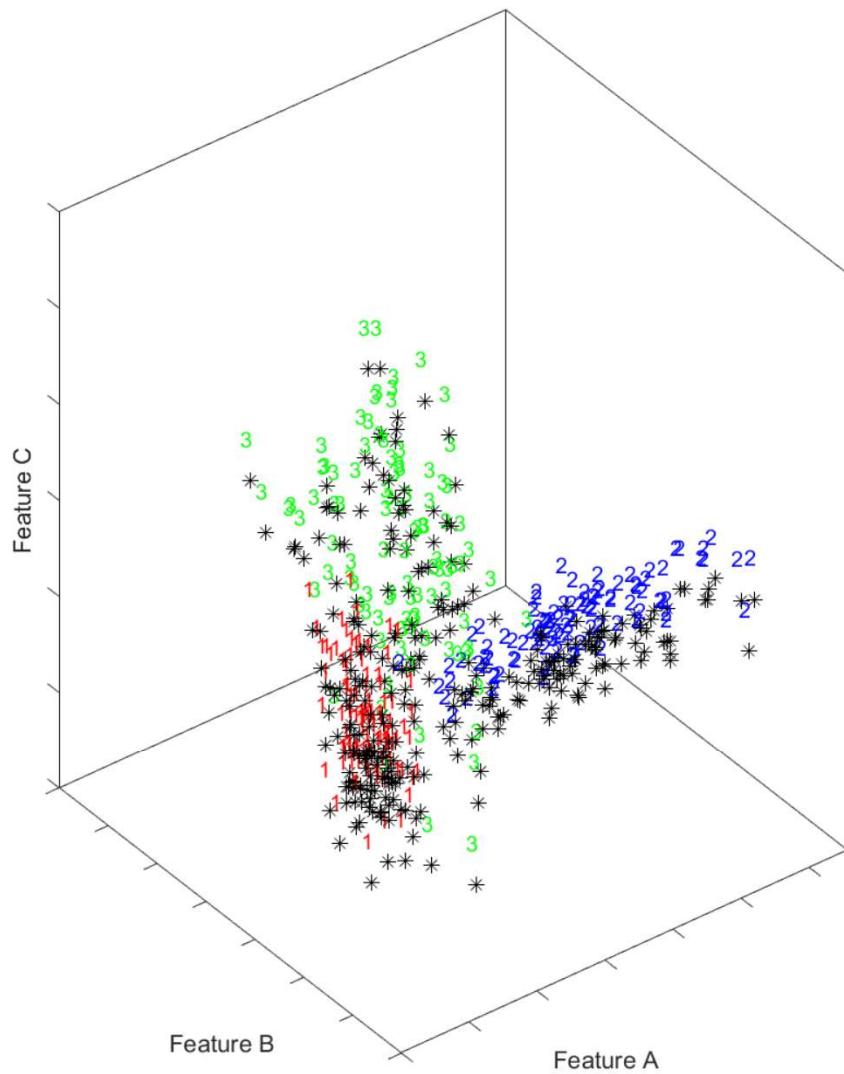
# Supervised Learning: Classification



# Data Set: Labeled Data

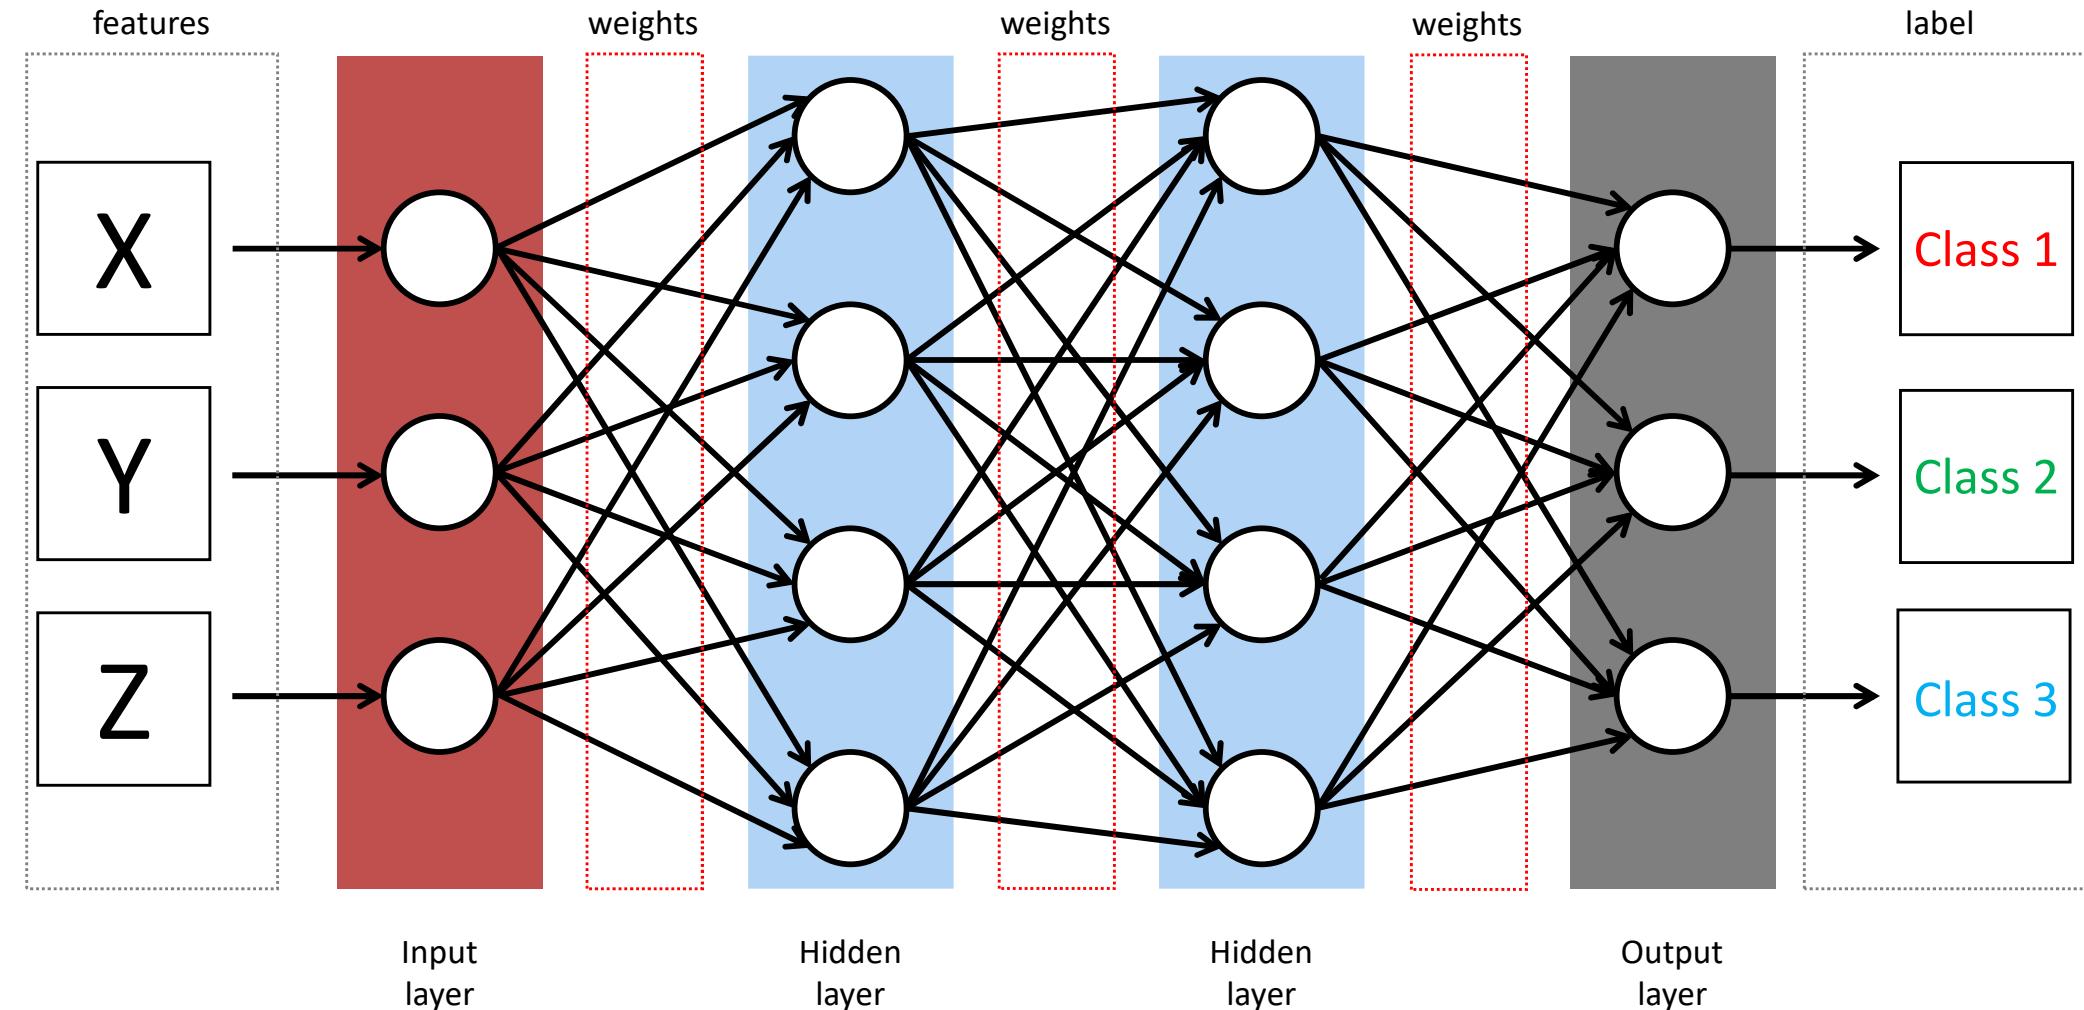


# Data Set: Labeled Data



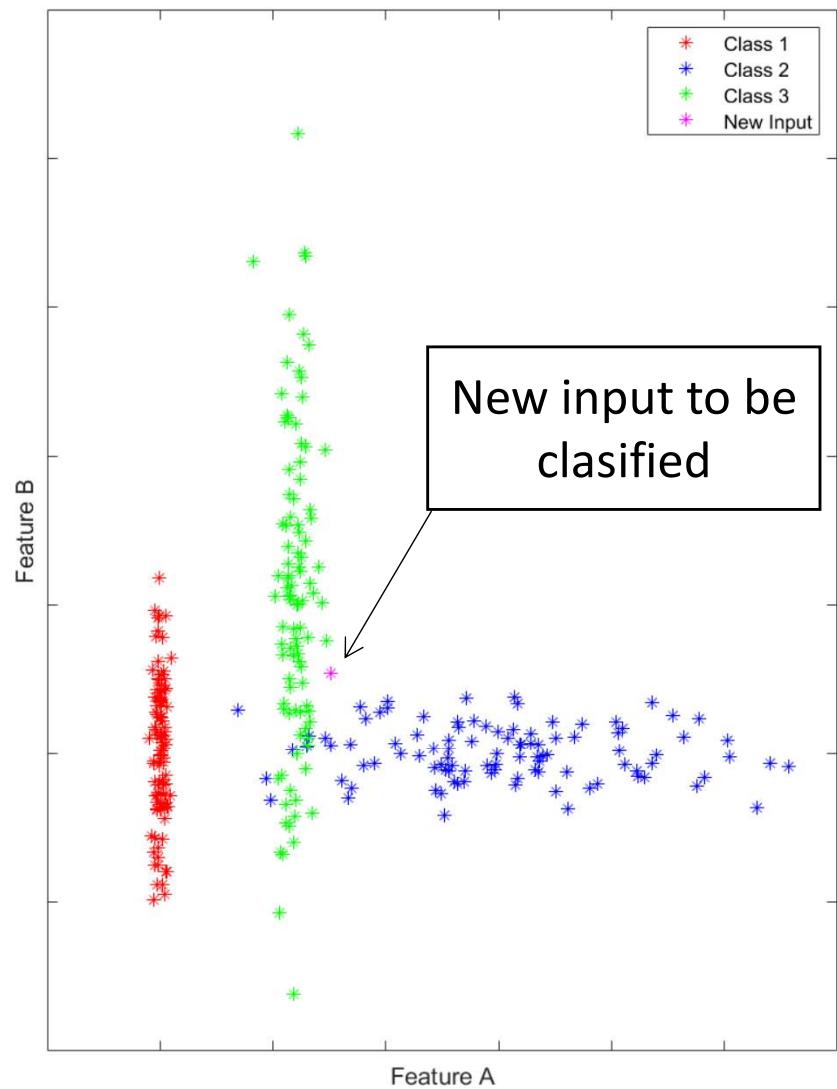
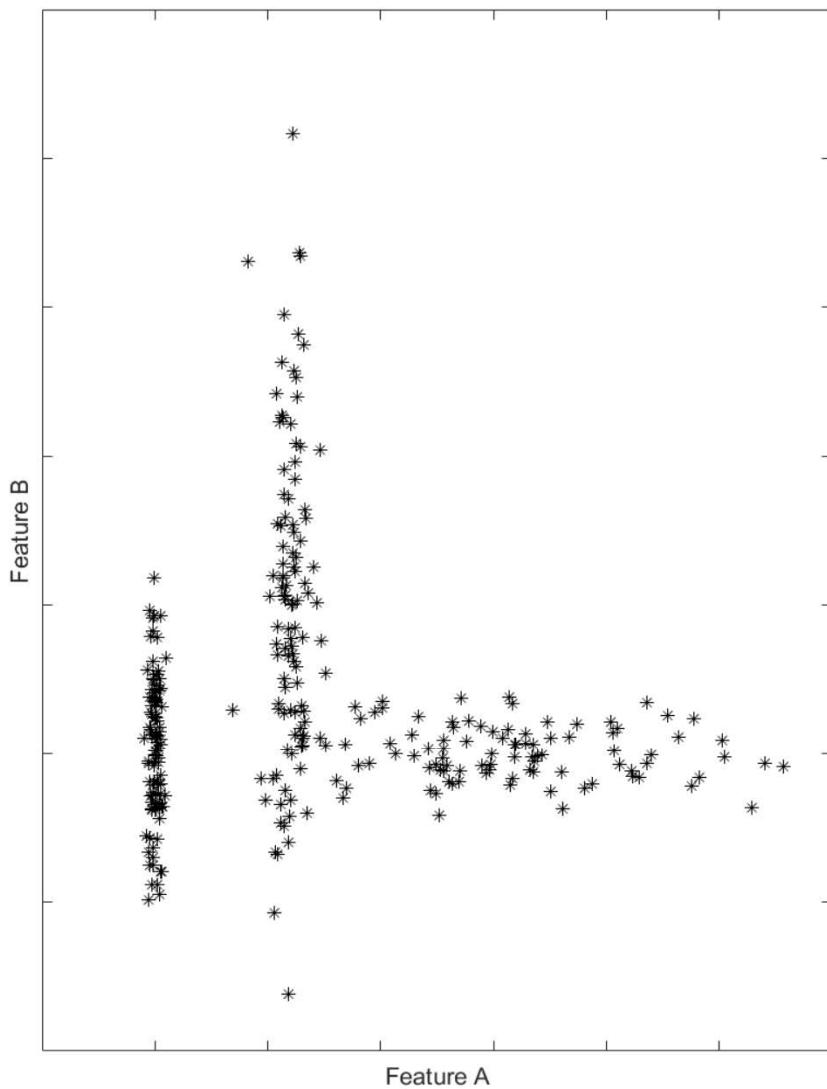
# ANN: Supervised Learning

In order to work properly a classifier **needs to be trained** first with **labeled data**.

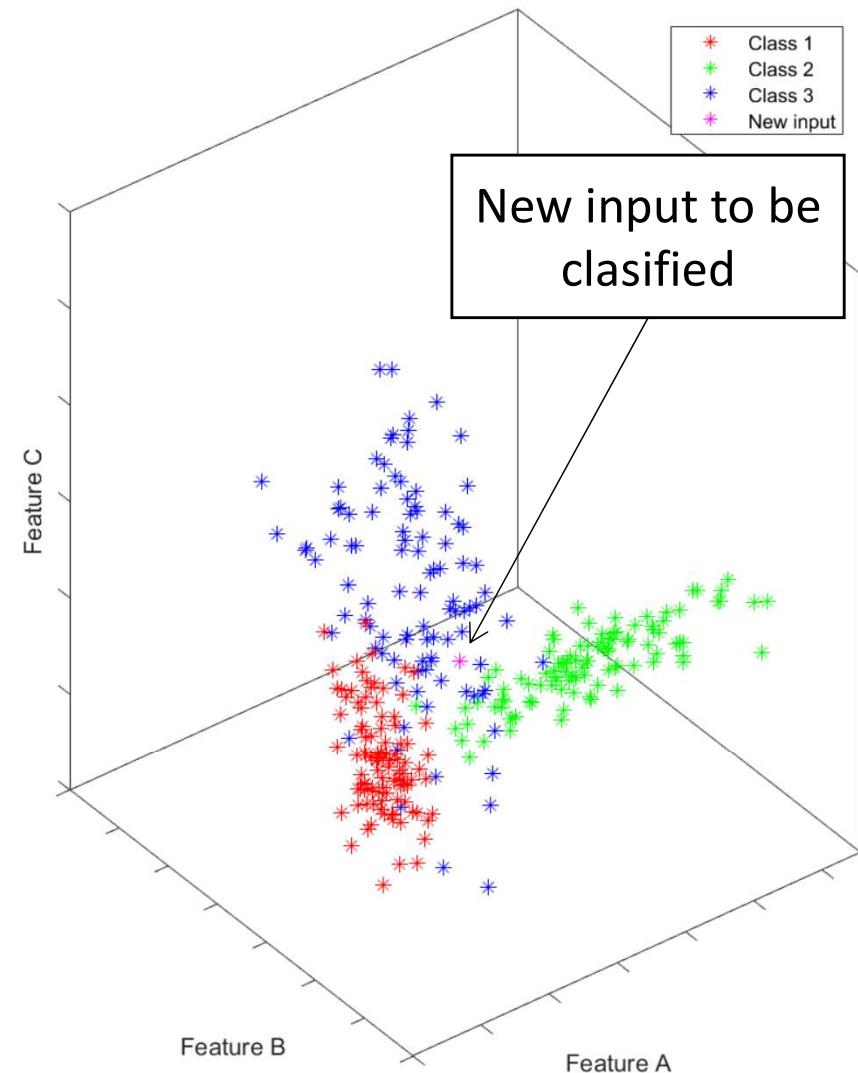
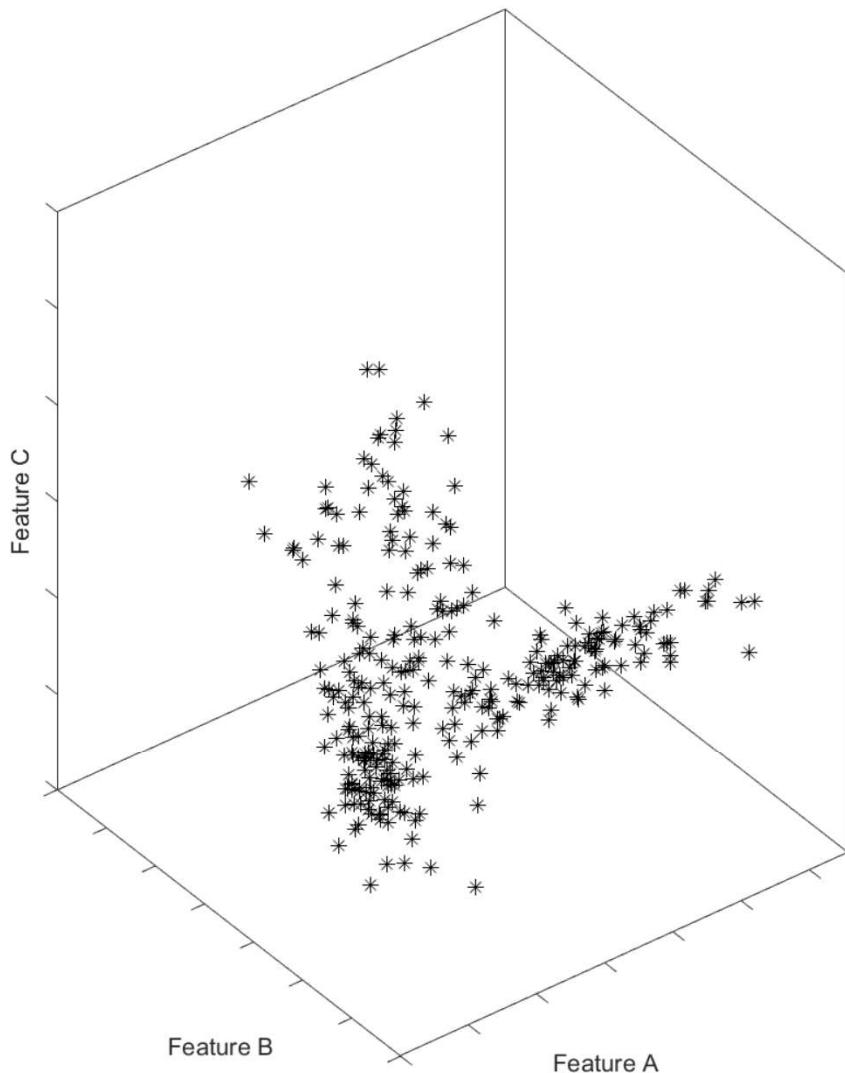


**Training** will adjust all the weights within this artificial neural network.

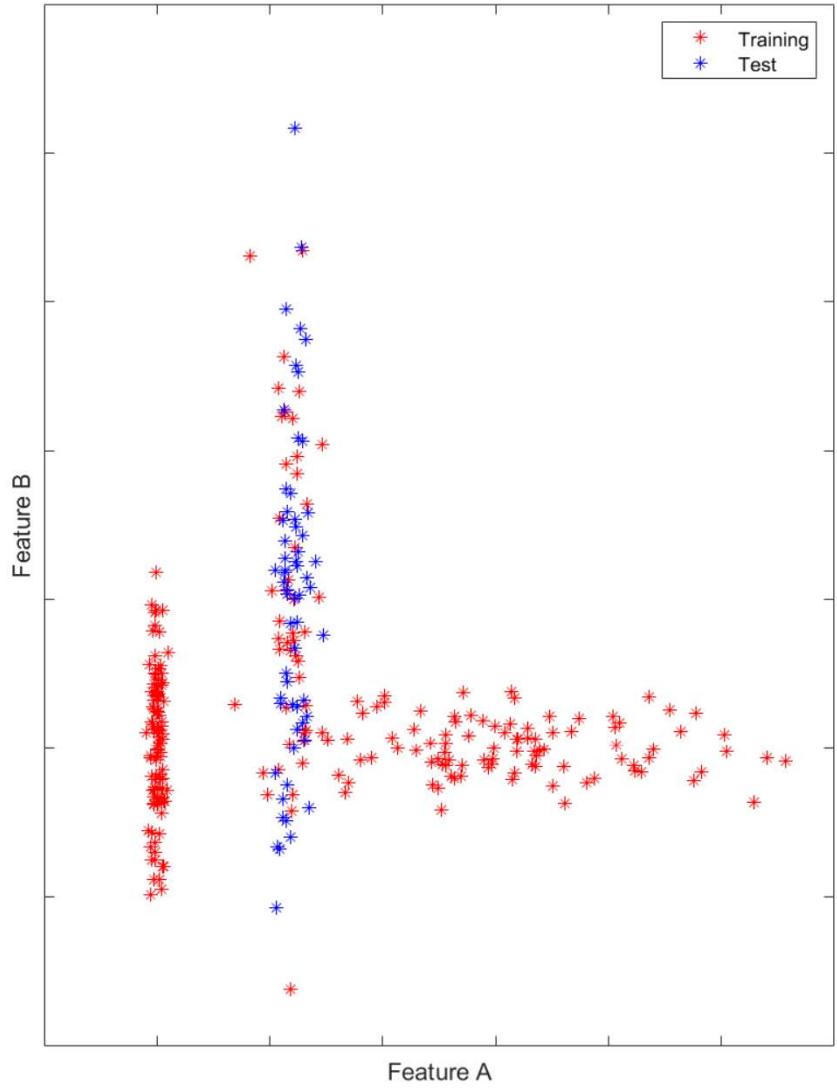
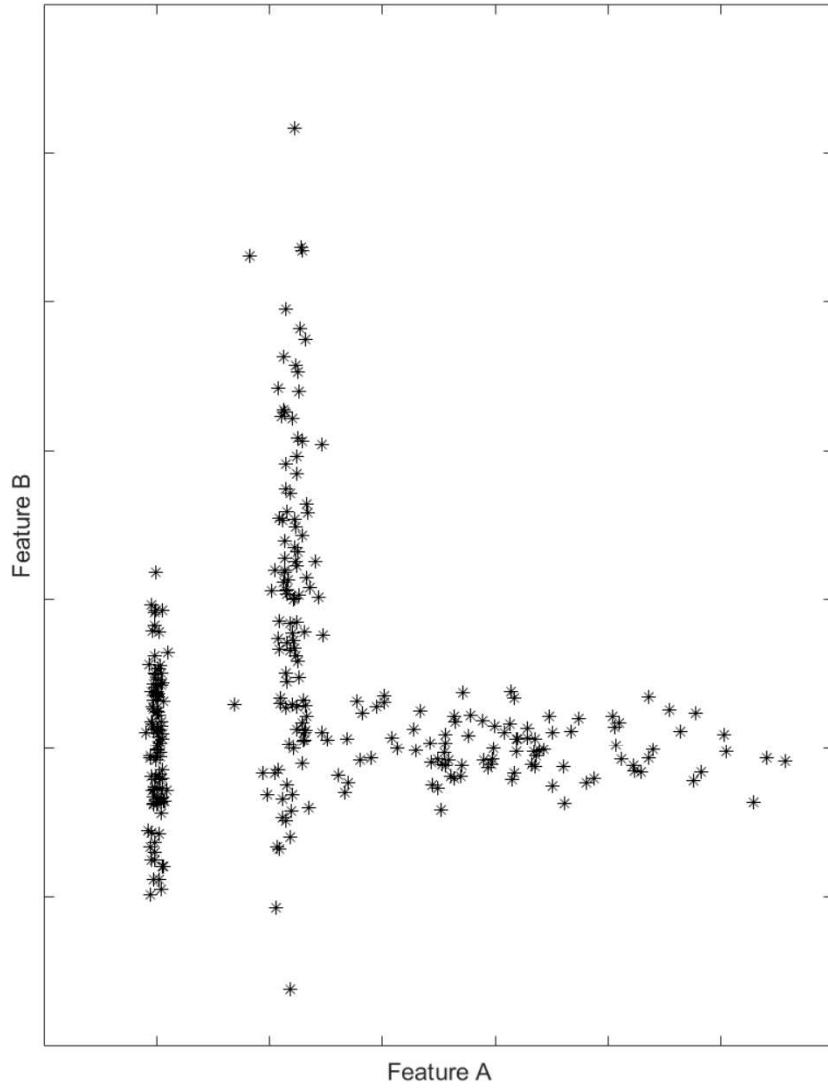
# Supervised Learning: New Input



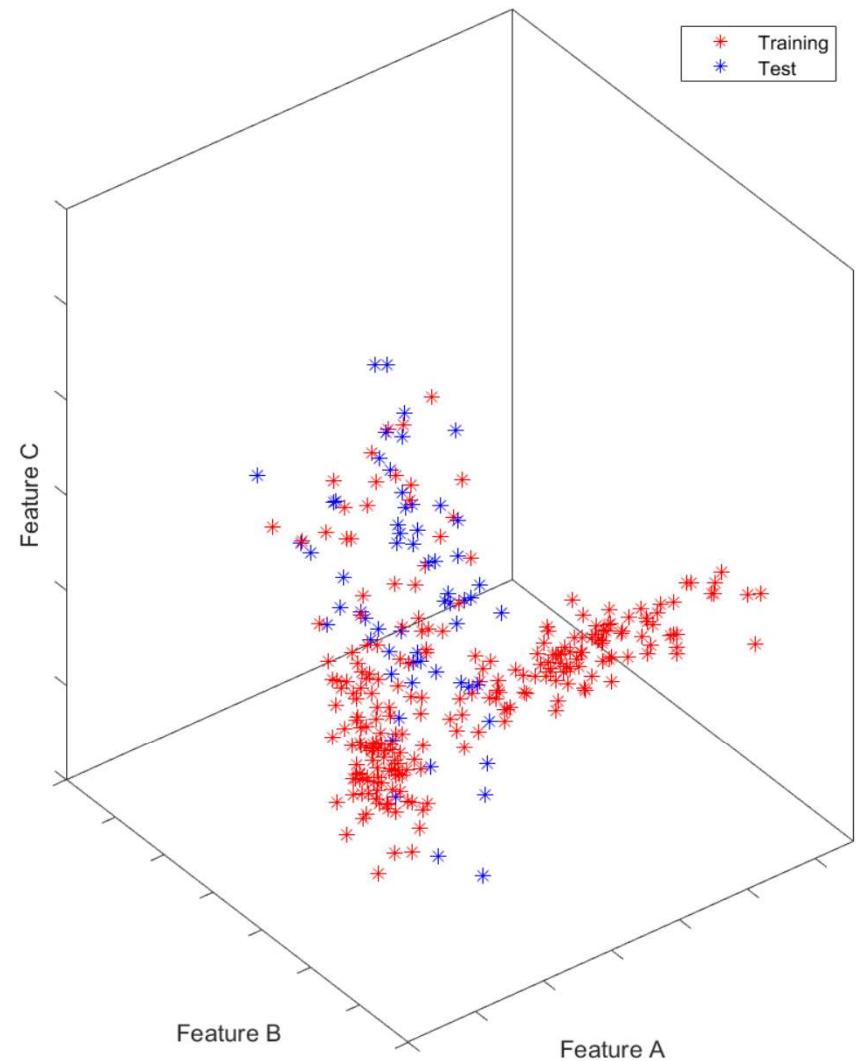
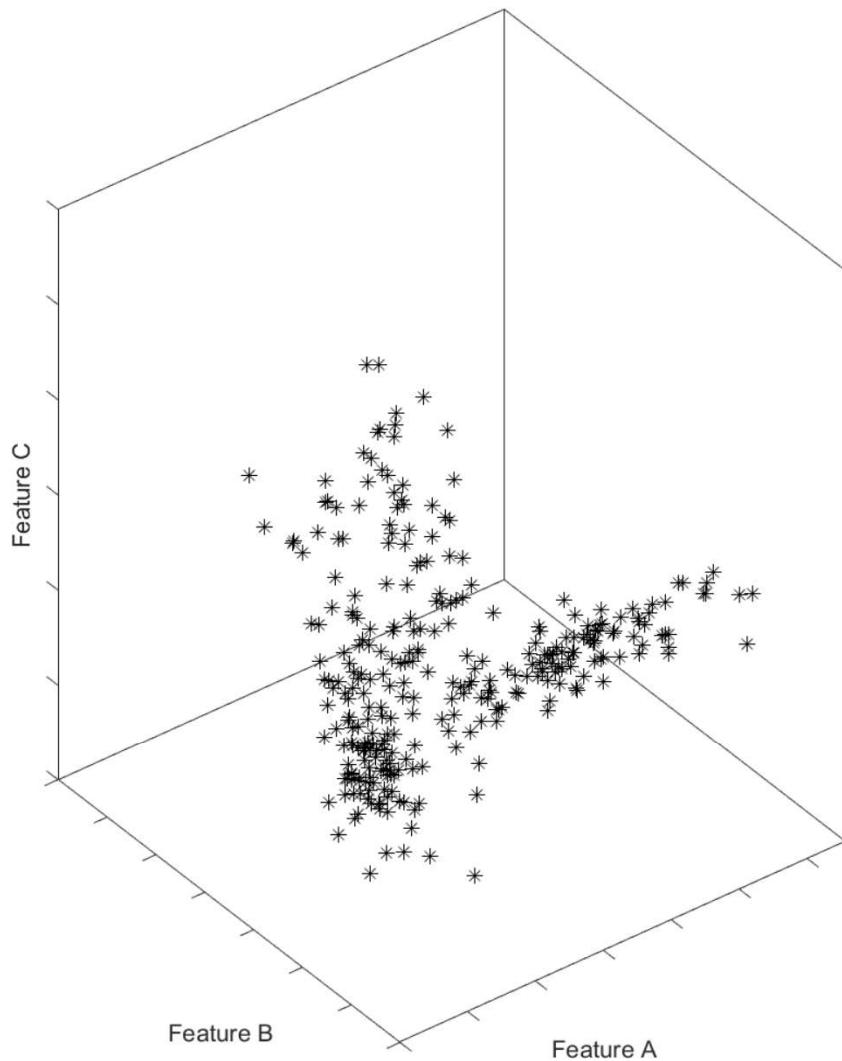
# Supervised Learning: New Input



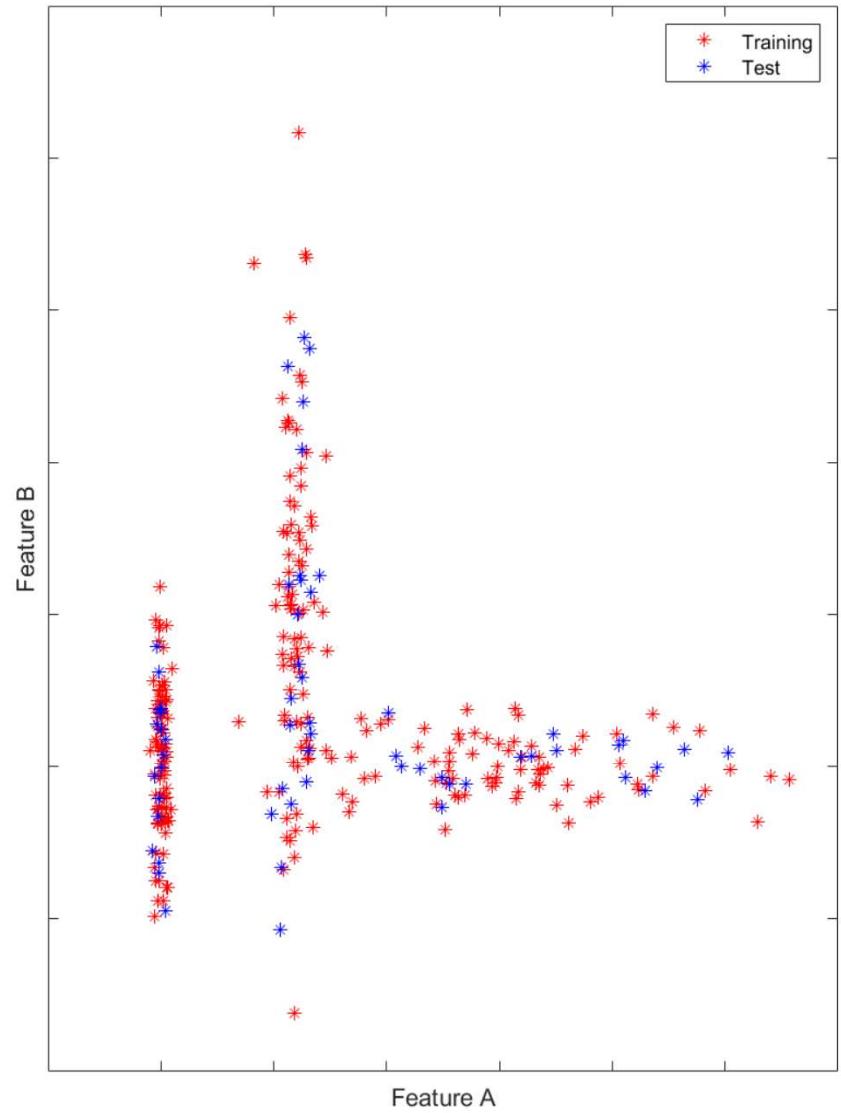
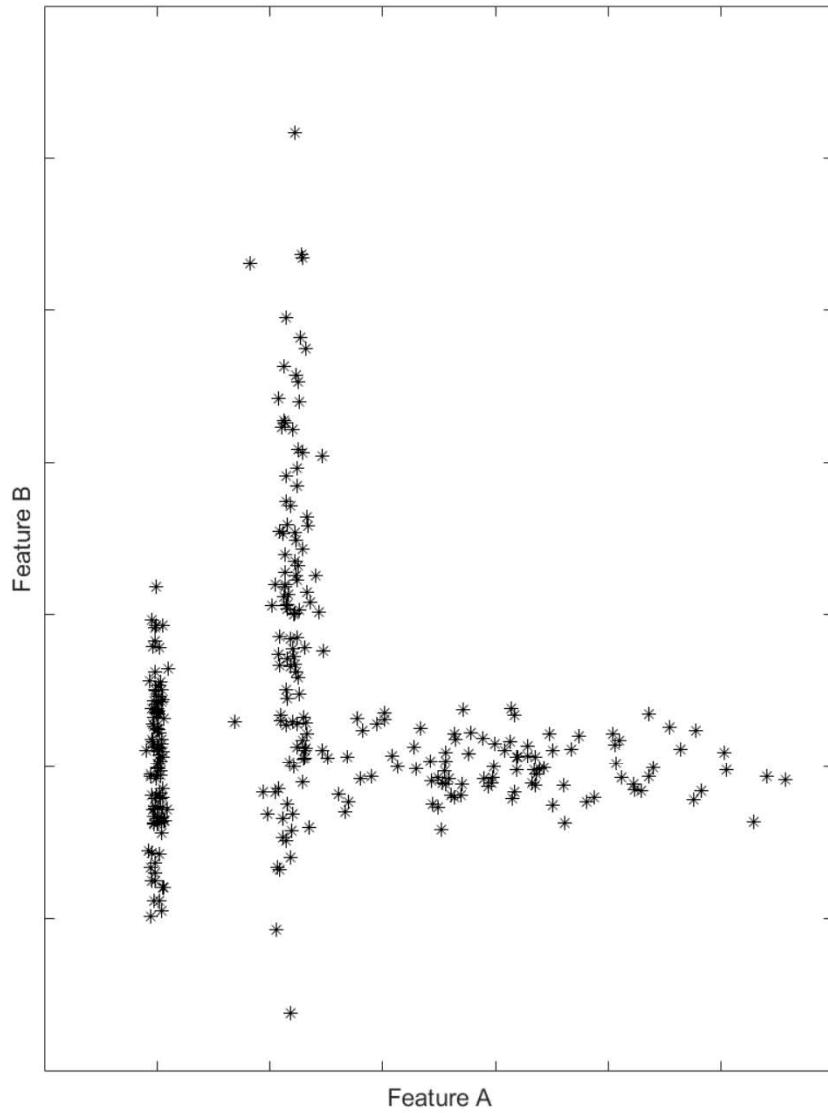
# Training / Test Sets: Poor Split



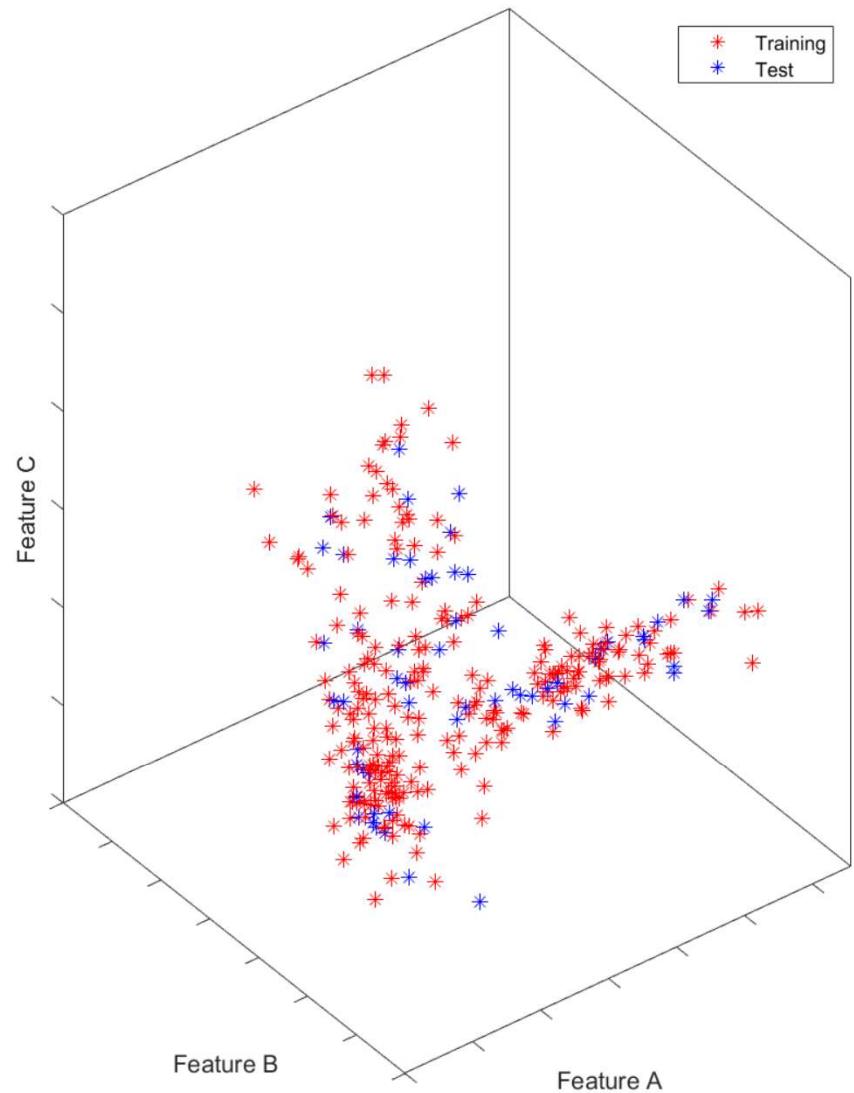
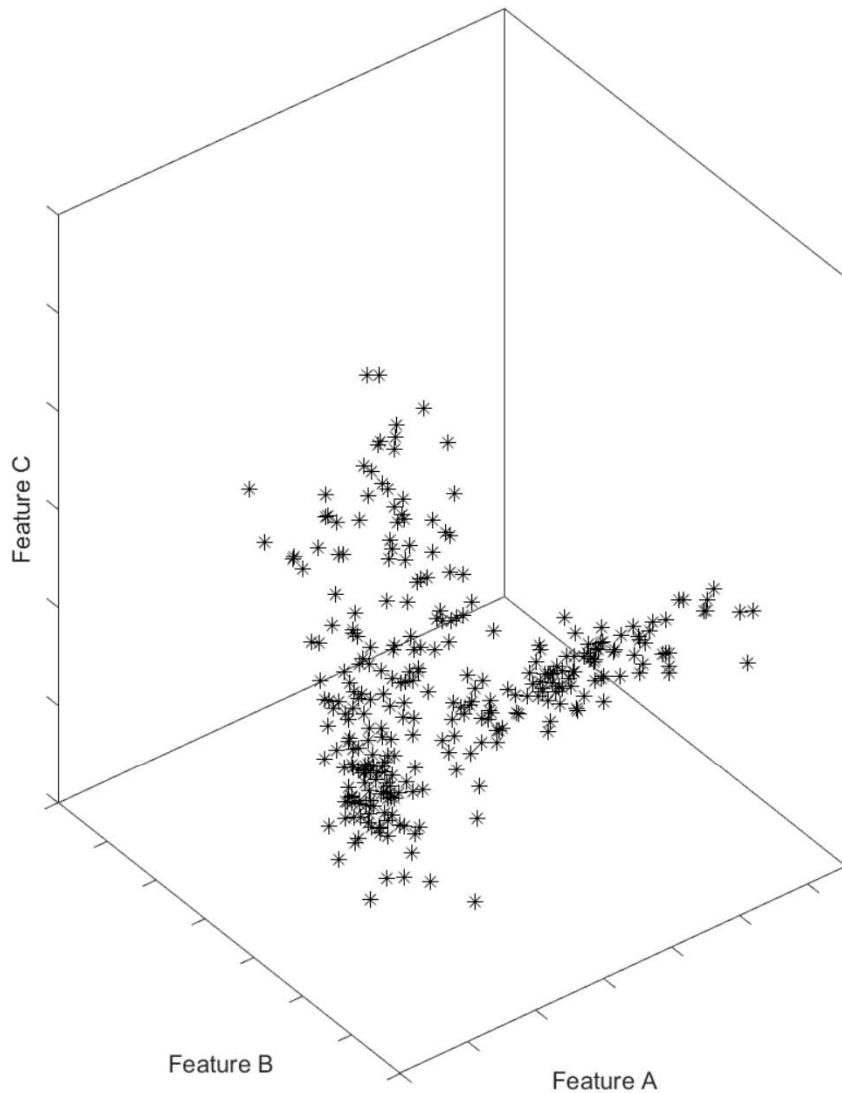
# Training / Test Sets: Poor Split



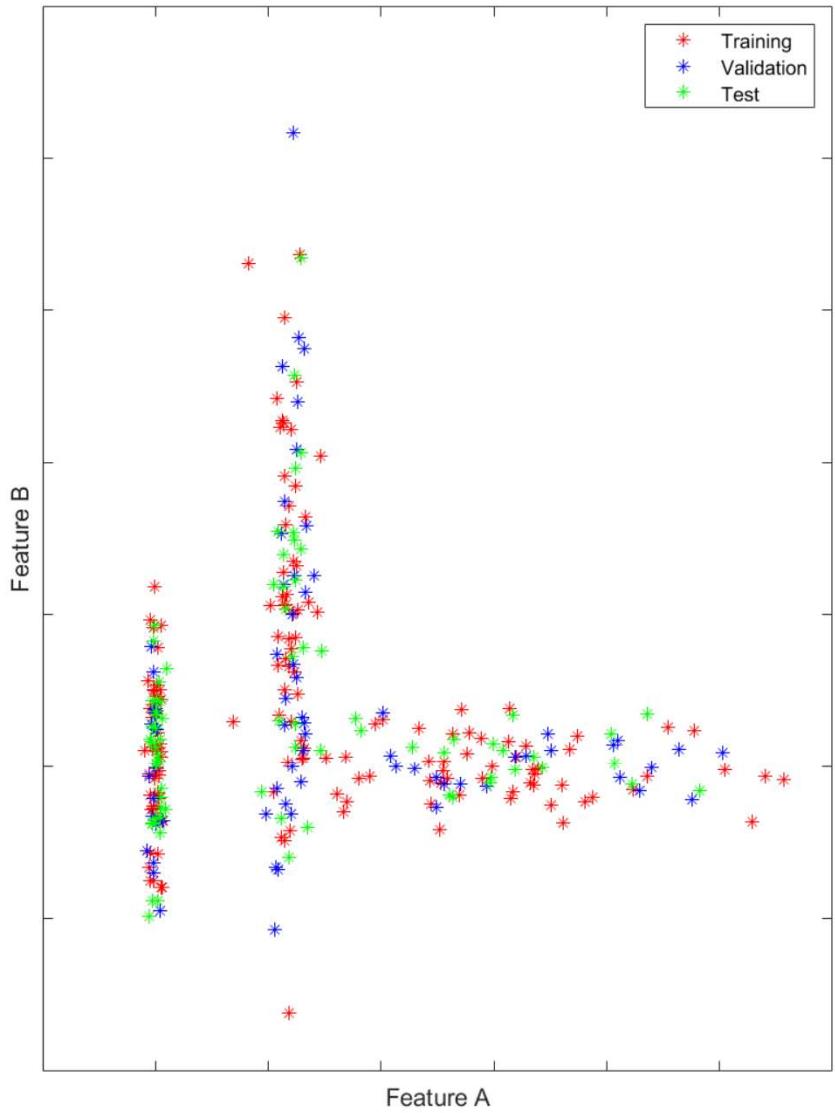
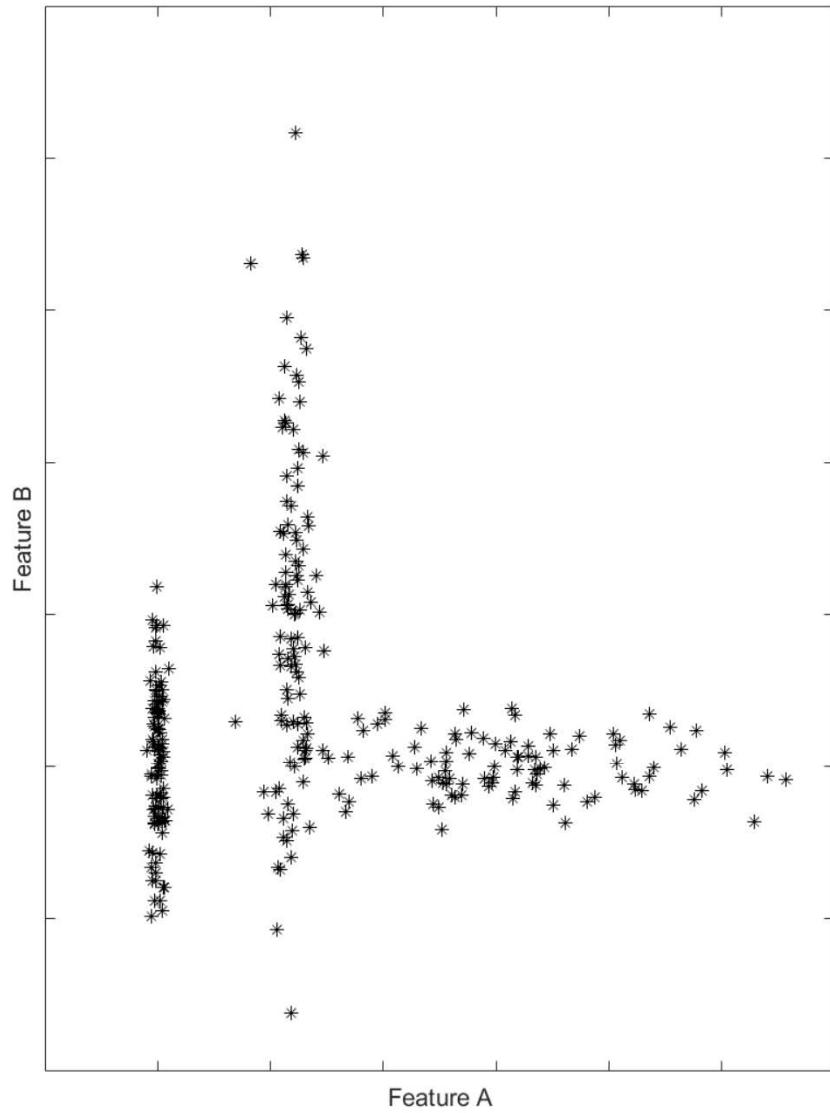
# Training / Test Sets: Better Split



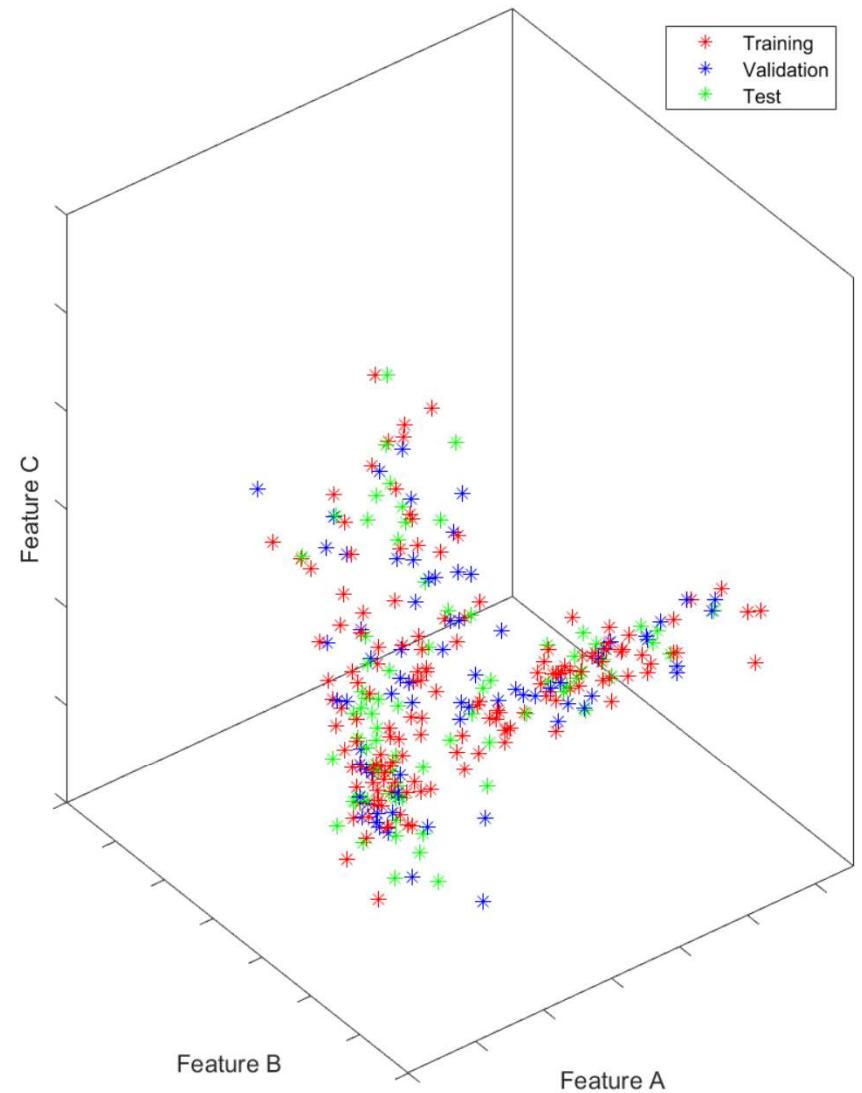
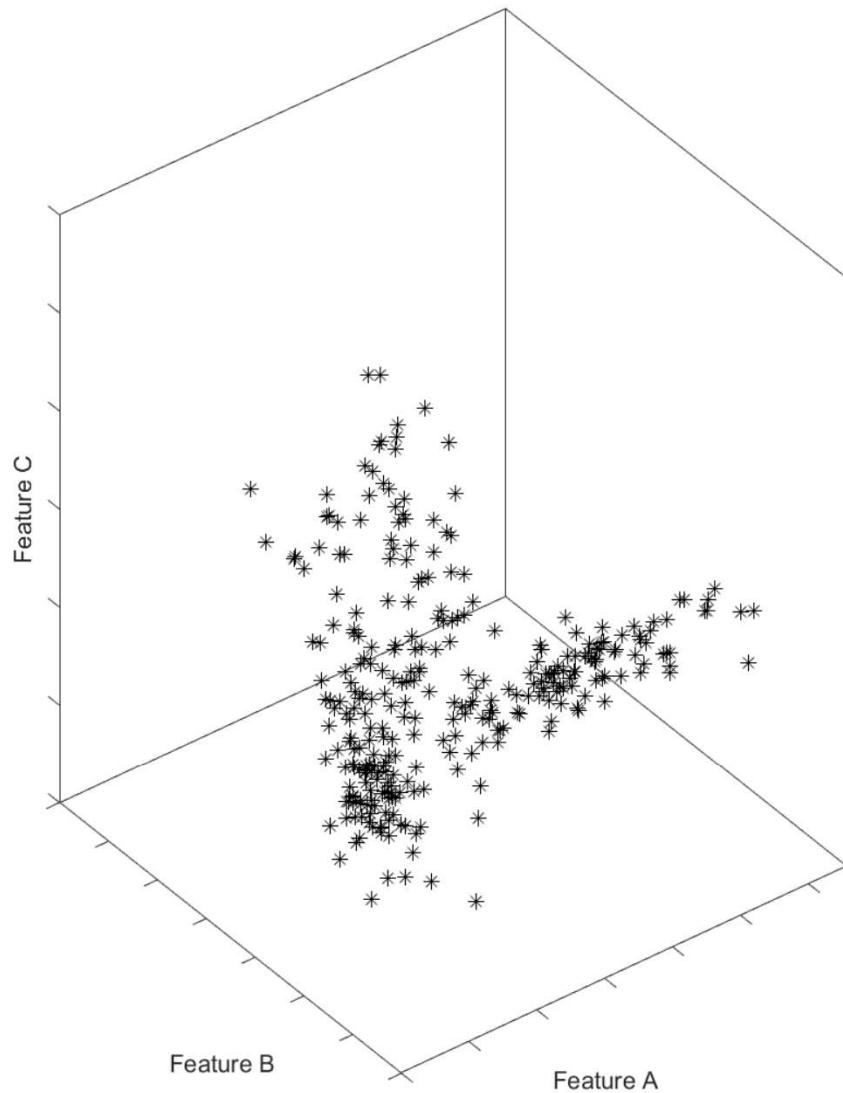
# Training / Test Sets: Better Split



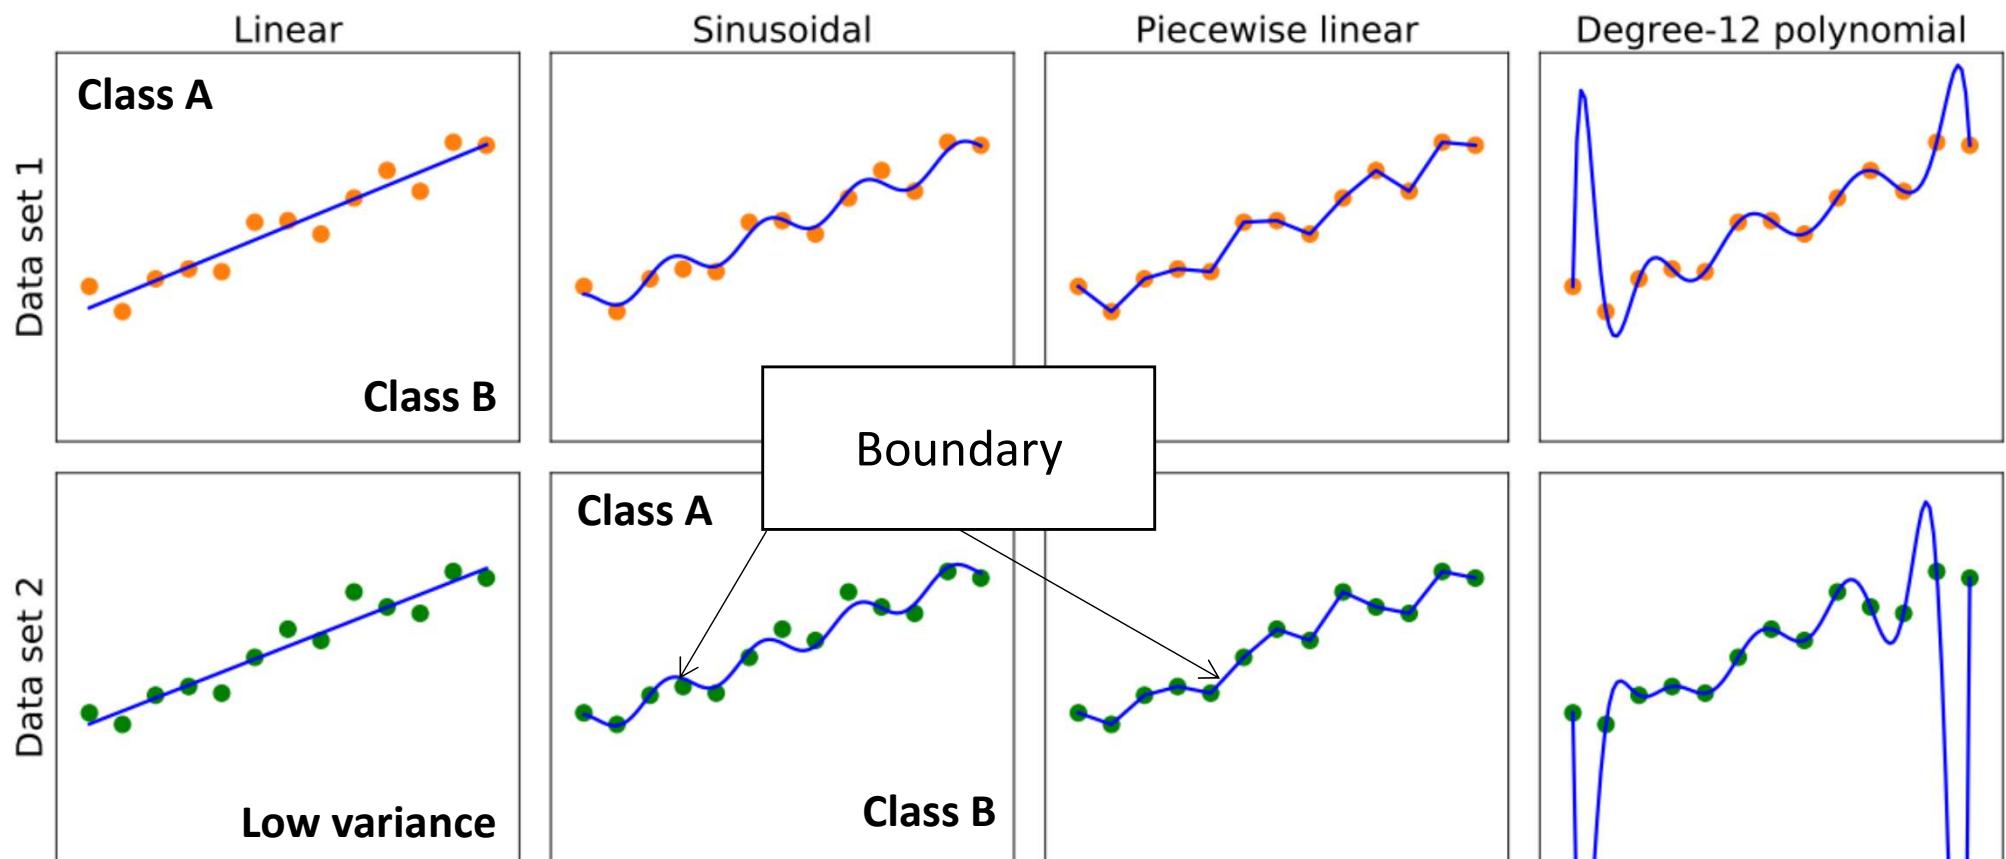
# Training / Validation / Test Sets



# Training / Validation / Test Sets



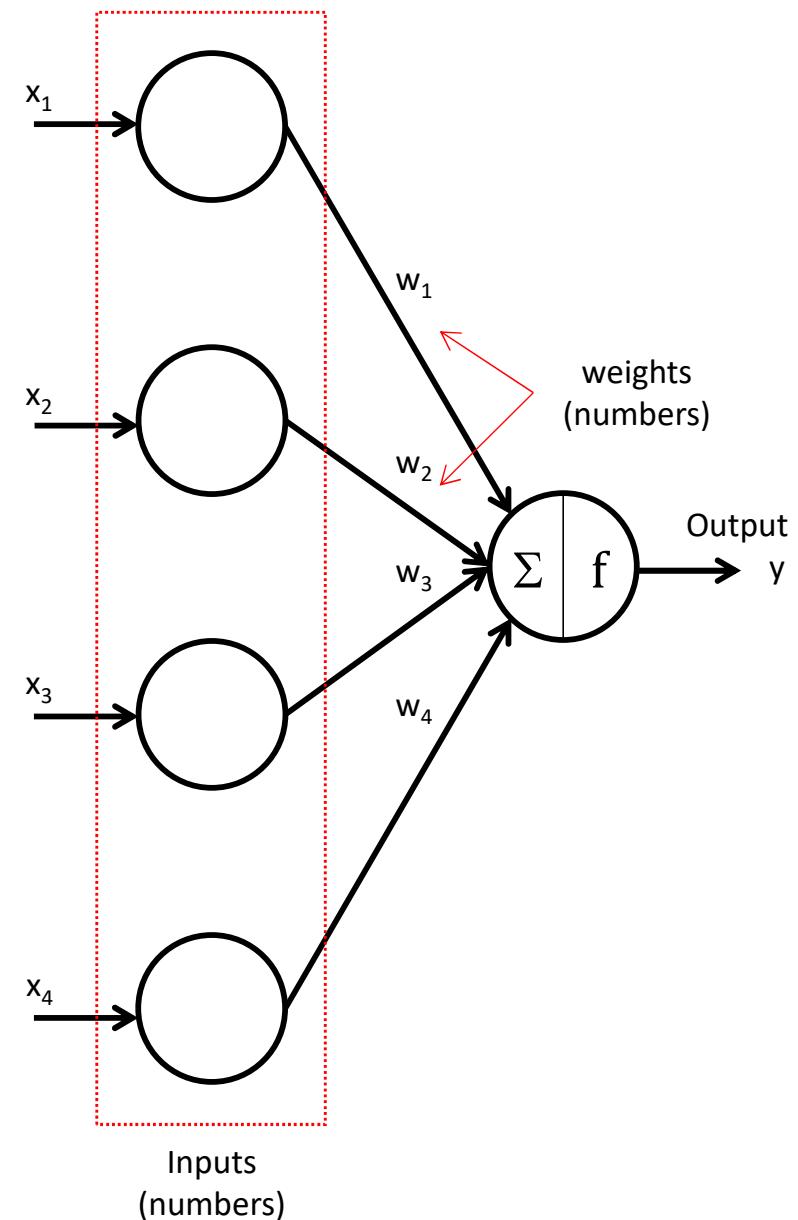
# Hypothesis: Decision “Boundary”



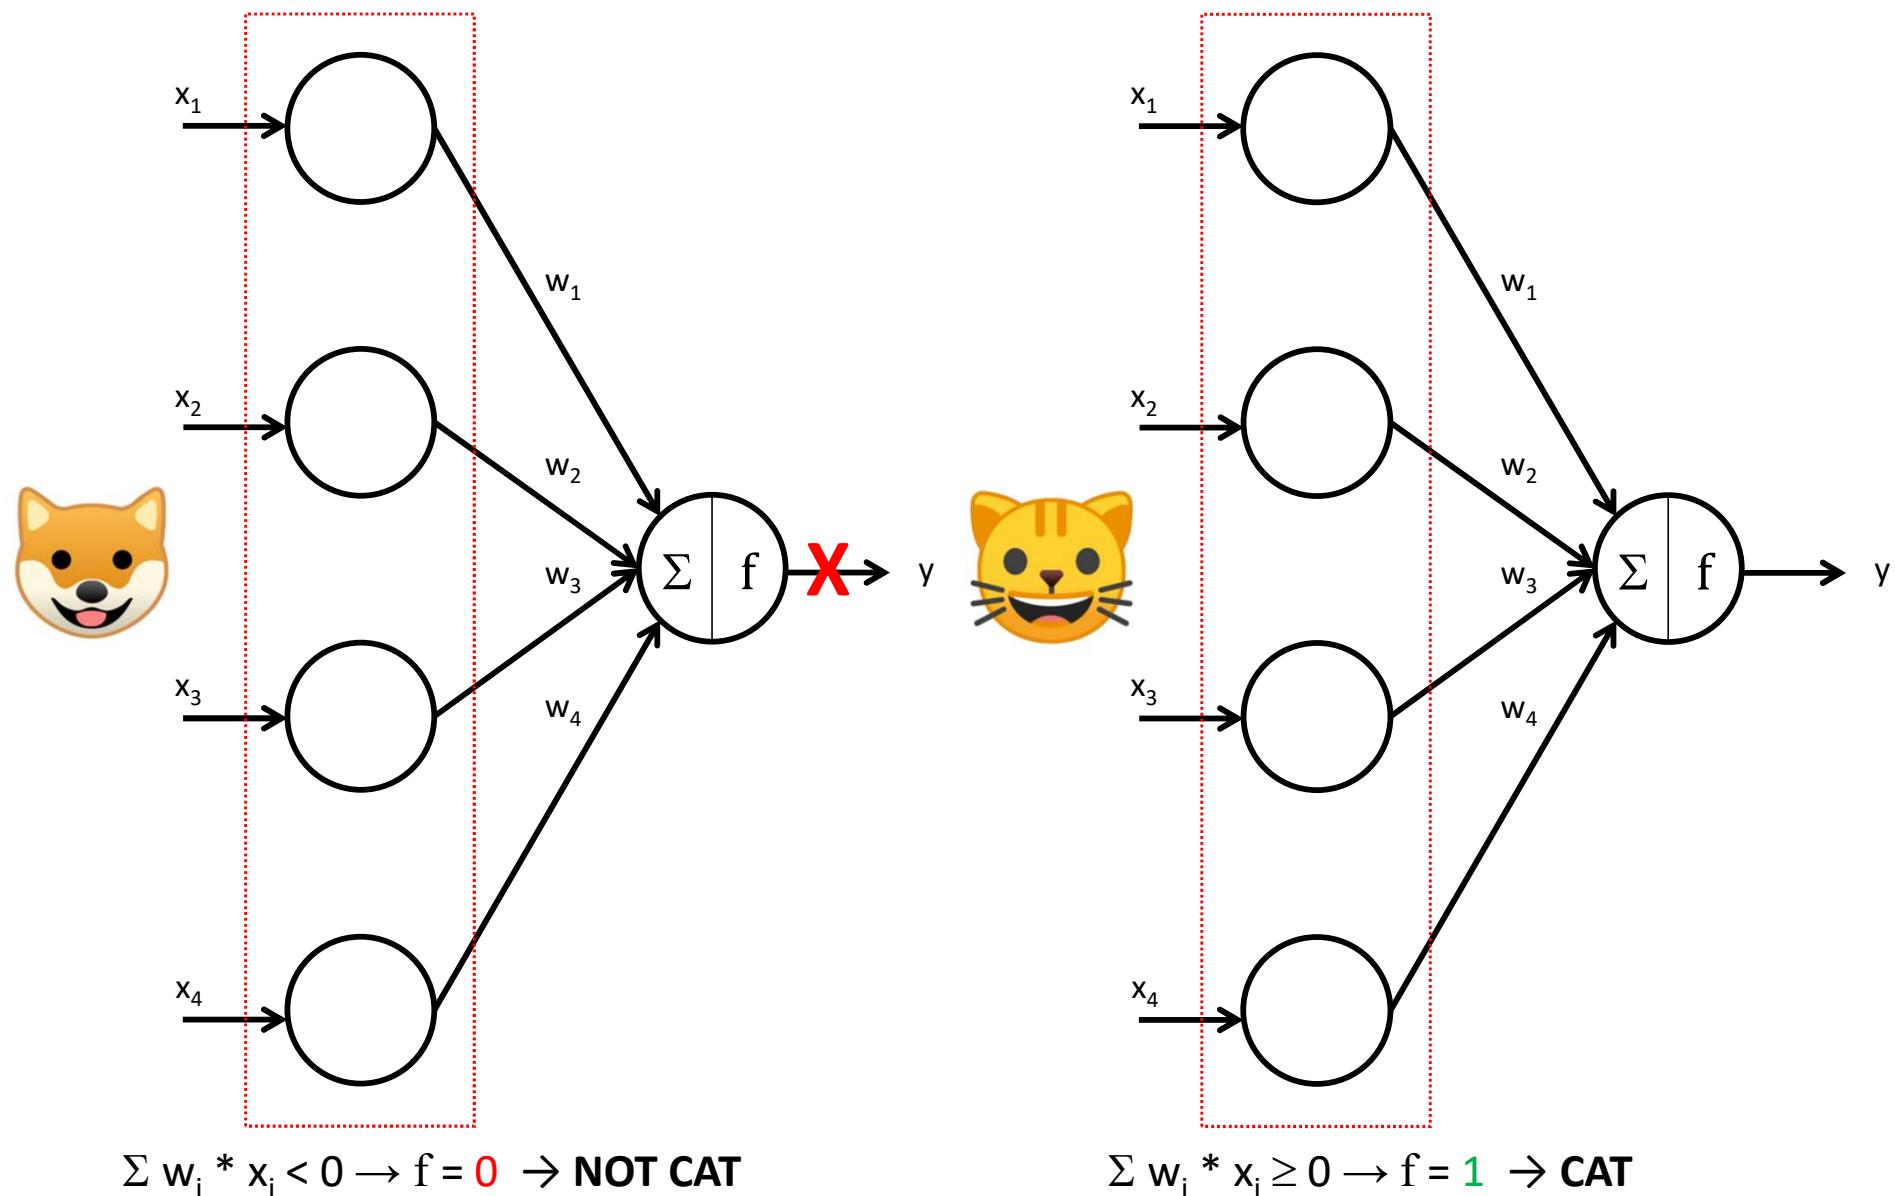
# Artificial Neuron (Perceptron)

A (single-layer) **perceptron** is a model of a biological neuron. It is made of the following components:

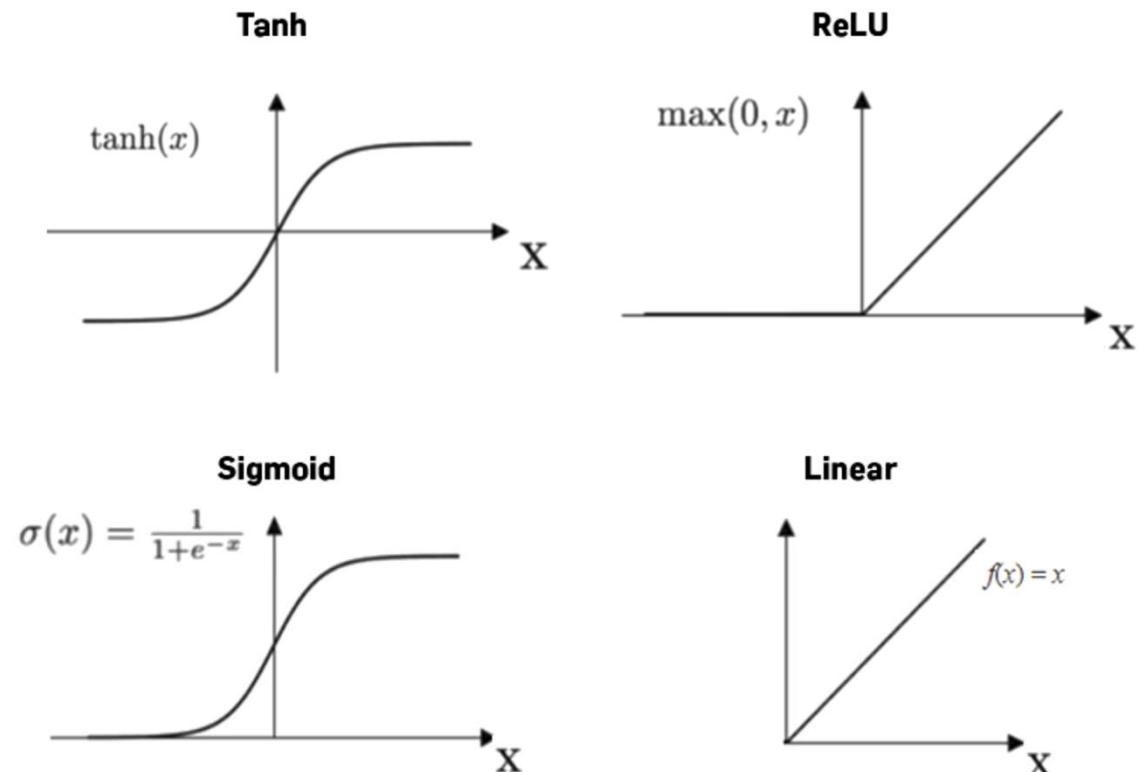
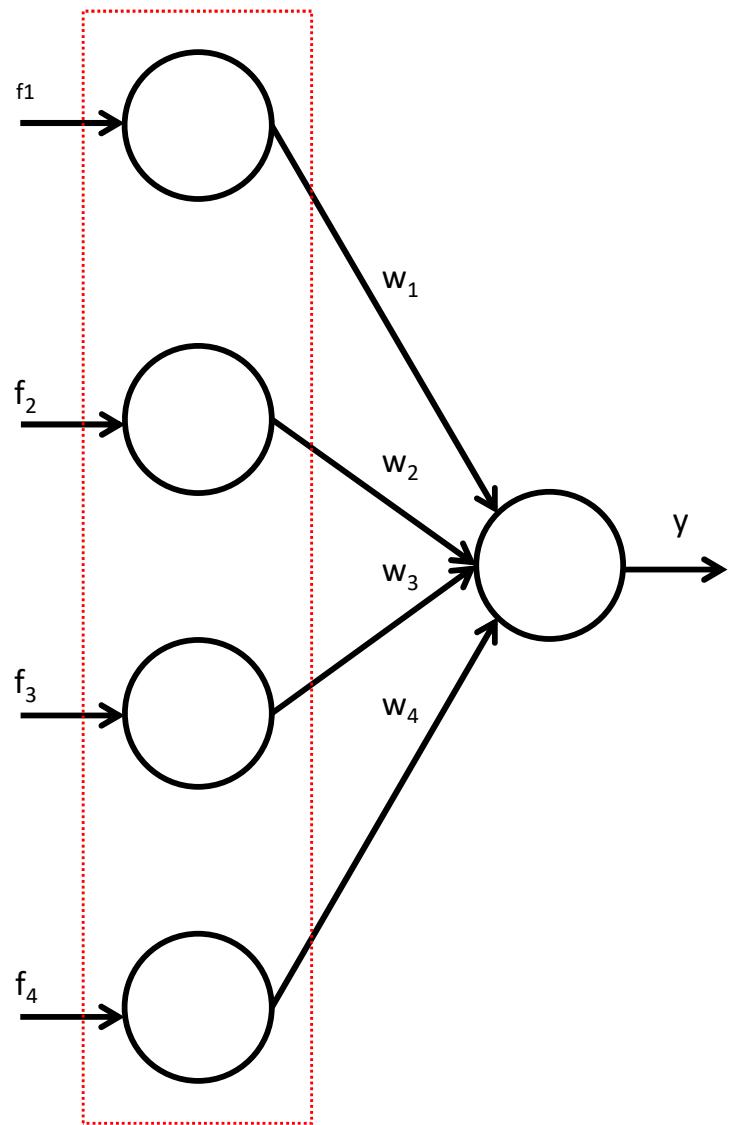
- inputs  $x_i$  - numerical values representing information
- weights  $w_i$  - numerical values representing how “important” corresponding input is
- weighted sum:  $\sum w_i * x_i$
- activation function  $f$  that decides if the neuron “fires”



# Single-layer Perceptron as a Classifier

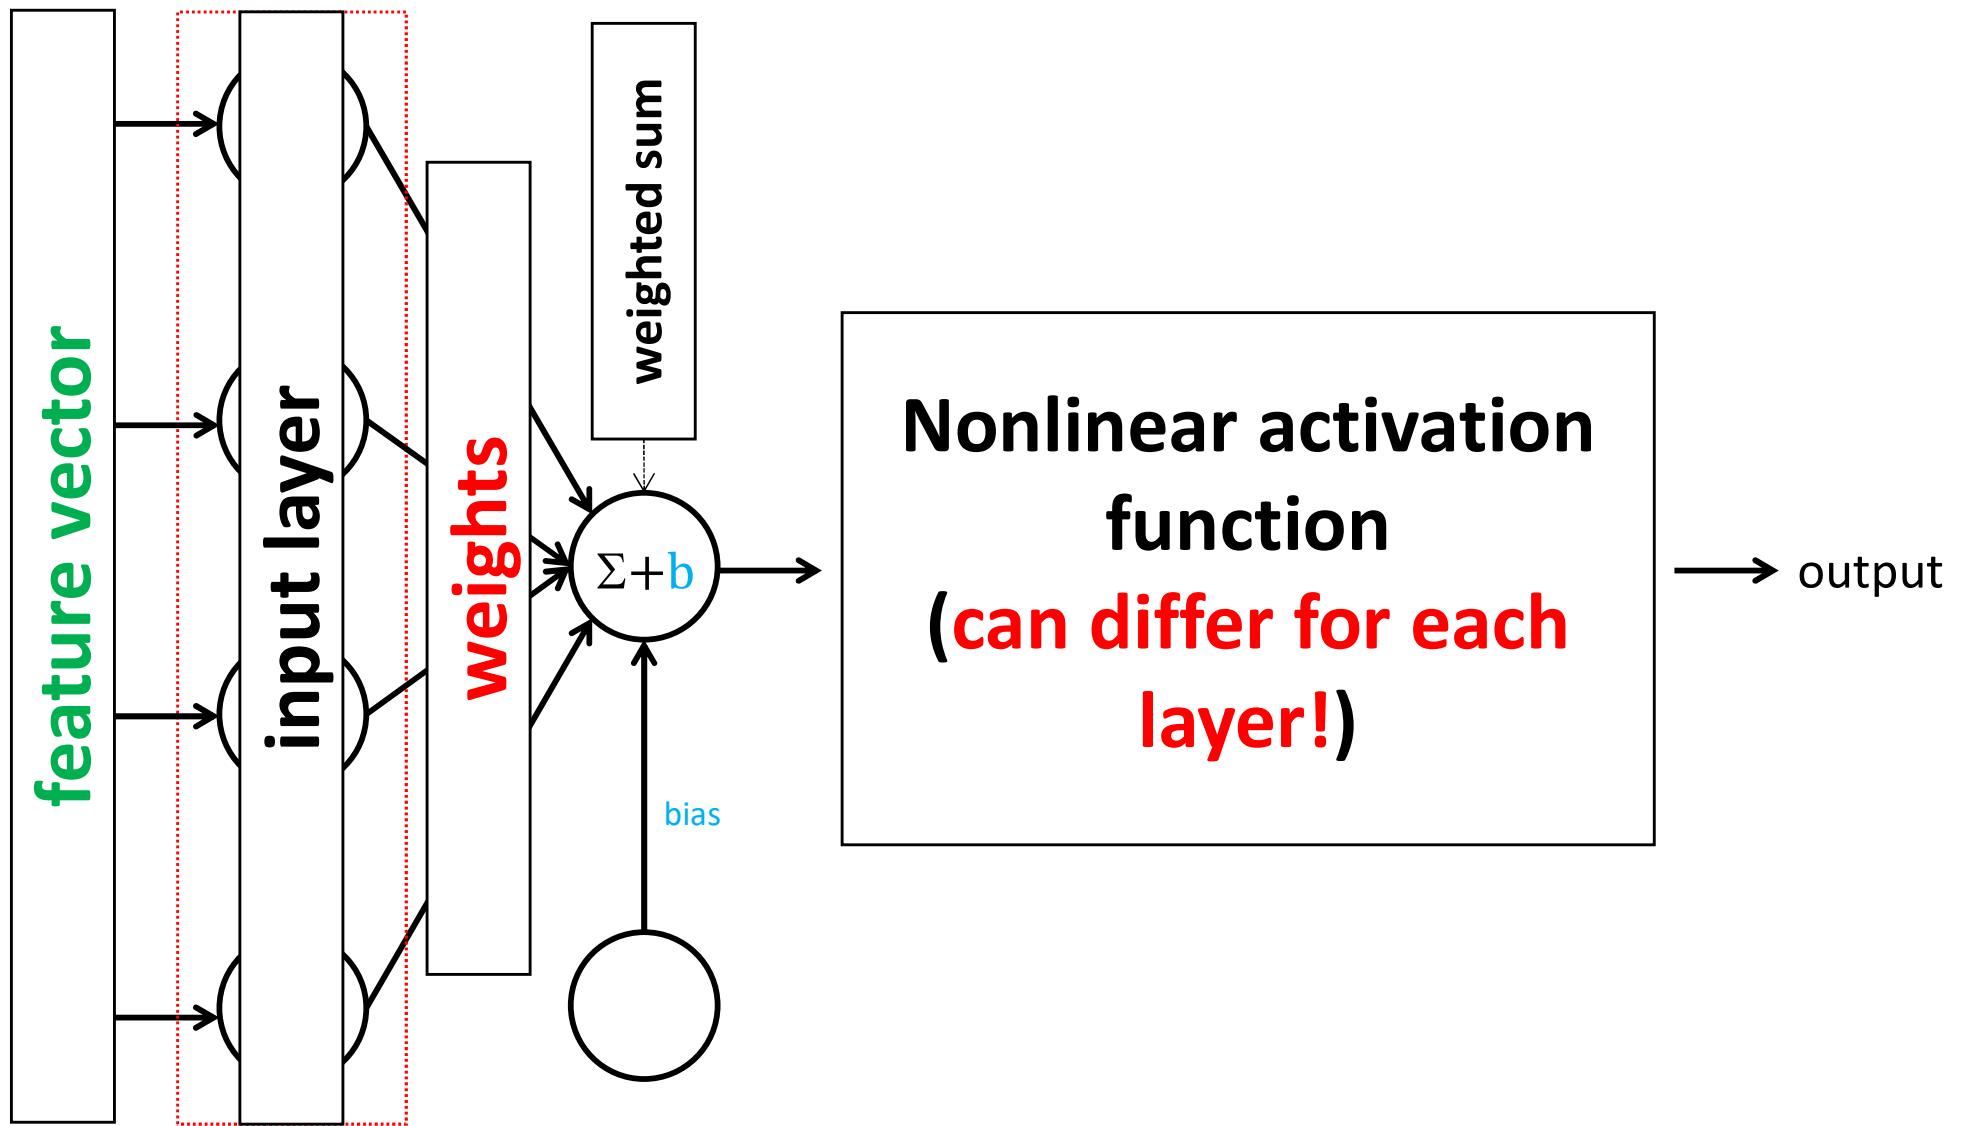


# Selected Activation Functions

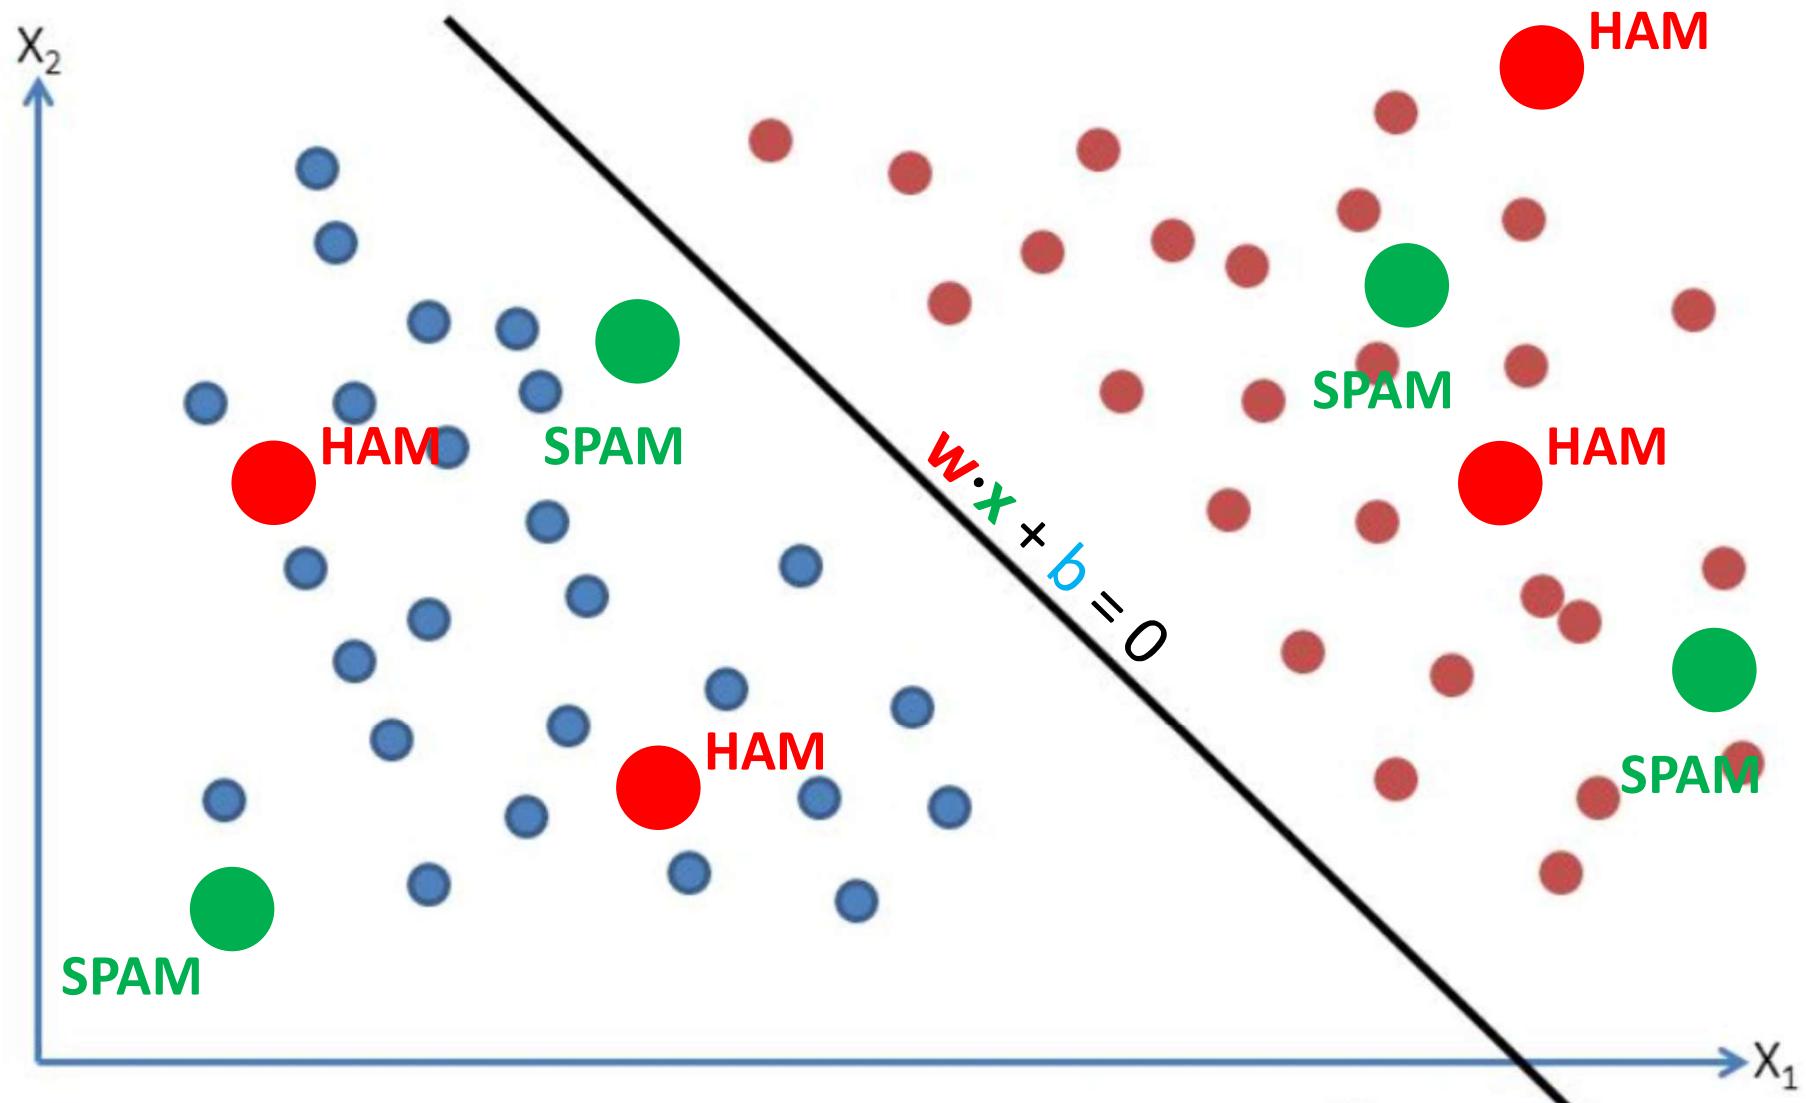


ReLU: Rectified Linear Unit

# Basic Neural Unit

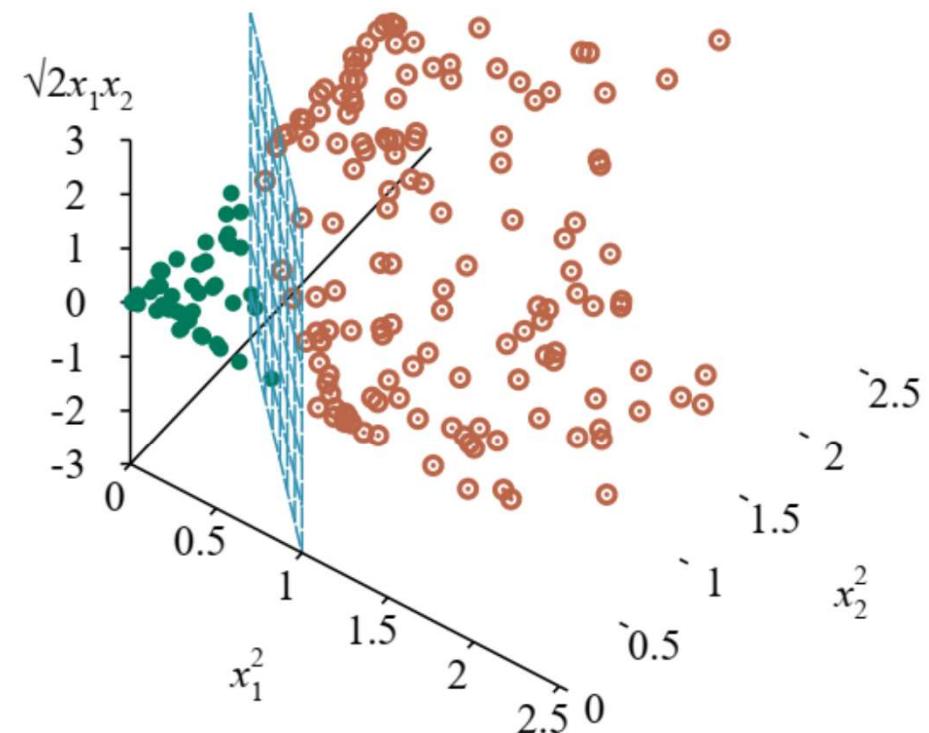
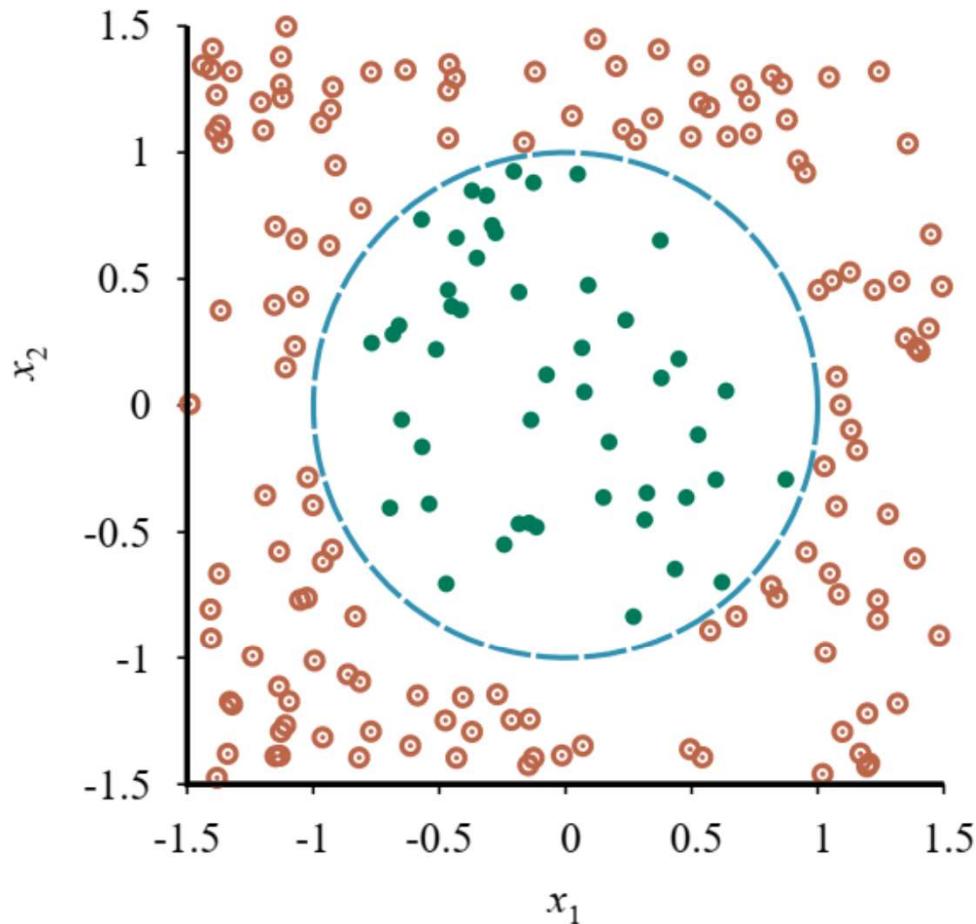


# Classification: Linear Separation?

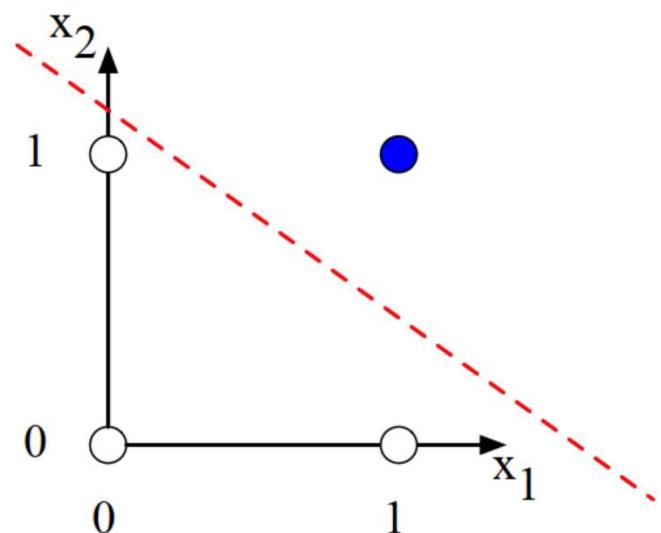


Sometimes **decision boundary CANNOT** be linear? Not linearly separable  $f()$

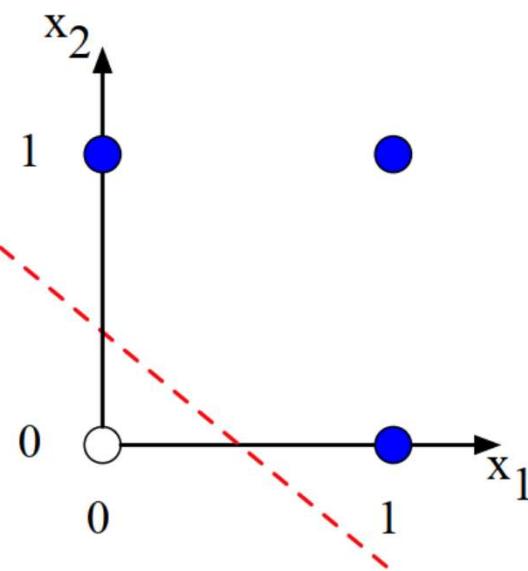
# Hypothesis: Classification “Boundary”



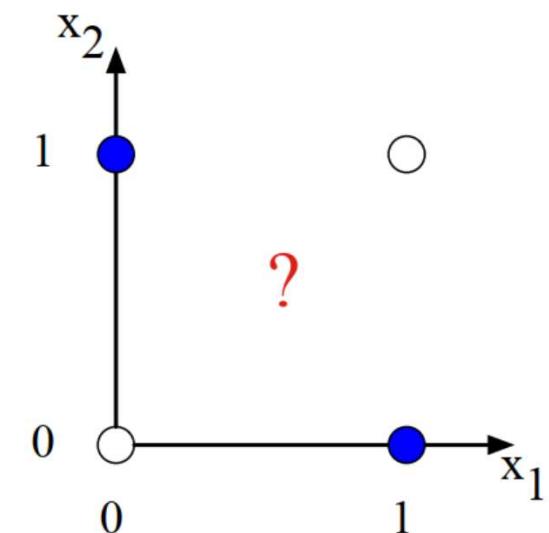
# XOR: Not a Linearly Separable f()



a)  $x_1$  AND  $x_2$



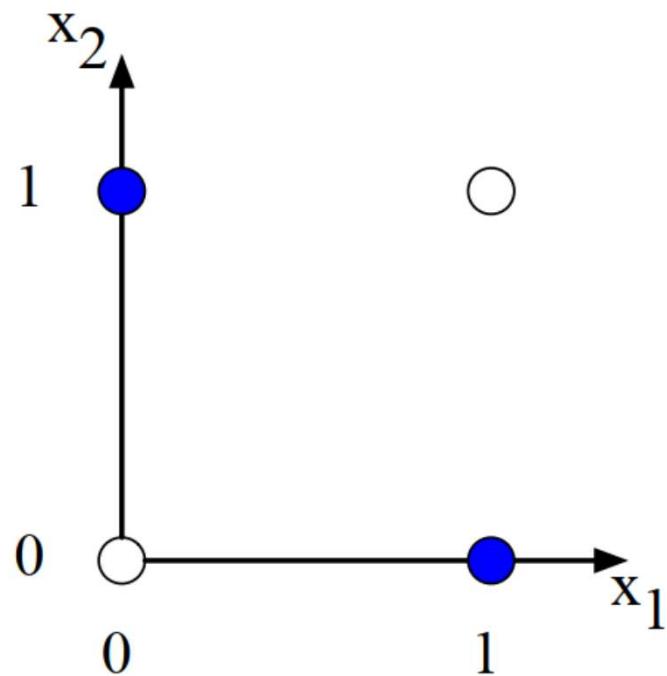
b)  $x_1$  OR  $x_2$



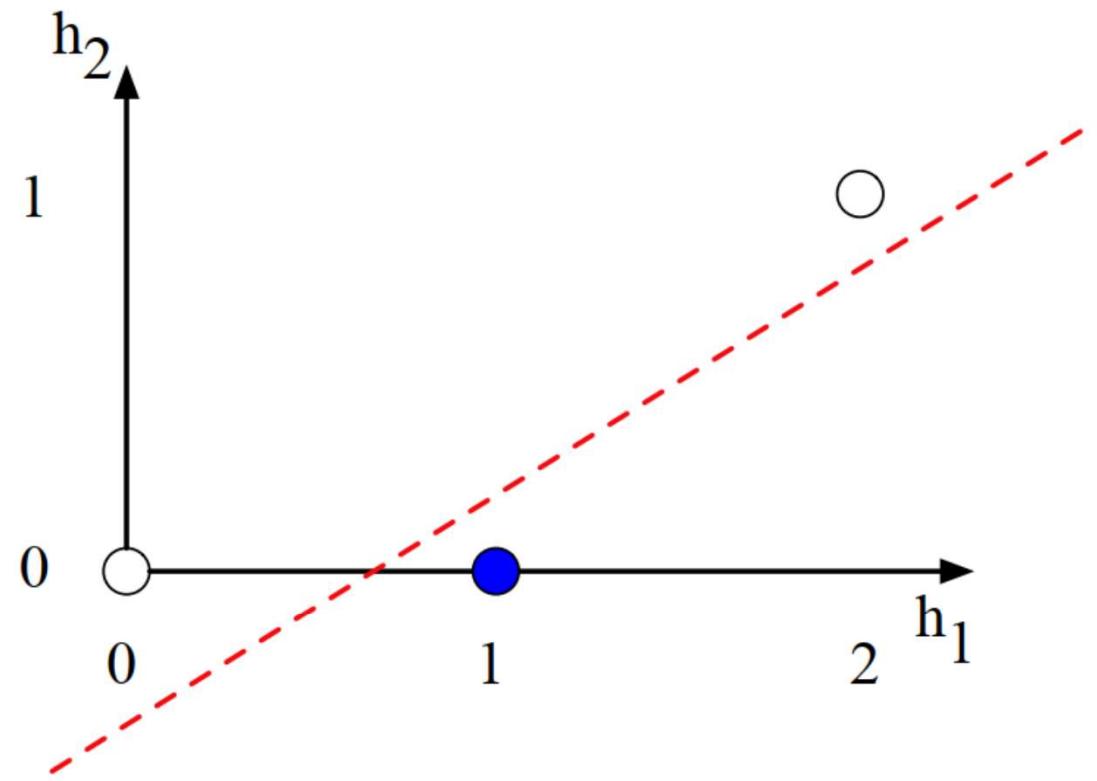
c)  $x_1$  XOR  $x_2$

Logical XOR is an example of a function that is **NOT linearly separable**

# XOR: Not a Linearly Separable $f()$ ?

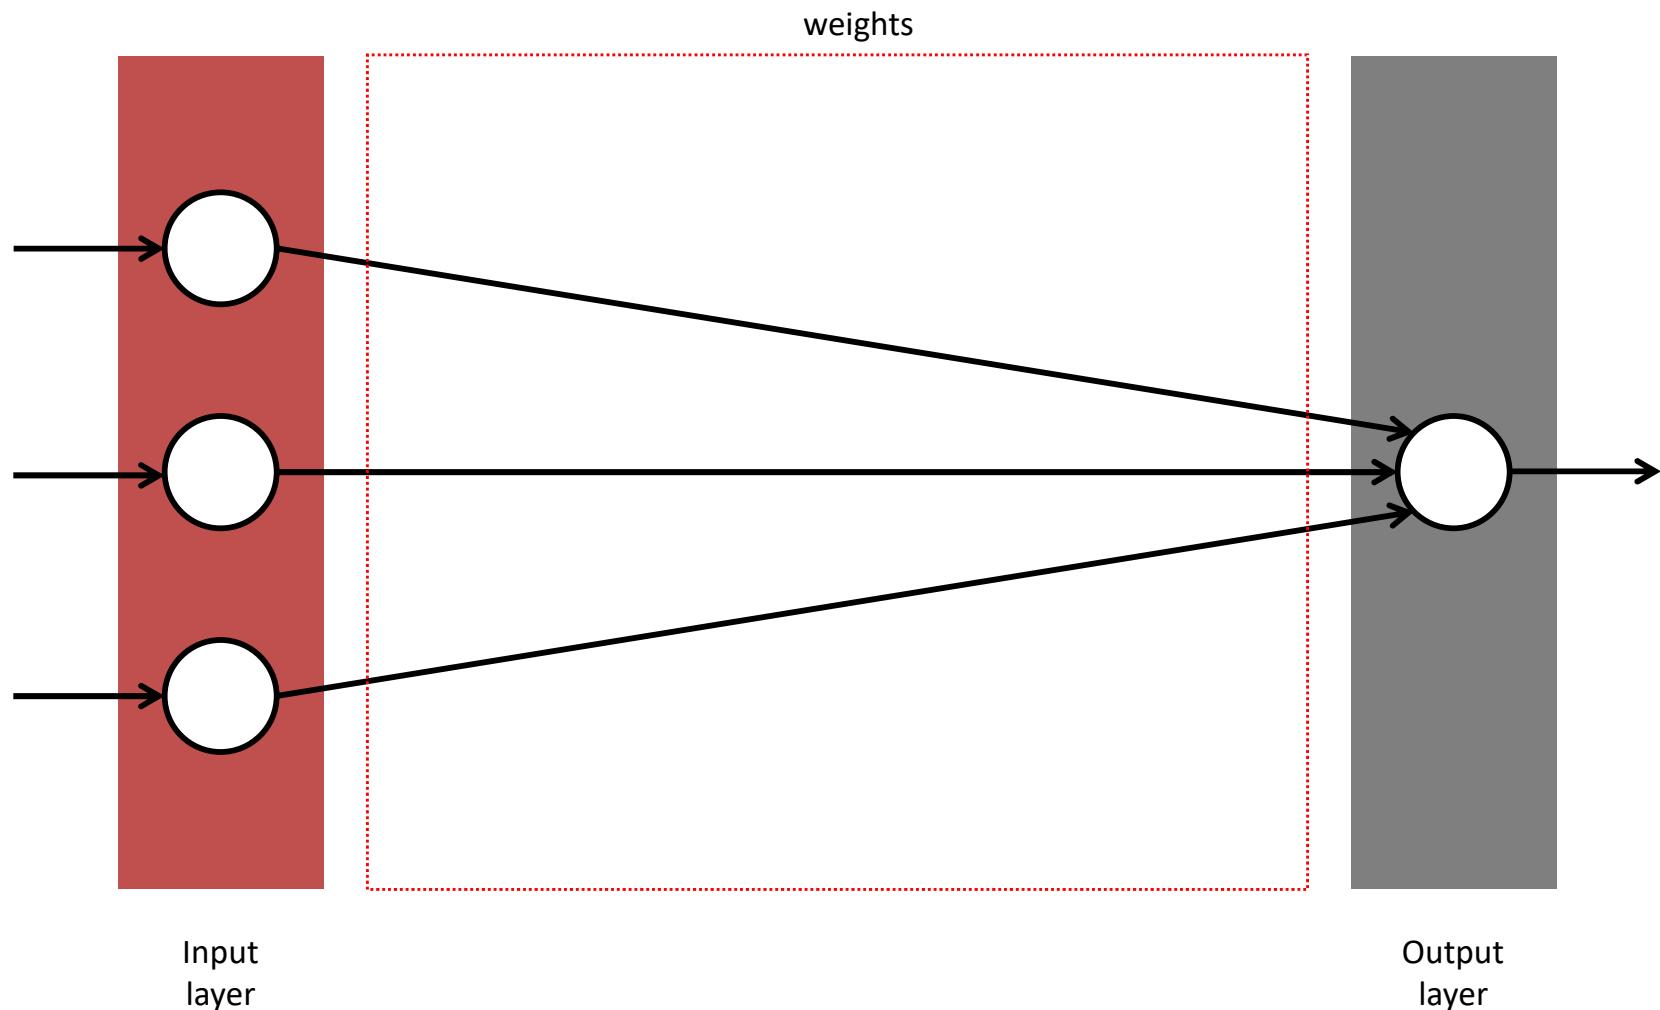


a) The original  $x$  space

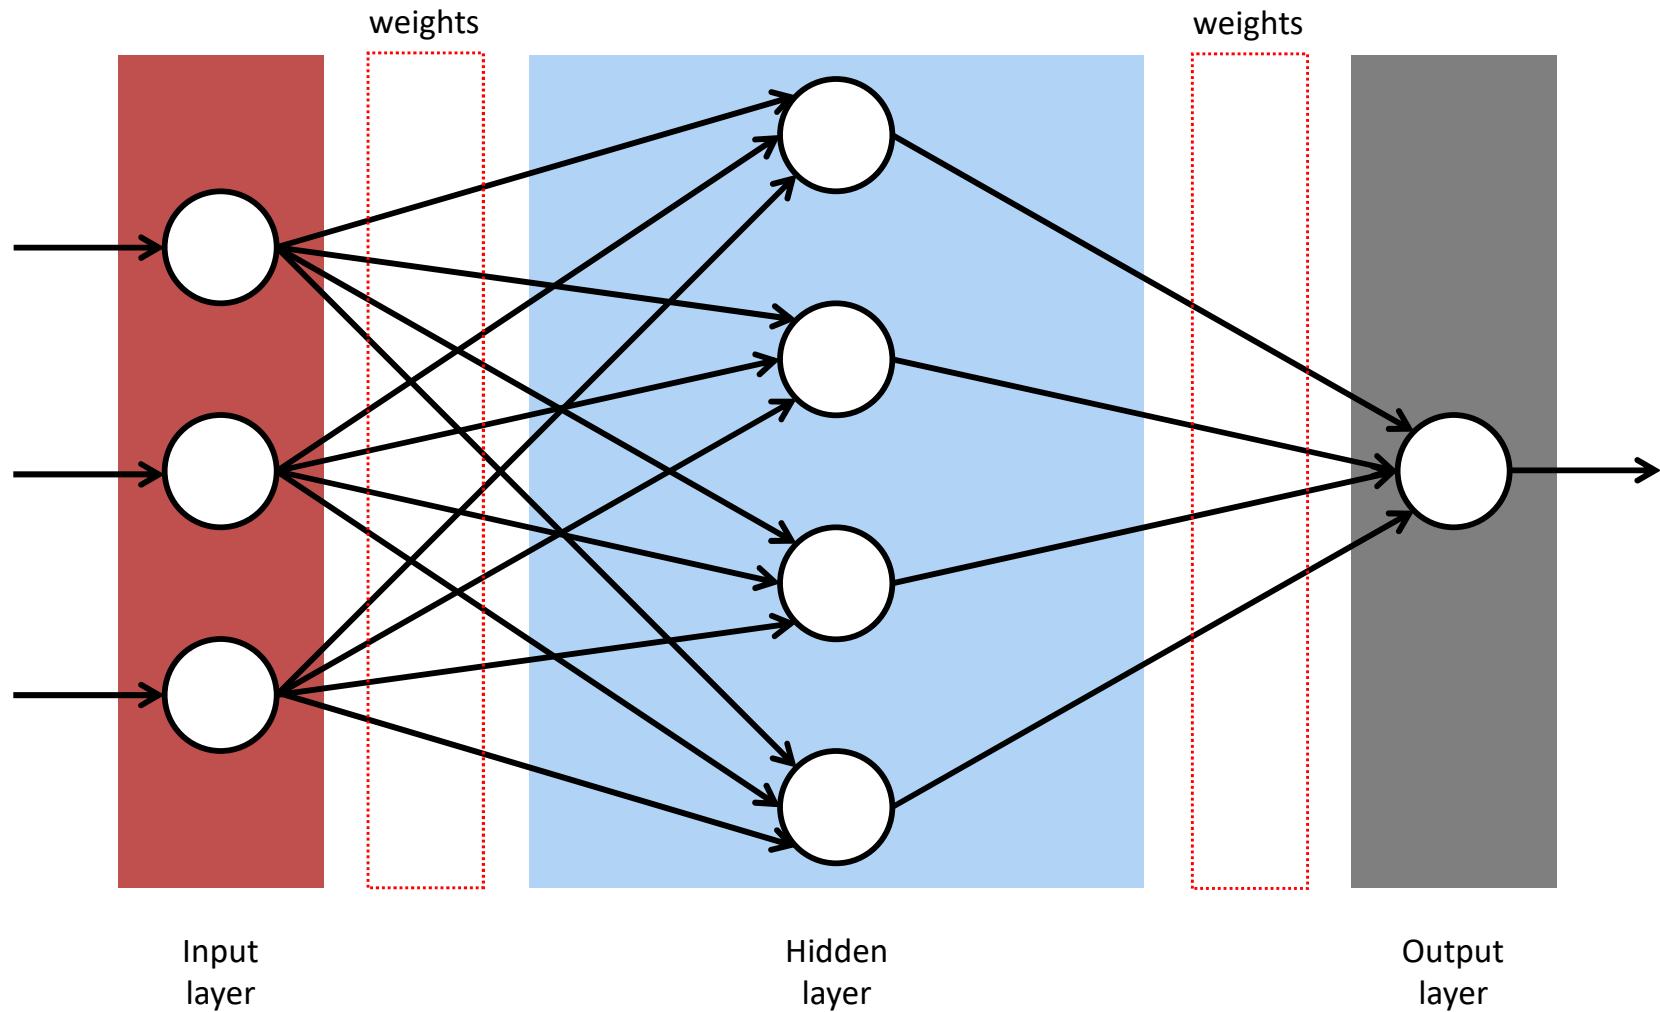


b) The new (linearly separable)  $h$  space

# Basic Neural Unit



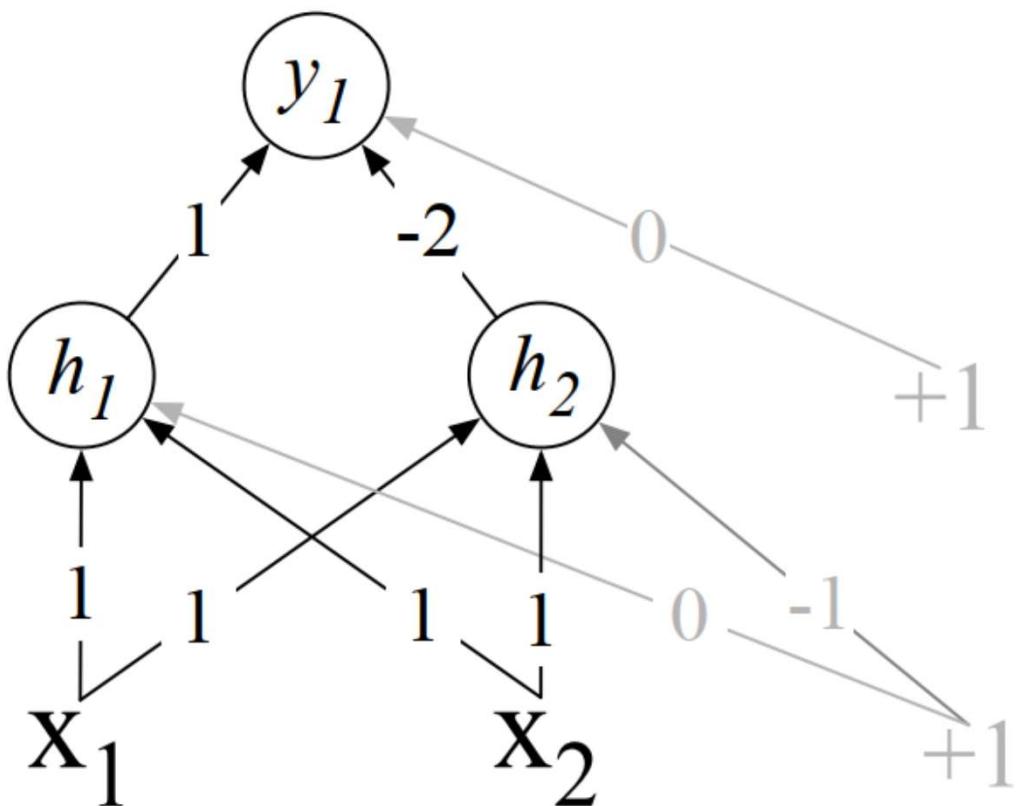
# Hidden Layer



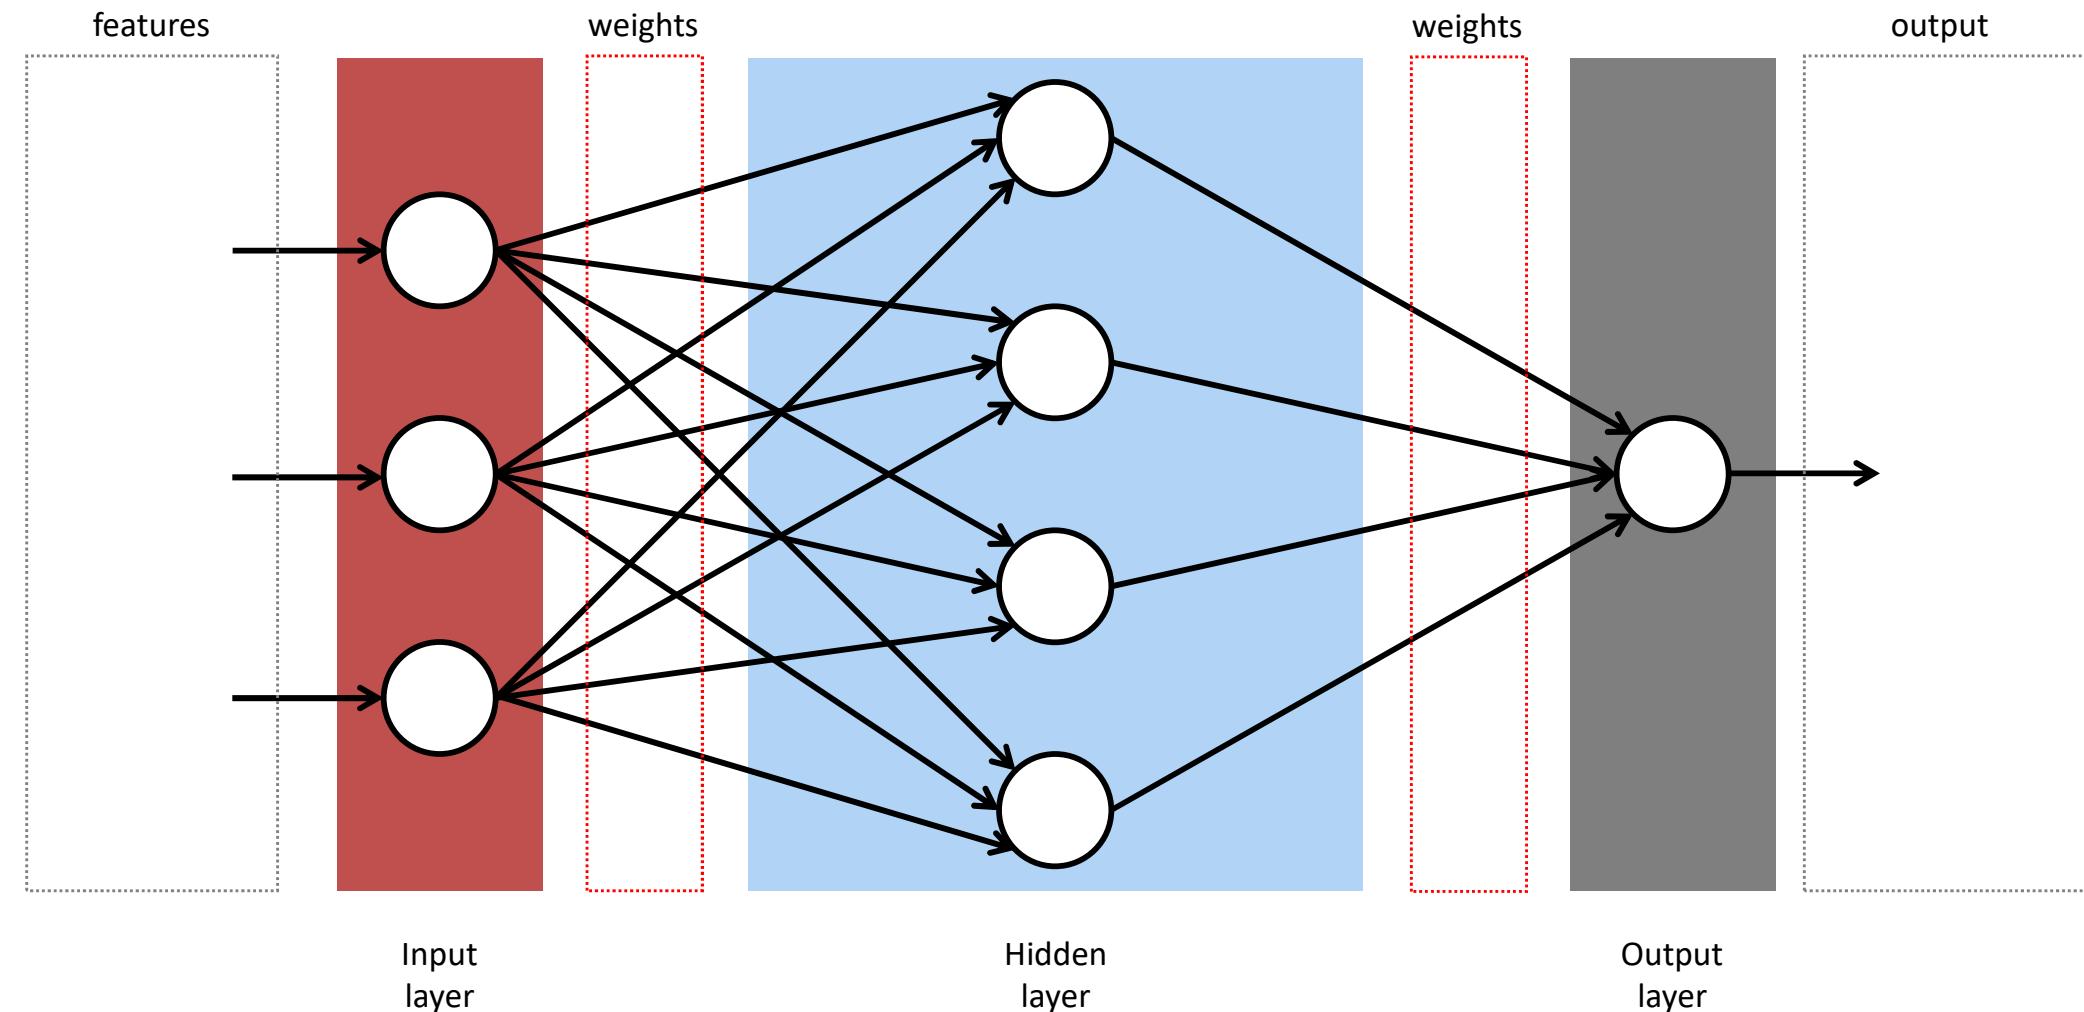
# XOR: Hidden Layer Approach

XOR

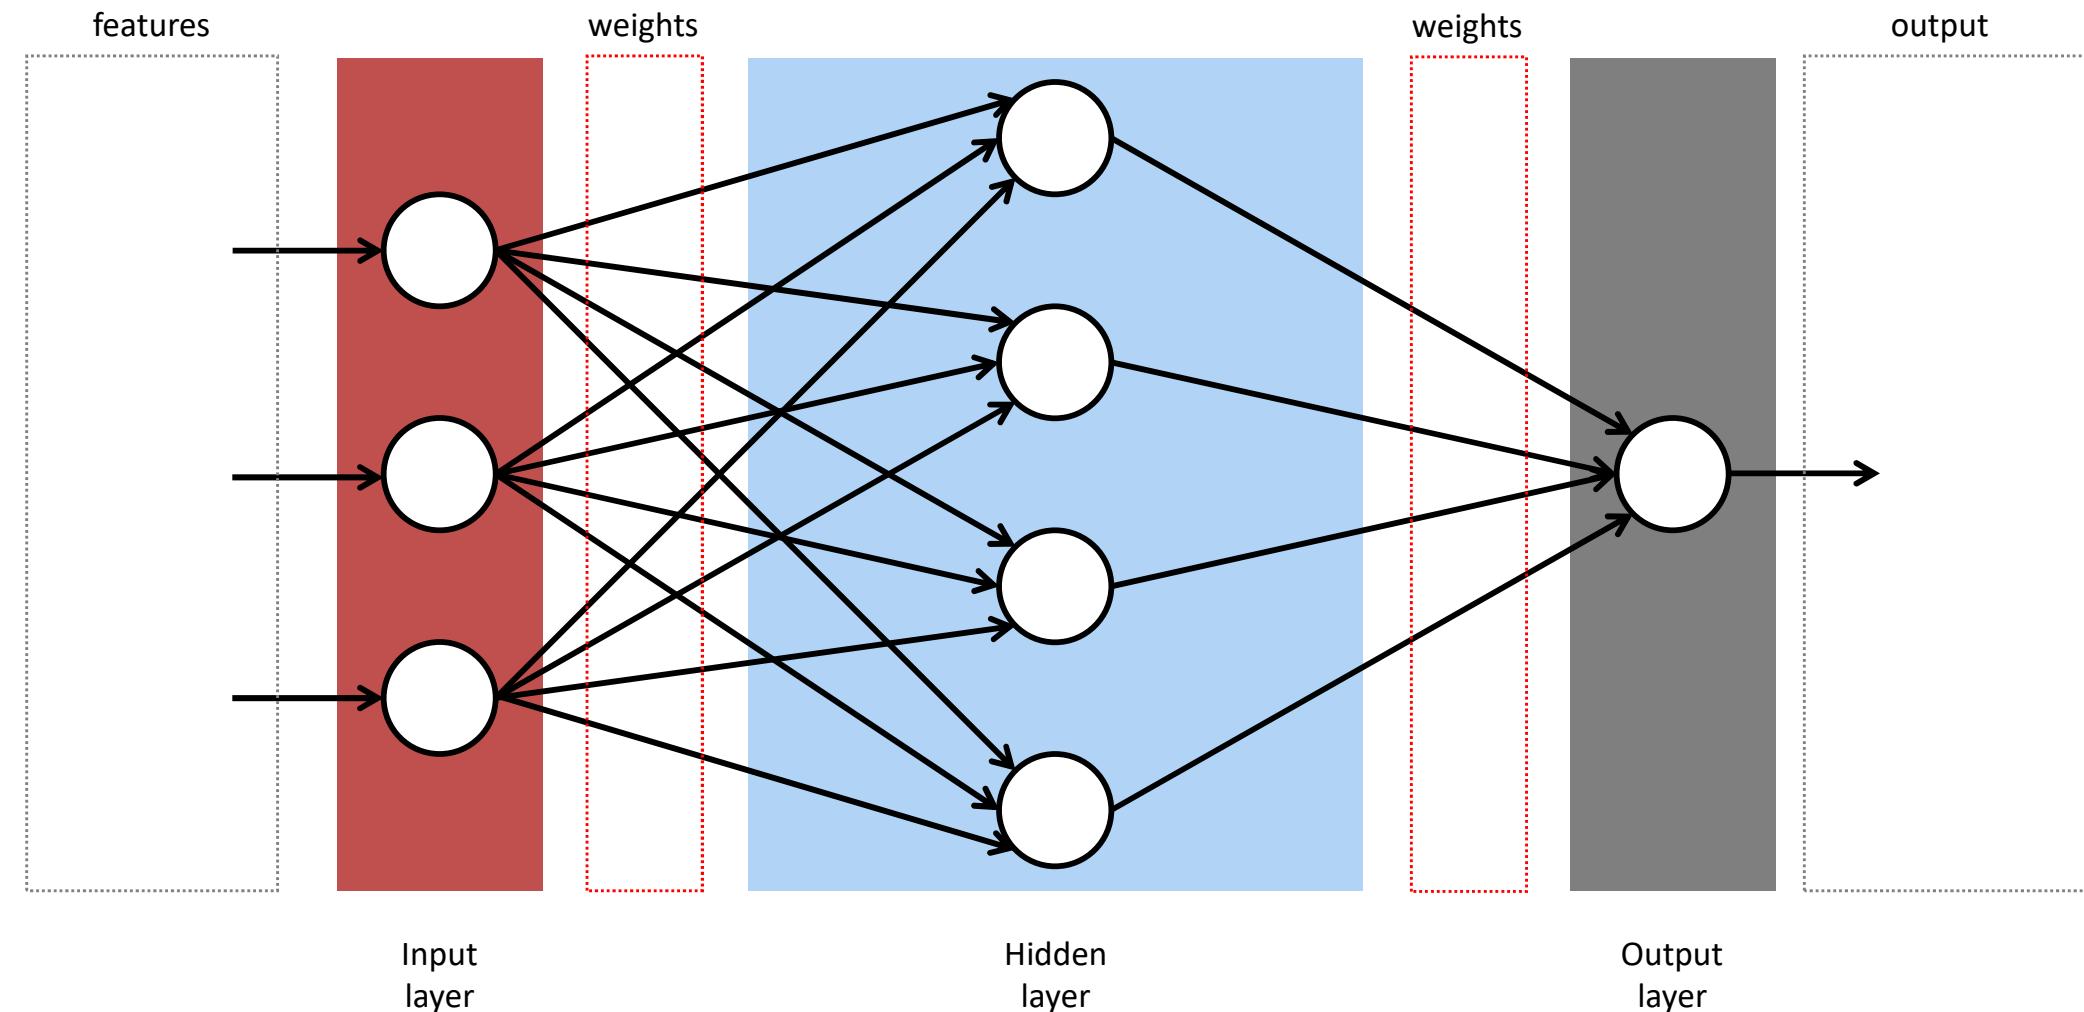
x1	x2	y
0	0	0
0	1	1
1	0	1
1	1	0



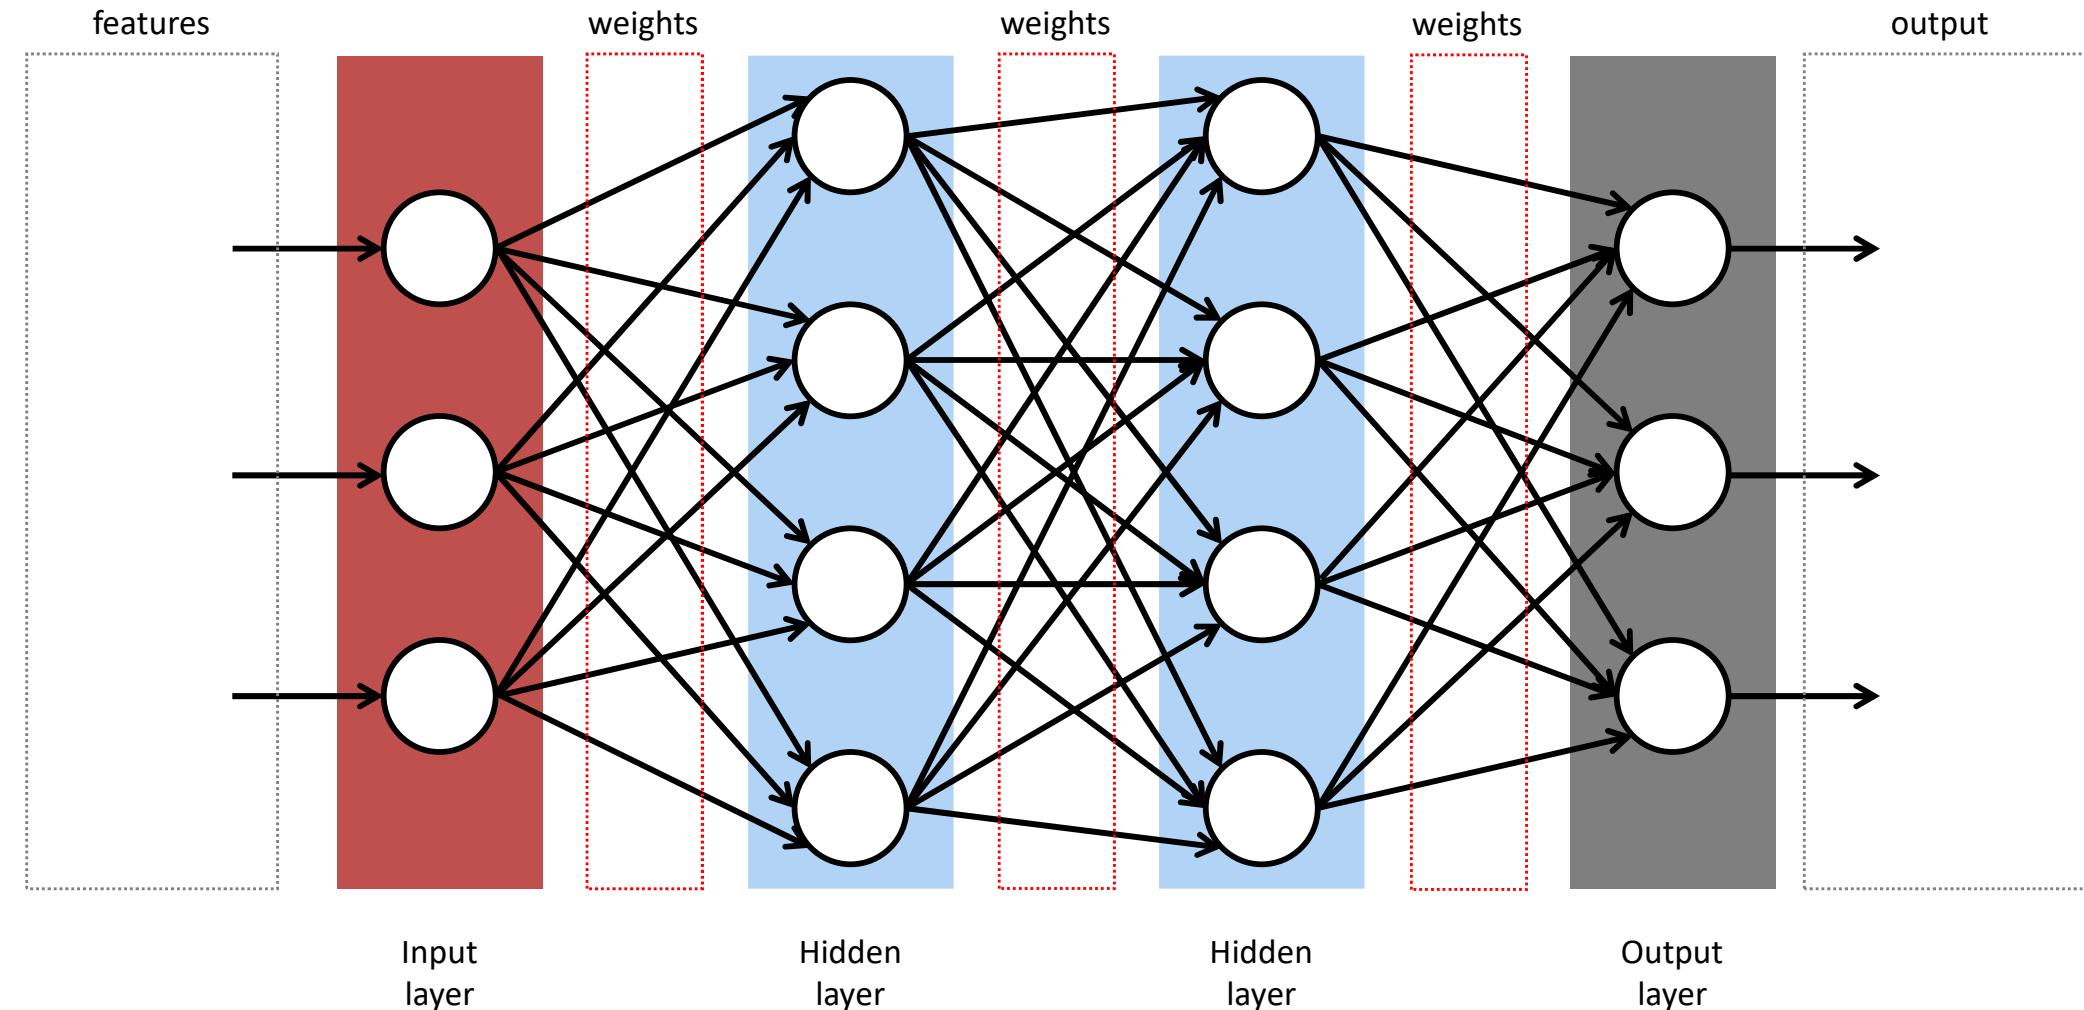
# Hidden Layer



# Hidden Layer



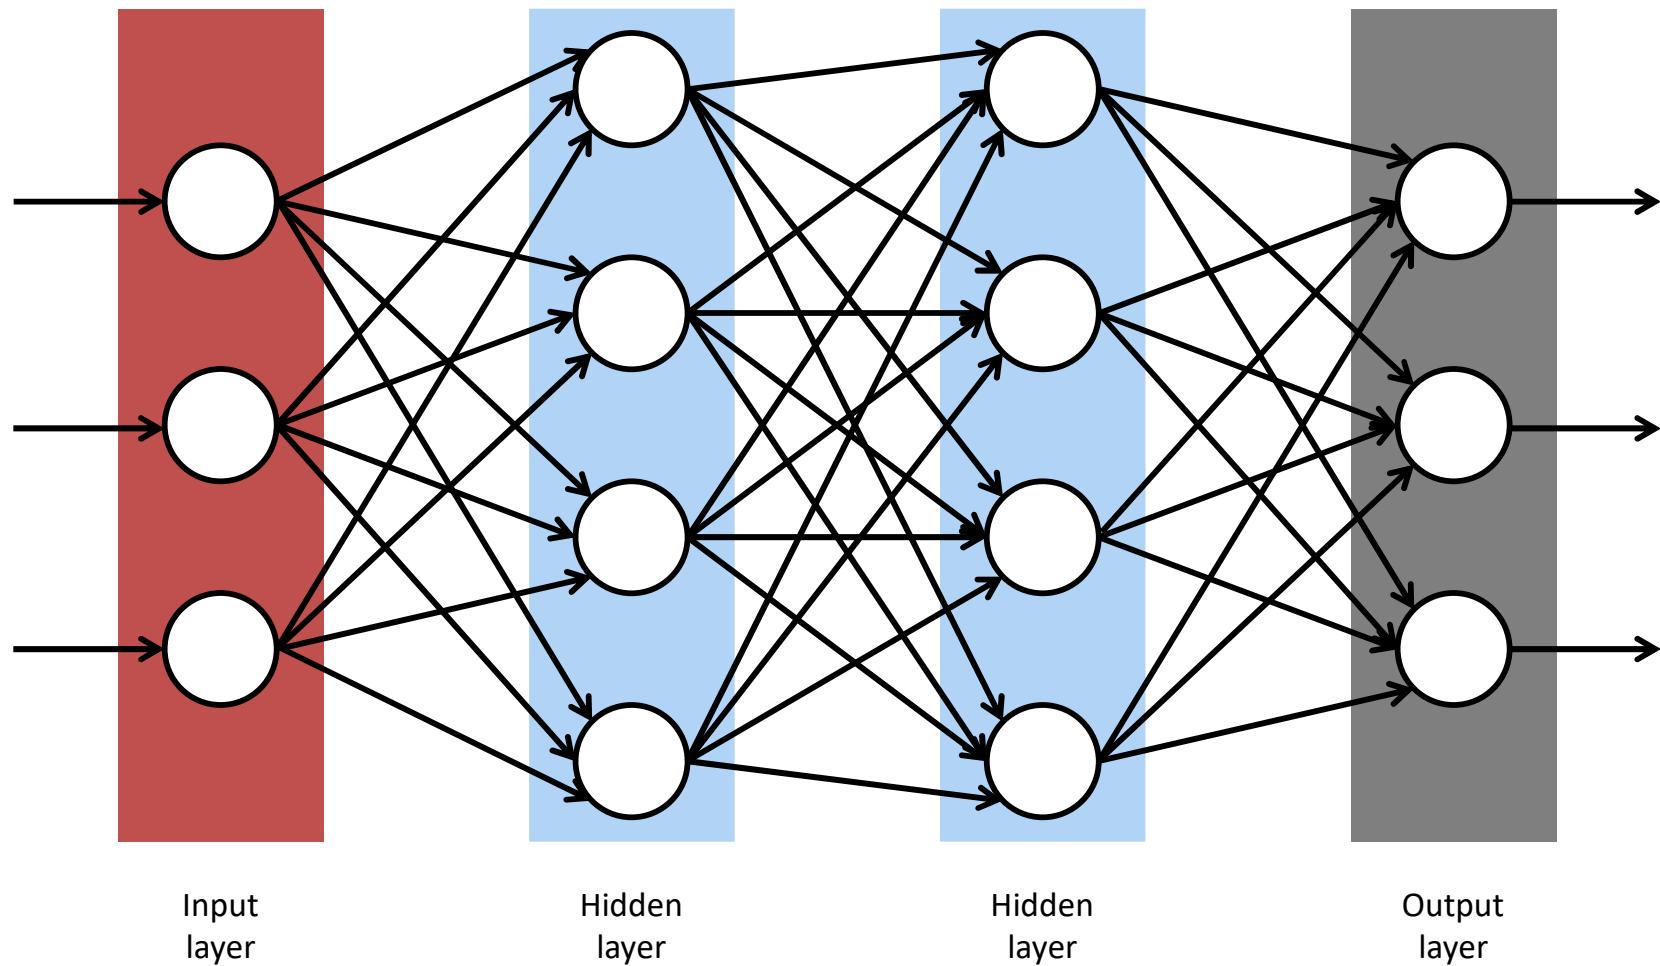
# Feedforward Neural Network



Also called (historically): **multi-layer perceptron**

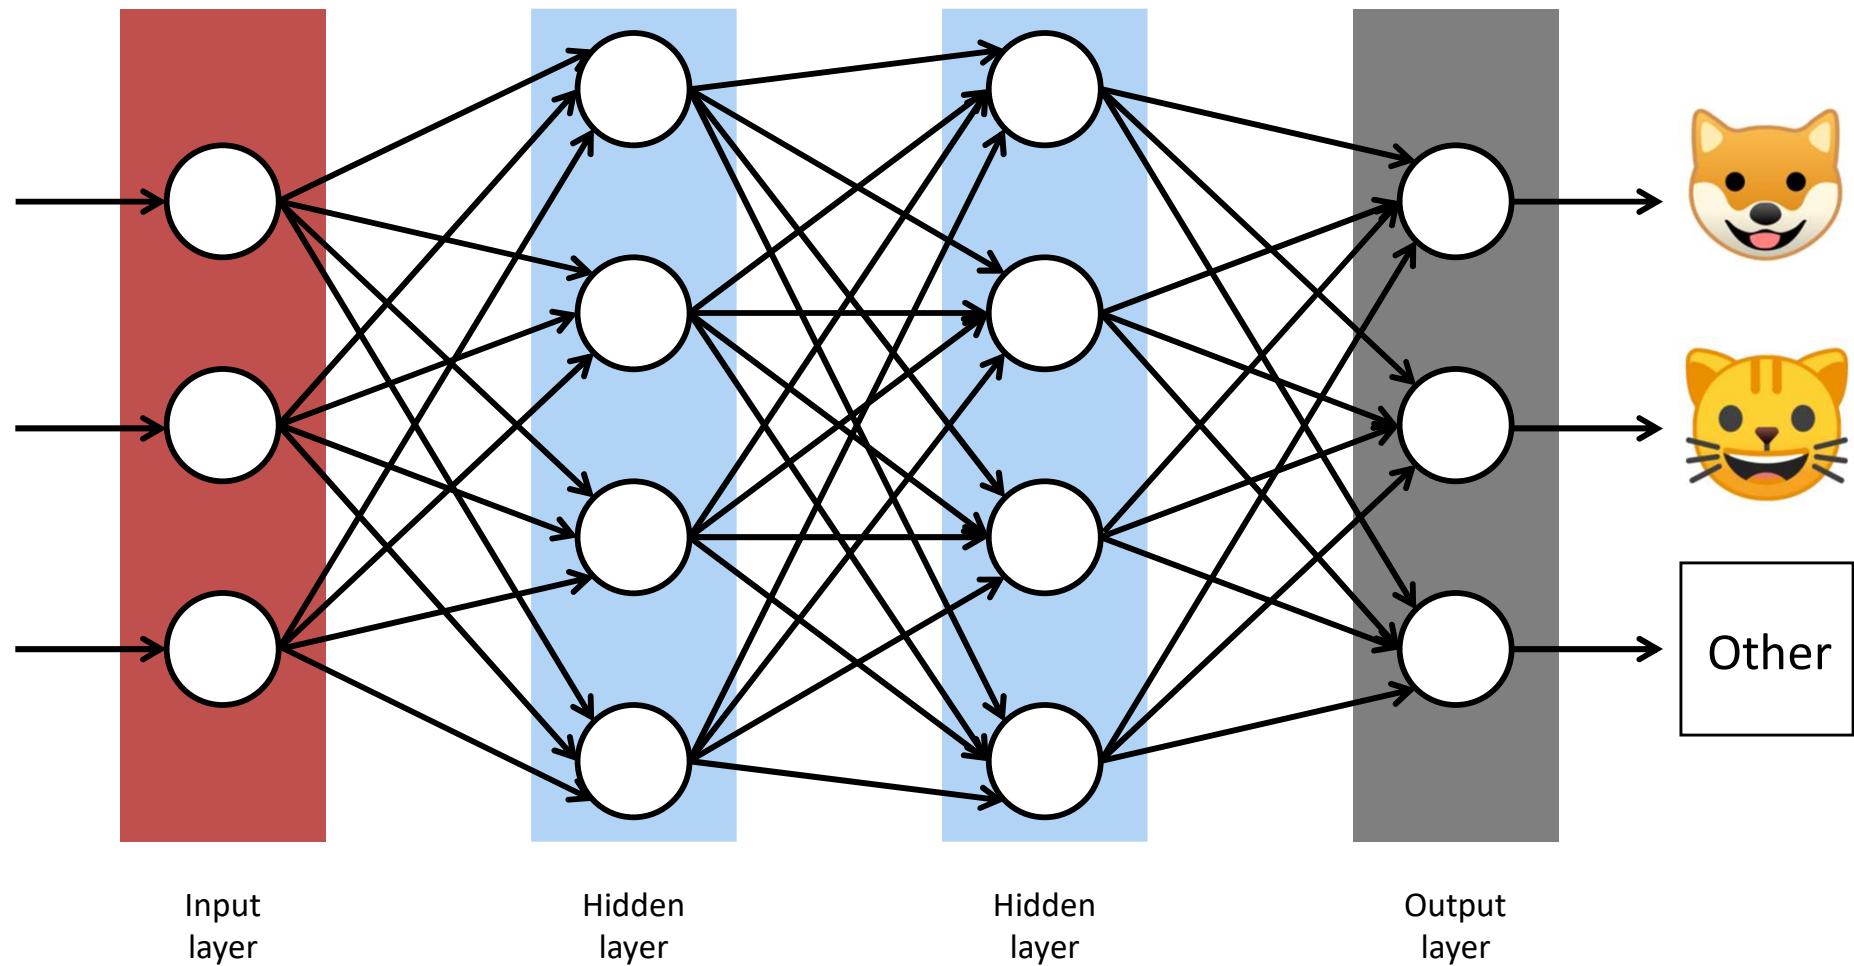
# Artificial Neural Network (ANN)

An artificial neural network is made of **multiple artificial neuron layers**.

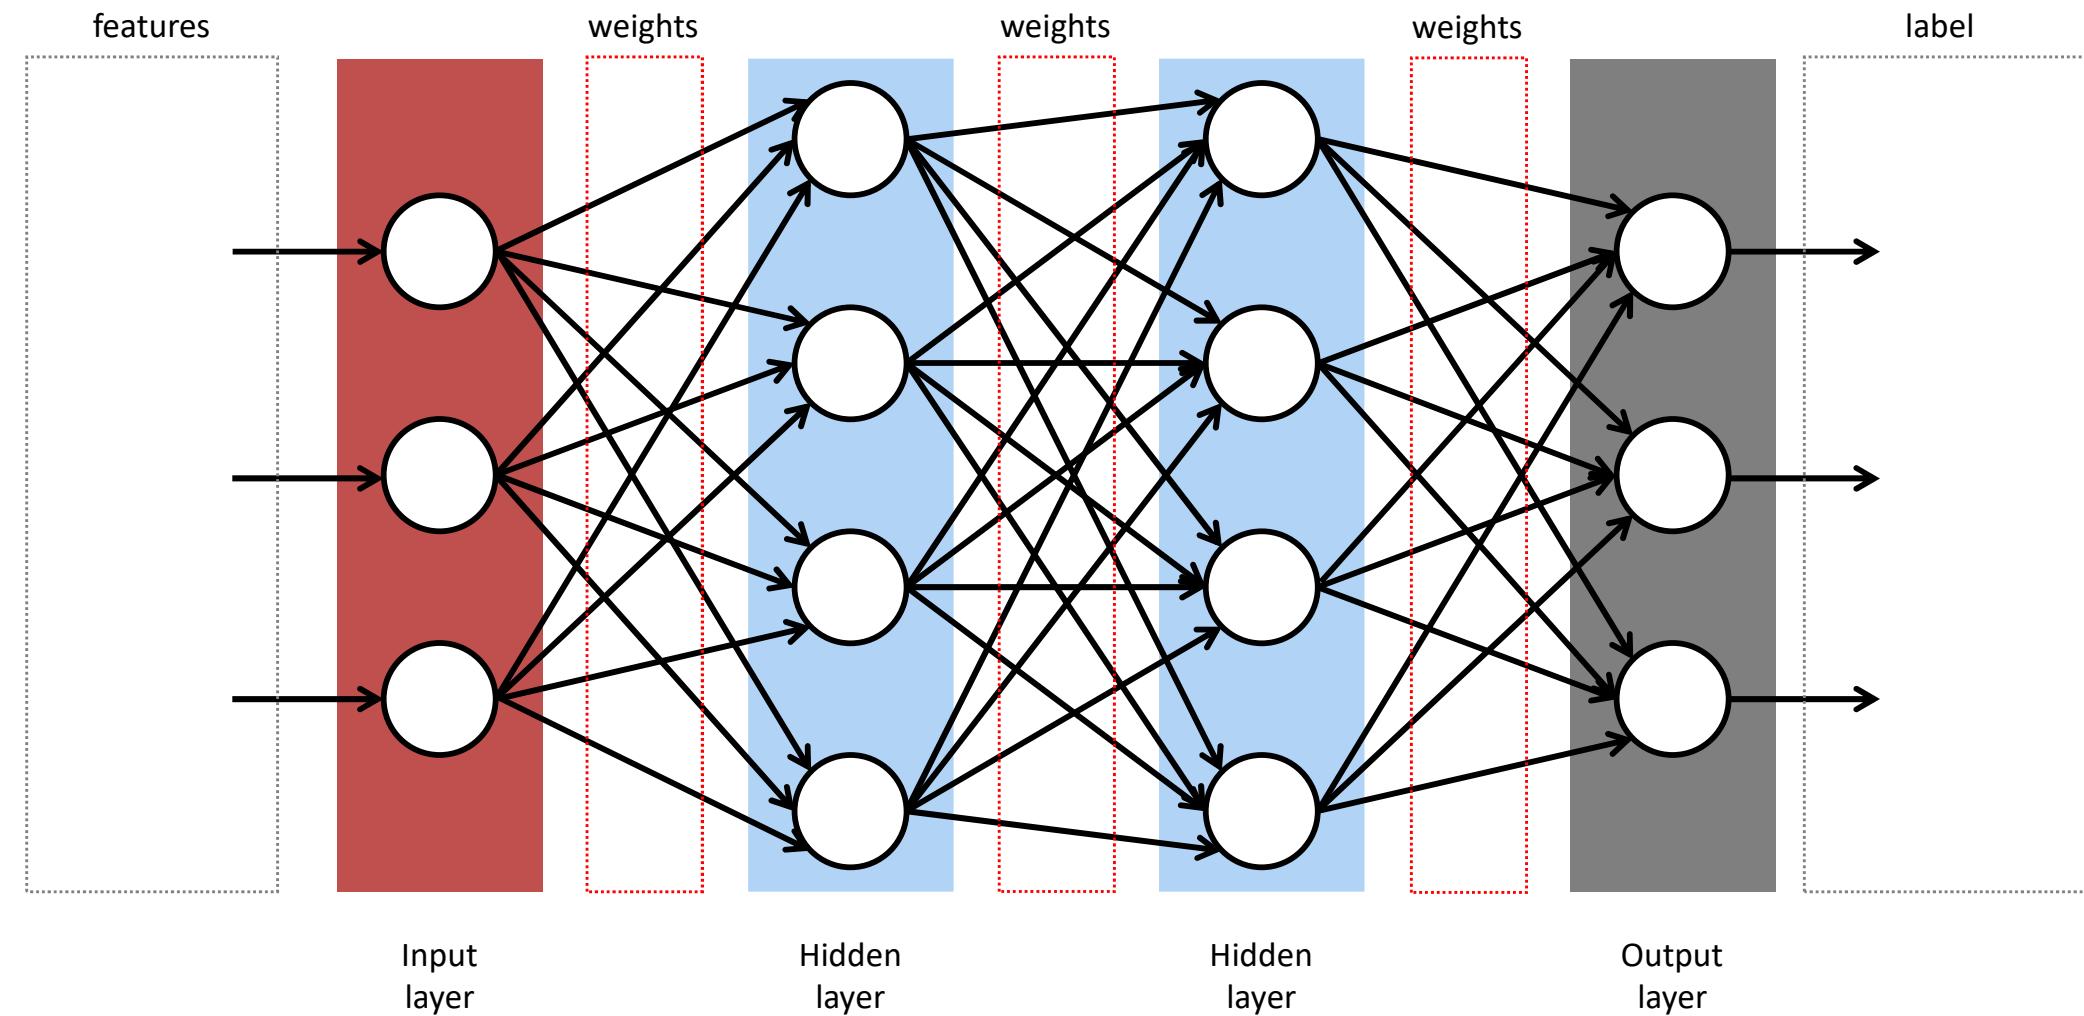


# ANN as an Image Classifier

An artificial neural network can be used as a **classifier** as well.

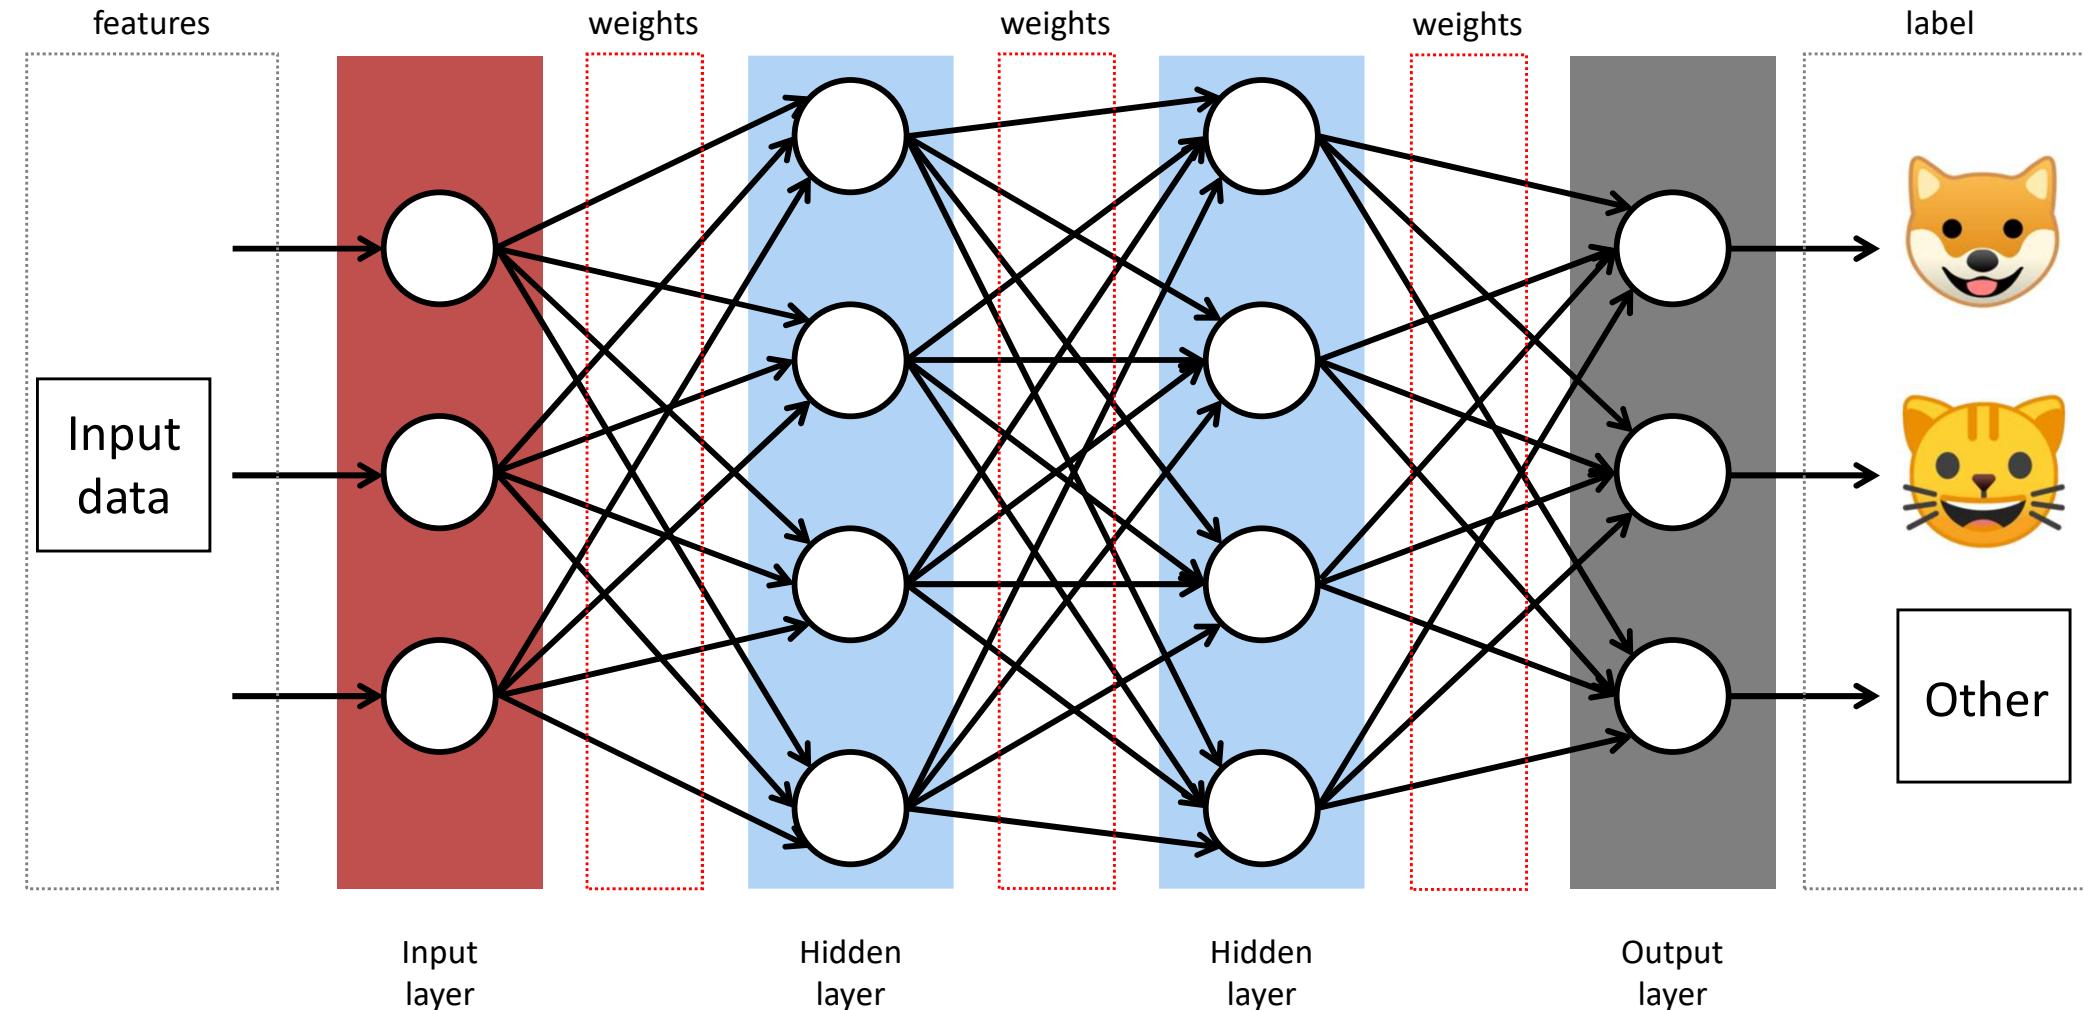


# ANN as a Classifier



# ANN: Supervised Learning

In order to work properly a classifier **needs to be trained** first with **labeled data**.



**Training** will **adjust all the weights** within this artificial neural network.

# Training Data: Features + Labels

Typically input data will be represented by a **limited set of features**.



Features:  
Wheels: 4  
Weight: 8 tons  
Passengers: 1

Label:  
**Truck**



Features:  
Wheels: 6  
Weight: 8 tons  
Passengers: 1

Label:  
**Truck**



Features:  
Wheels: 4  
Weight: 1 ton  
Passengers: 4

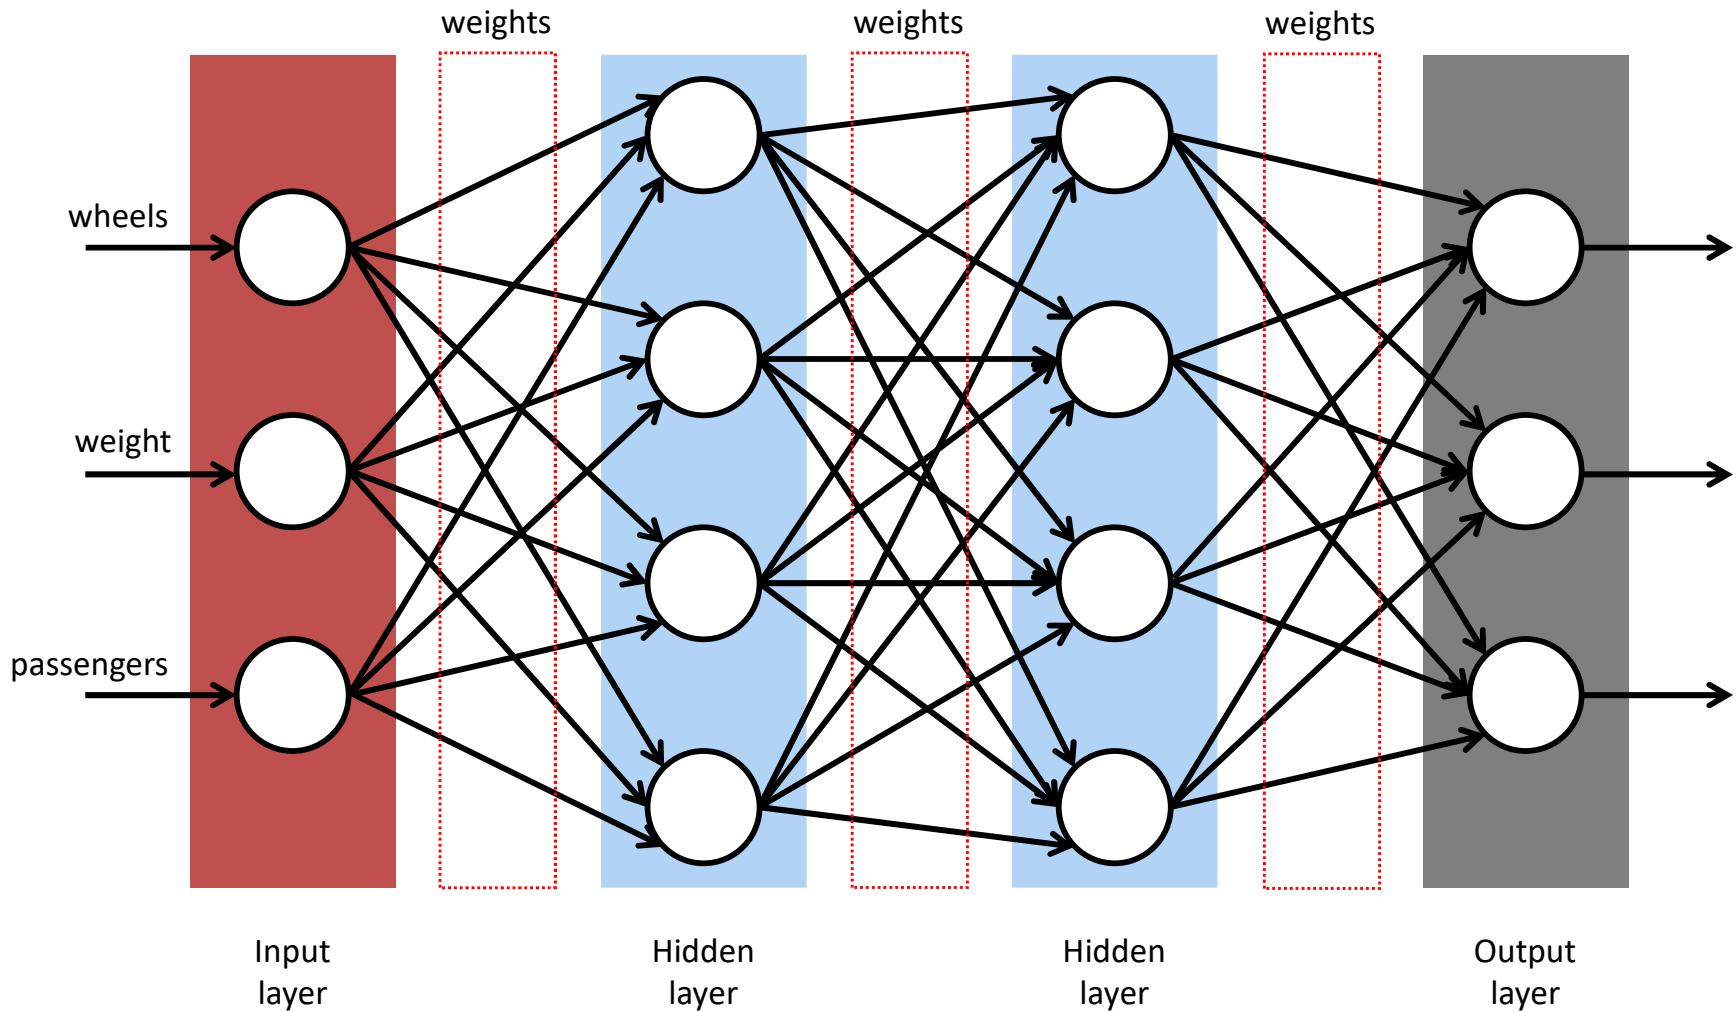
Label:  
**Car**



Features:  
Wheels: 4  
Weight: 2 tons  
Passengers: 4

Label:  
**Car**

# ANN: Supervised Learning



# Training Data: Images + Labels

A classifier **needs to be “shown” thousands of labeled examples to learn.**



Label:  
**BUS**



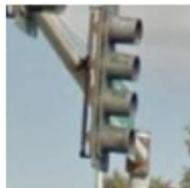
Label:  
**CAR**



Label:  
**BRIDGE**



Label:  
**PALM**



Label:  
**TRAFFIC LIGHT**



Label:  
**TAXI**



Label:  
**CROSSWALK**



Label:  
**CHIMNEY**



Label:  
**MOTORCYCLE**



Label:  
**STREET SIGN**



Label:  
**HYDRANT**



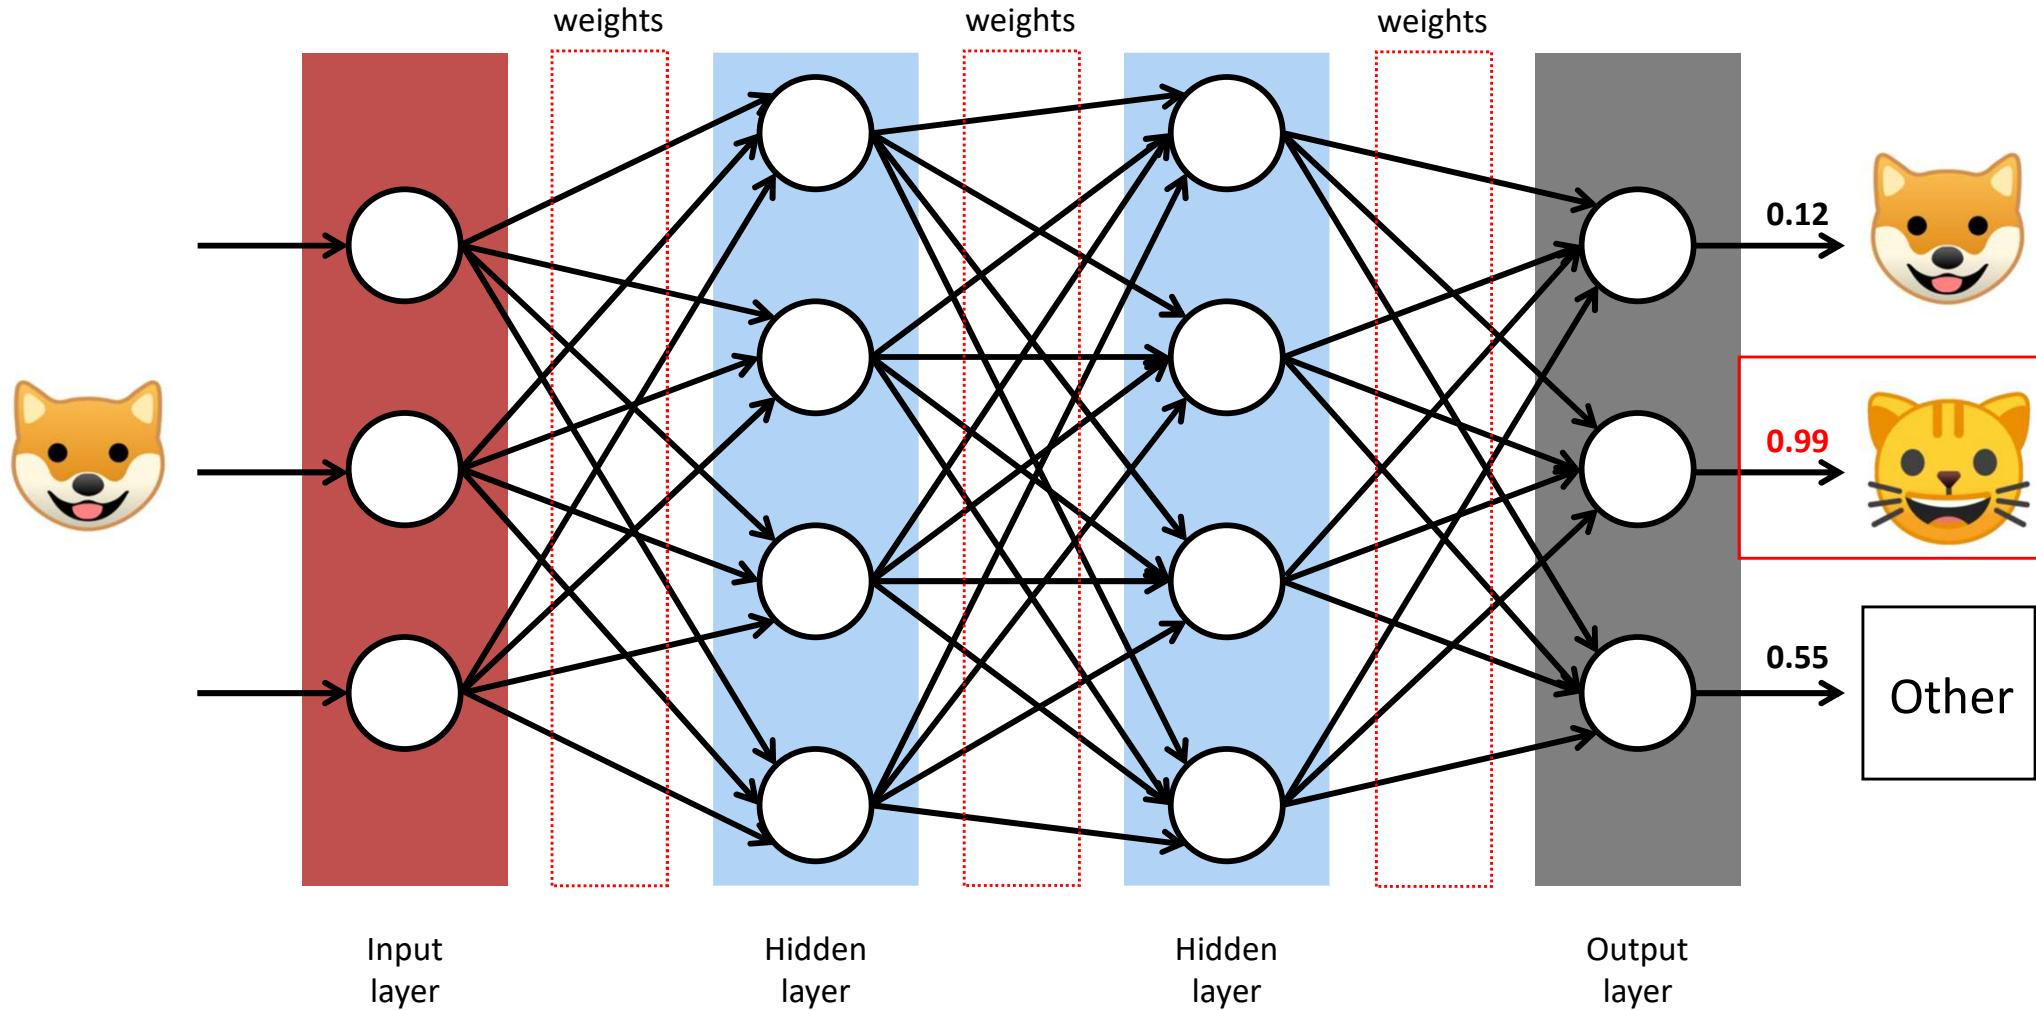
Label:  
**BICYCLE**

*Note how some images are “incomplete” and “flawed”.*



# ANN: Supervised Learning

An **untrained classifier** will **NOT** label input data correctly.

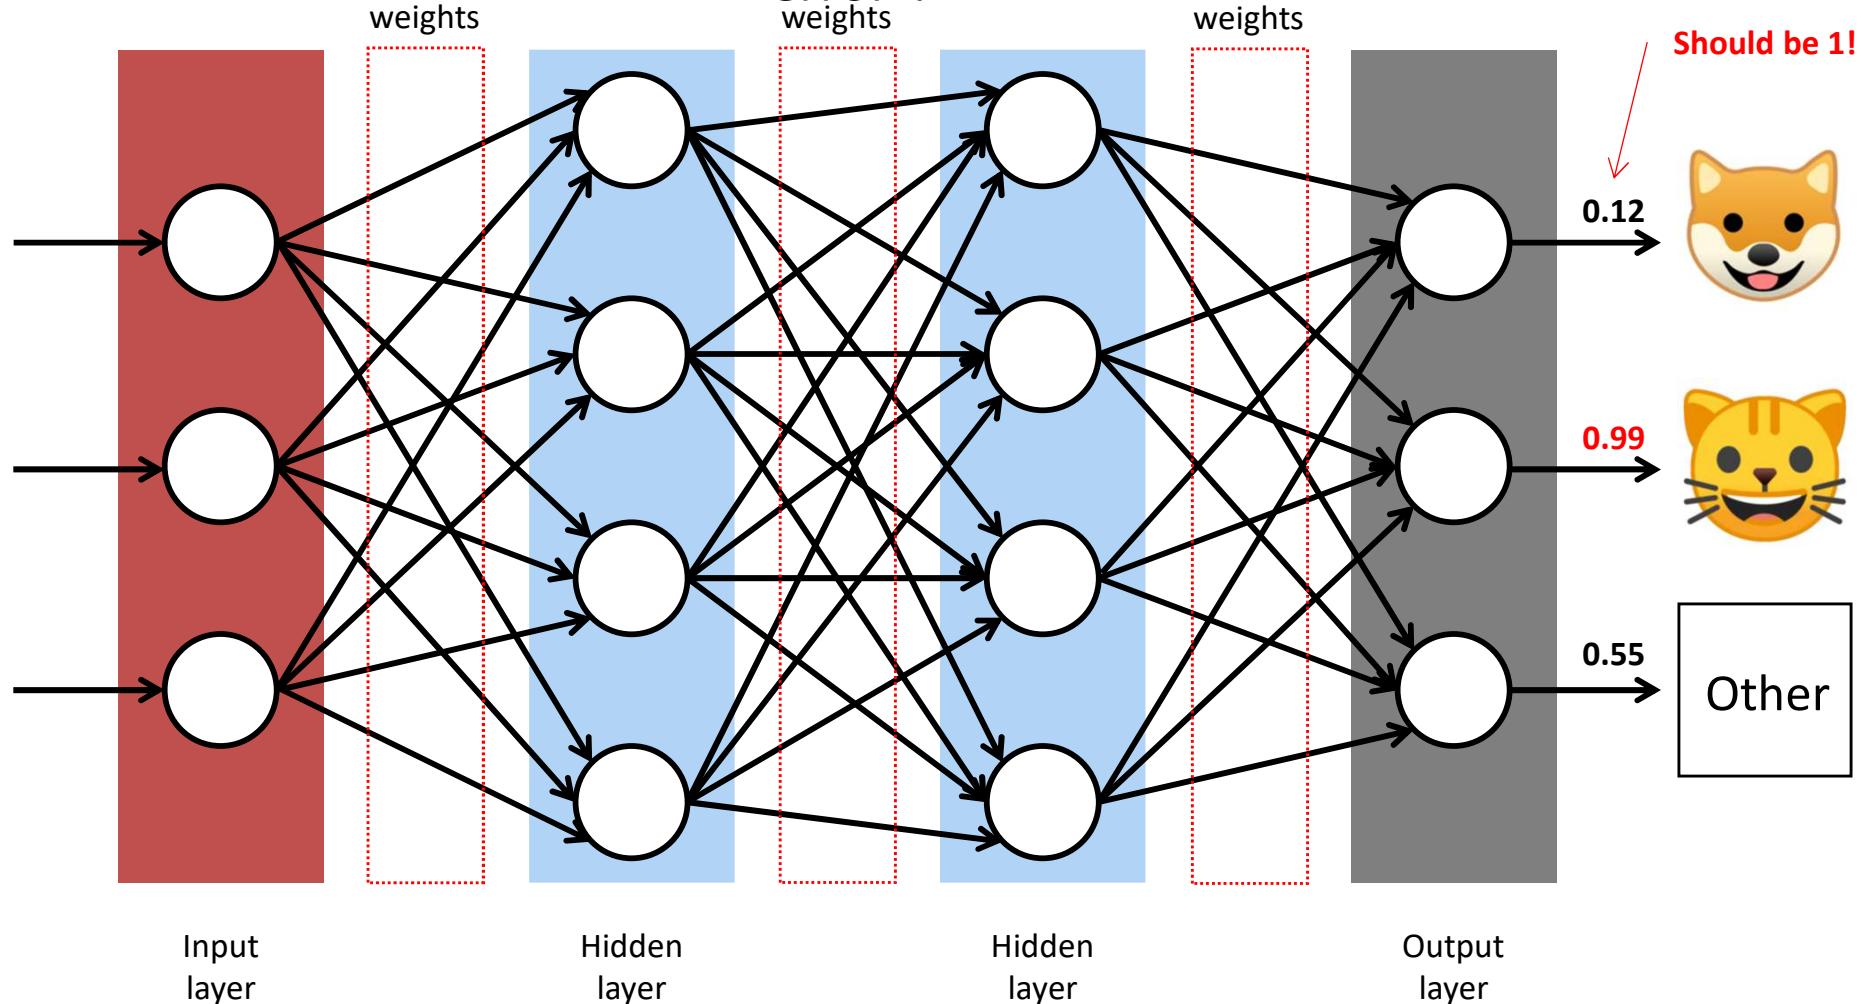


# ANN: Training

Given: input d 

and it's corresponding **expected** label: DOG calculate

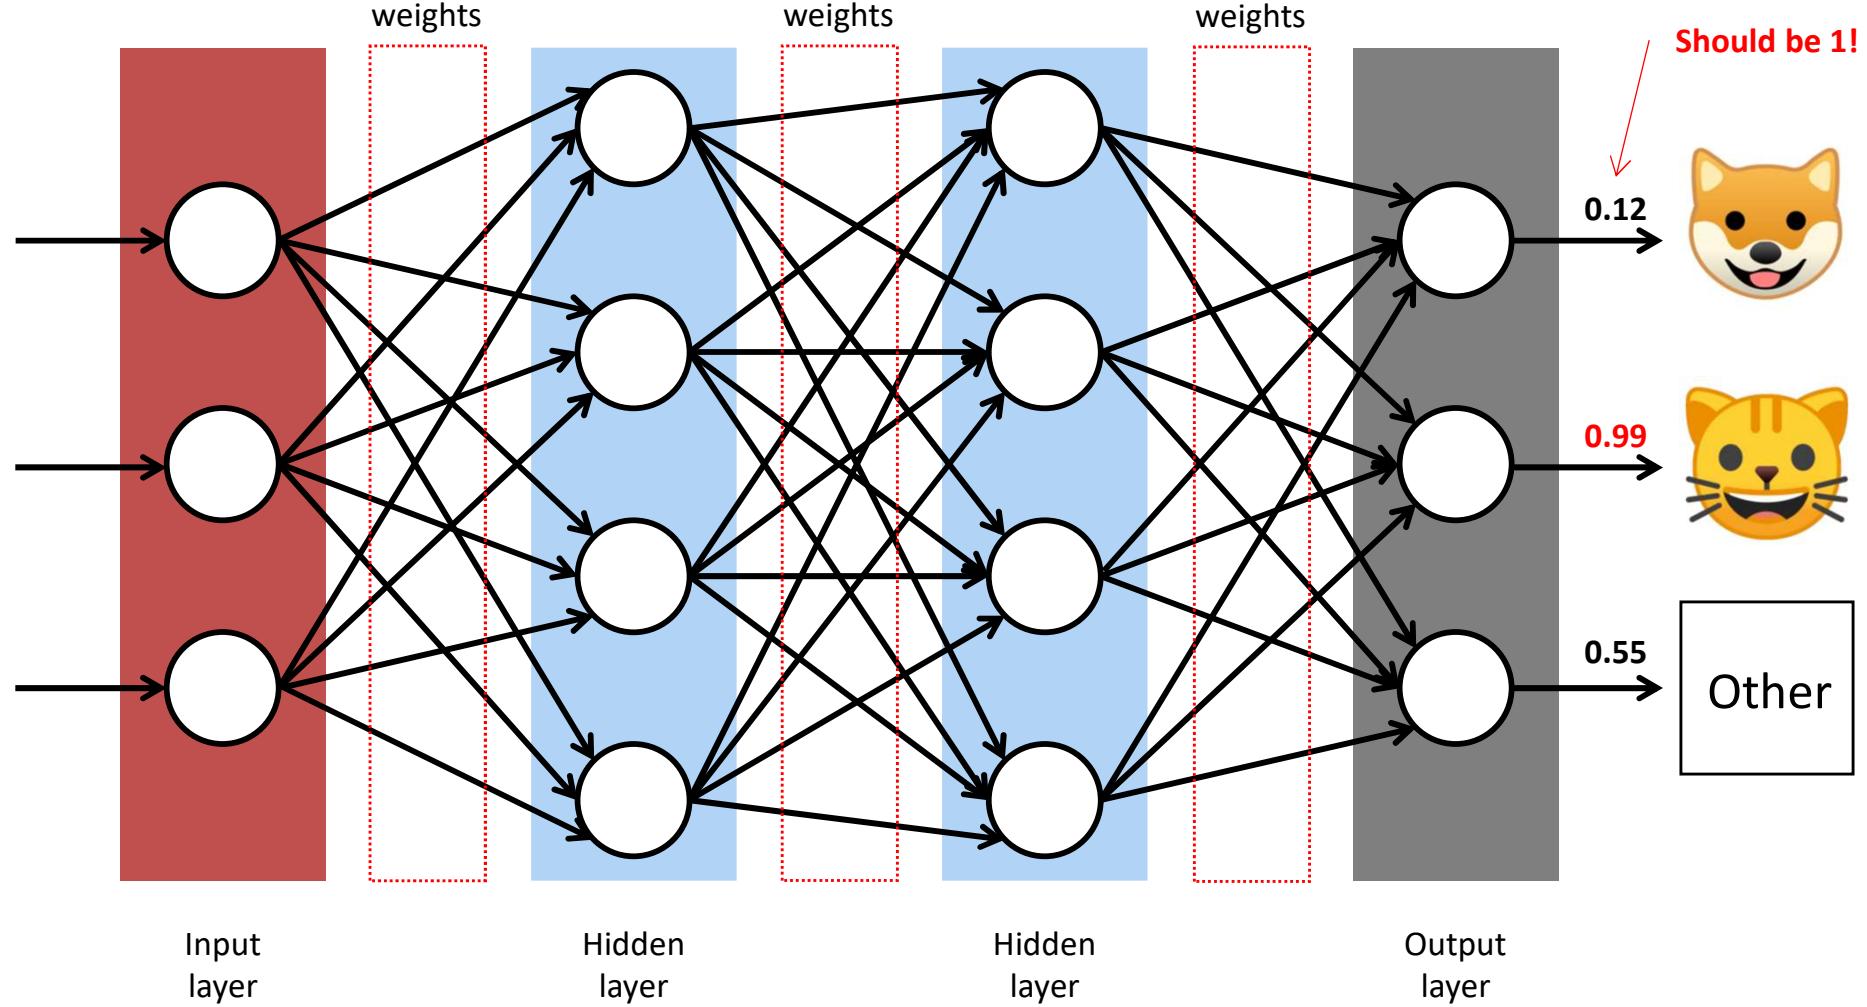
"error".



"Error" = 0.88. Go back and **adjust all the weights** to ensure it is lower next time.

# ANN: Training

Show data / label pair:  / DOG.



← **Correct all the weights.** Repeat many times.

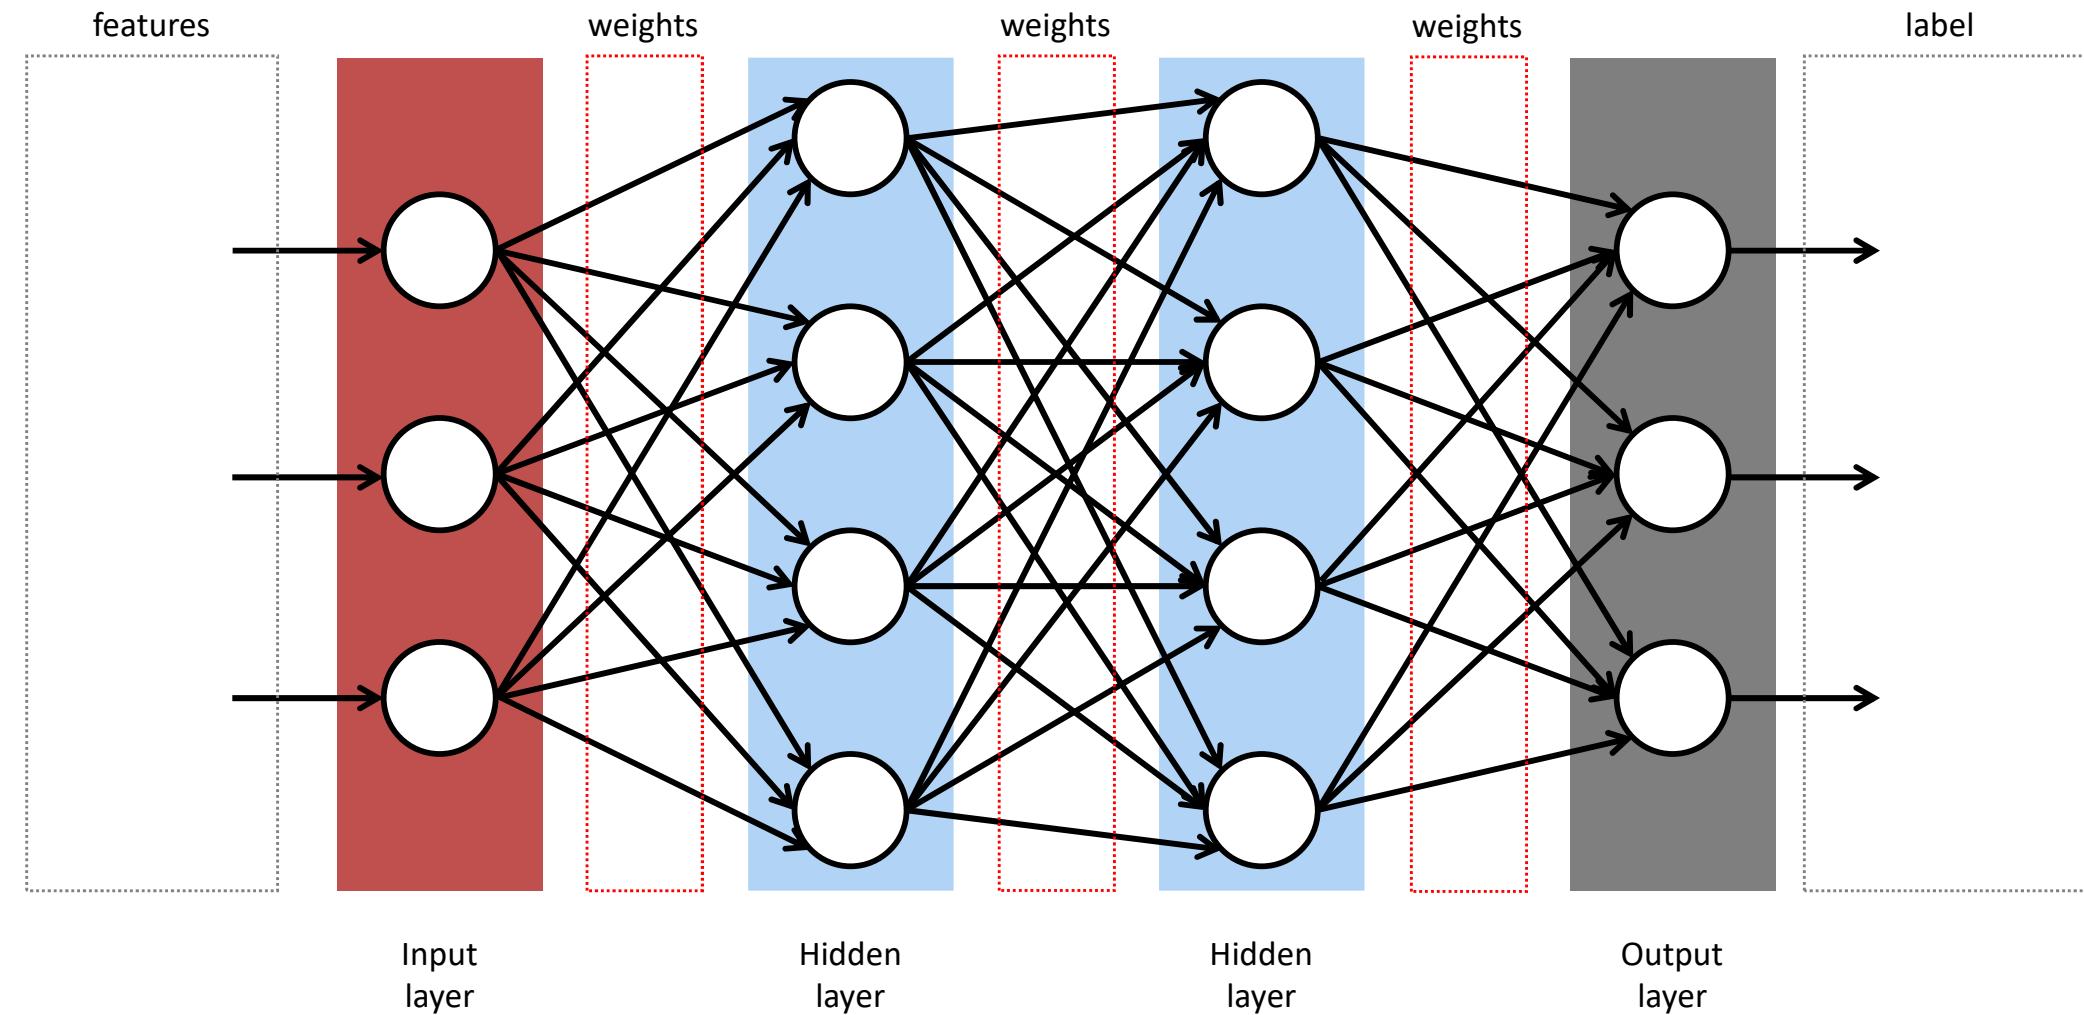
# **Exercise: ANN Demo**

**<http://playground.tensorflow.org/>**

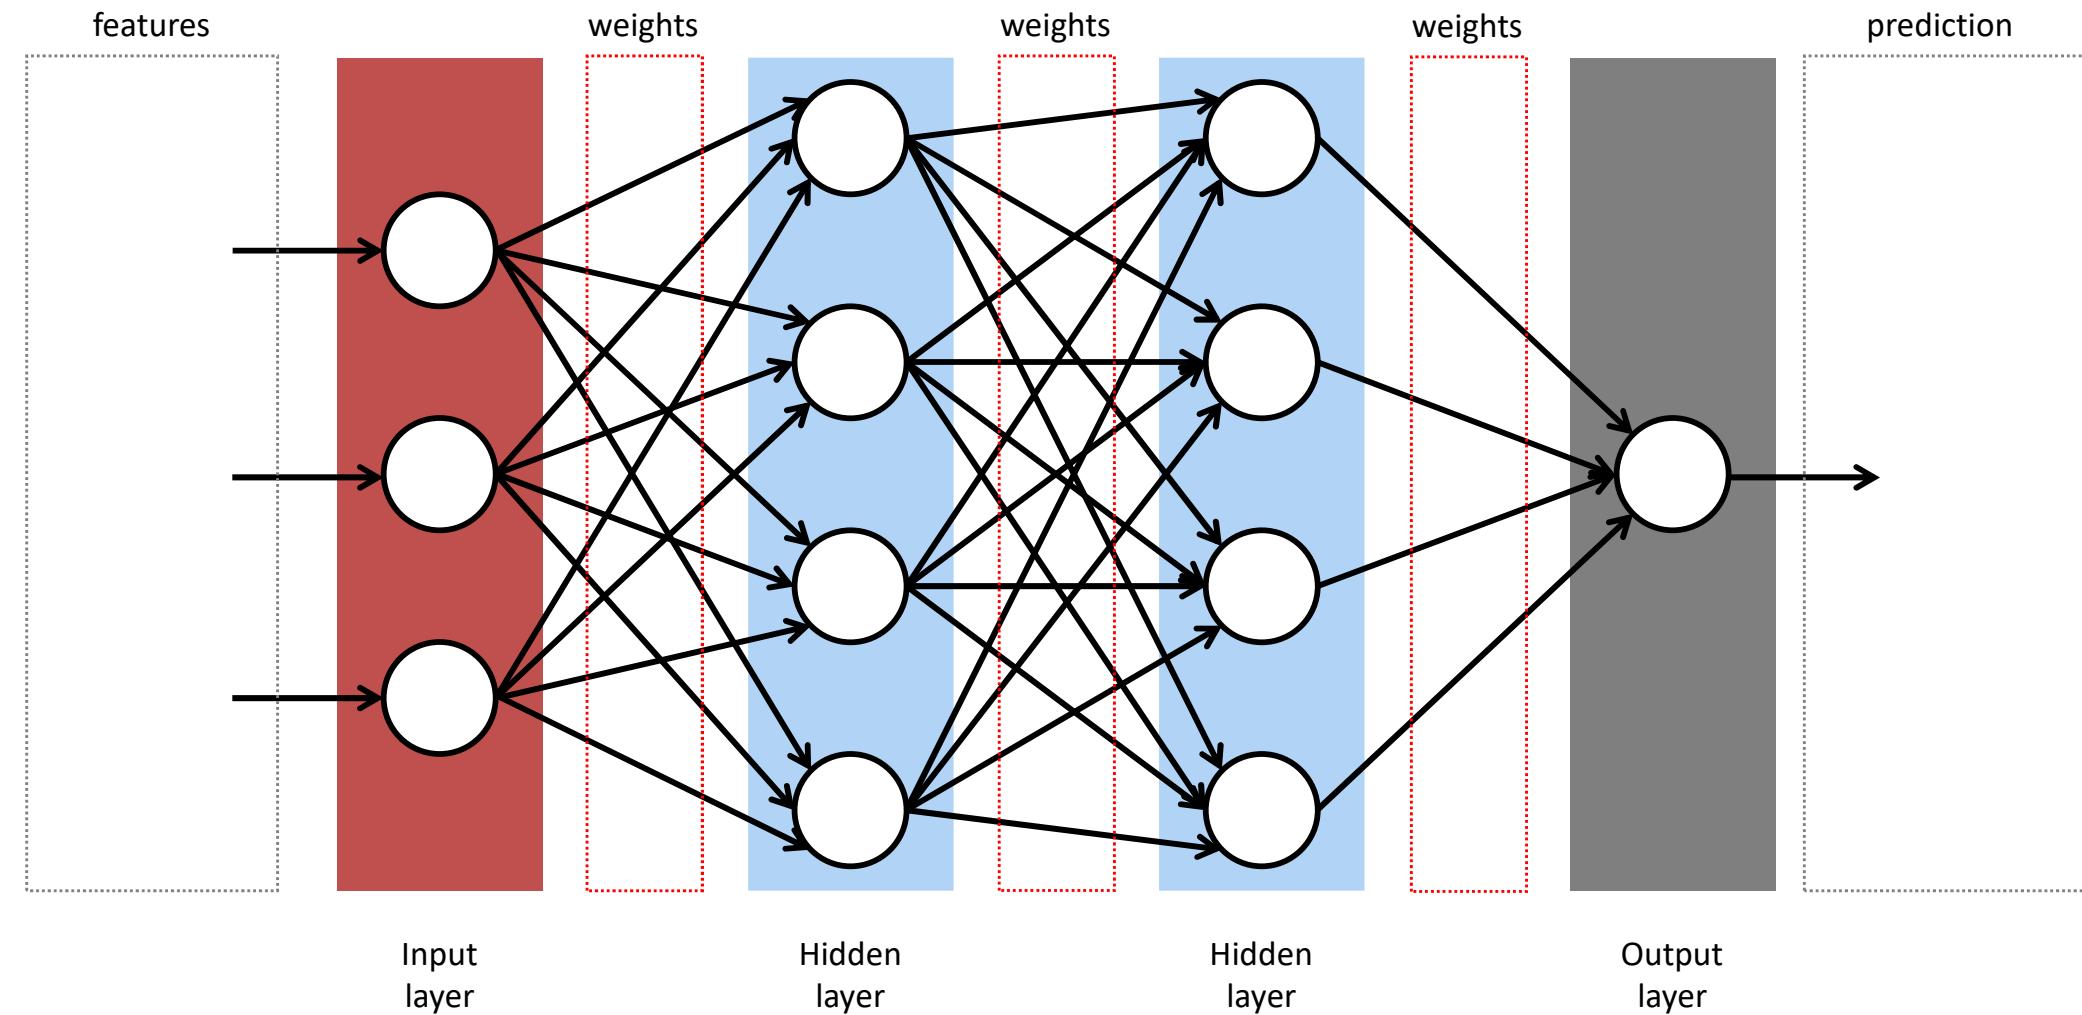
# **Exercise: Train a Classifier!**

**<https://teachablemachine.withgoogle.com/>**

# ANN for Classification



# ANN for Regression



# ANN for Regression: Used Car Price

Used car price predictor: train it first with used **car data - price** pairs.

